

*The Open Group Standard*

**Additional APIs for the Base Specifications Issue 8, Part 2**



Copyright © 2022, The Open Group

The Open Group hereby authorizes you to use this document for any purpose, PROVIDED THAT any copy of this document, or any part thereof, which you make shall retain all copyright and other proprietary notices contained herein.

This document may contain other proprietary notices and copyright information.

Nothing contained herein shall be construed as conferring by implication, estoppel, or otherwise any license or right under any patent or trademark of The Open Group or any third party. Except as expressly provided above, nothing contained herein shall be construed as conferring any license or right under any copyright of The Open Group.

Note that any product, process, or technology in this document may be the subject of other intellectual property rights reserved by The Open Group, and may not be licensed hereunder.

This document is provided “AS IS” WITHOUT WARRANTY OF ANY KIND, EITHER EXPRESSED OR IMPLIED, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY, FITNESS FOR A PARTICULAR PURPOSE, OR NON-INFRINGEMENT. Some jurisdictions do not allow the exclusion of implied warranties, so the above exclusion may not apply to you.

Any publication of The Open Group may include technical inaccuracies or typographical errors. Changes may be periodically made to these publications; these changes will be incorporated in new editions of these publications. The Open Group may make improvements and/or changes in the products and/or the programs described in these publications at any time without notice.

Should any viewer of this document respond with information including feedback data, such as questions, comments, suggestions, or the like regarding the content of this document, such information shall be deemed to be non-confidential and The Open Group shall have no obligation of any kind with respect to such information and shall be free to reproduce, use, disclose, and distribute the information to others without limitation. Further, The Open Group shall be free to use any ideas, concepts, know-how, or techniques contained in such information for any purpose whatsoever including but not limited to developing, manufacturing, and marketing products incorporating such information.

If you did not obtain this copy through The Open Group, it may not be the latest version. For your convenience, the latest version of this publication may be downloaded at [www.opengroup.org/library](http://www.opengroup.org/library).

The Open Group Standard

**Additional APIs for the Base Specifications Issue 8, Part 2**

ISBN: 1-957866-06-2

Document Number: C228

Published by The Open Group, December 2022

Comments relating to the material contained in this document may be submitted to:

The Open Group, Apex Plaza, Forbury Road, Reading, Berkshire, RG1 1AX, United Kingdom

or by electronic mail to:

[ogspeccs@opengroup.org](mailto:ogspeccs@opengroup.org)

## Contents

1	Introduction.....	1
1.1	Scope.....	1
1.2	Relationship to Other Formal Standards.....	1
2	Application Program Interfaces .....	2
2.1	Change Bars.....	2
2.2	Reference Pages.....	2

# Preface

## The Open Group

The Open Group is a global consortium that enables the achievement of business objectives through technology standards. Our diverse membership of more than 870 organizations includes customers, systems and solutions suppliers, tools vendors, integrators, academics, and consultants across multiple industries.

The mission of The Open Group is to drive the creation of Boundaryless Information Flow™ achieved by:

- Working with customers to capture, understand, and address current and emerging requirements, establish policies, and share best practices
- Working with suppliers, consortia, and standards bodies to develop consensus and facilitate interoperability, to evolve and integrate specifications and open source technologies
- Offering a comprehensive set of services to enhance the operational efficiency of consortia
- Developing and operating the industry's premier certification service and encouraging procurement of certified products

Further information on The Open Group is available at [www.opengroup.org](http://www.opengroup.org).

The Open Group publishes a wide range of technical documentation, most of which is focused on development of Standards and Guides, but which also includes white papers, technical studies, certification and testing documentation, and business titles. Full details and a catalog are available at [www.opengroup.org/library](http://www.opengroup.org/library).

## This Document

This document has been prepared by The Open Group Base Working Group. The Open Group Base Working Group is considering submitting a number of additional APIs to the Austin Group as input to the Issue 8 revision of the Base Specifications.

This document contains the second set of these APIs.

## Trademarks

ArchiMate, DirecNet, Making Standards Work, Open O logo, Open O and Check Certification logo, Platform 3.0, The Open Group, TOGAF, UNIX, UNIXWARE, and the Open Brand X logo are registered trademarks and Boundaryless Information Flow, Build with Integrity Buy with Confidence, Commercial Aviation Reference Architecture, Dependability Through Assuredness, Digital Practitioner Body of Knowledge, DPBoK, EMMM, FACE, the FACE logo, FHIM Profile Builder, the FHIM logo, FPB, Future Airborne Capability Environment, IT4IT, the IT4IT logo, O-AA, O-DEF, O-HERA, O-PAS, Open Agile Architecture, Open FAIR, Open Footprint, Open Process Automation, Open Subsurface Data Universe, Open Trusted Technology Provider, OSDU, Sensor Integration Simplified, SOSA, and the SOSA logo are trademarks of The Open Group.

All other brands, company, and product names are used for identification purposes only and may be trademarks that are the sole property of their respective owners.

## Acknowledgements

The Open Group gratefully acknowledges the contribution of the following in the development of this document:

- The Open Group Base Working Group
- The Austin Group

The Open Group gratefully acknowledges the following reviewers who participated in the Company Review of this document:

- Eric Ackermann
- Eric Blake
- Geoff Clare
- Donald W. Cragun
- Andrew Josey
- Nick Stoughton
- Mark Ziegast

# 1 Introduction

---

## 1.1 Scope

The purpose of this document is to define a set of additional APIs for inclusion in the Issue 8 revision of the Base Specifications of the Single UNIX<sup>®</sup> Specification.

The additional APIs proposed by participants in the Austin Group that The Open Group has agreed to sponsor are as follows:

### Header

`<libintl.h>`

### Functions

<i>bind_textdomain_codeset()</i>	<i>getresgid()</i>
<i>bindtextdomain()</i>	<i>getresuid()</i>
<i>dcgettext()</i>	<i>gettext()</i>
<i>dcgettext_l()</i>	<i>gettext_l()</i>
<i>dcngettext()</i>	<i>ngettext()</i>
<i>dcngettext_l()</i>	<i>ngettext_l()</i>
<i>dgettext()</i>	<i>setresgid()</i>
<i>dgettext_l()</i>	<i>setresuid()</i>
<i>dngettext()</i>	<i>textdomain()</i>
<i>dngettext_l()</i>	

### Utilities

<i>gettext</i>	<i>realpath</i>
<i>msgfmt</i>	<i>timeout</i>
<i>ngettext</i>	<i>xgettext</i>
<i>readlink</i>	

## 1.2 Relationship to Other Formal Standards

This Standard is being forwarded to the Austin Group for consideration as input to the Issue 8 revision of the Base Specifications.

## **2 Application Program Interfaces**

---

The following pages are extracted from a complete draft of the Base Specifications in which the proposed changes have been applied, with change bars showing the differences from Issue 8 Draft 2.1. Only pages with technical changes are included – editorial changes such as additions to SEE ALSO and CHANGE HISTORY sections have been omitted (unless they appear on the same page as a technical change). The complete draft is also being made available for reference.

As a consequence of the change to NLSPATH in XBD Section 8.2, a change will also need to be made to the NLSPATH description on all existing utility reference pages. These changes are not included here but will be made during the preparation of Issue 8 Draft 3.

### **2.1 Change Bars**

Changed lines are marked with a '|' in the right-hand margin, new lines with a '+', and deleted lines with a '-'.

Note that sometimes the placement of change bars is slightly inaccurate. In particular, changes may extend into a line following a set of change-barred lines. Also, changes within tables do not have change bars.

### **2.2 Reference Pages**

The reference pages for the new header, function, and utility additions, and pages with related changes follow.



### 3.106 Dot

In the context of naming files, the filename consisting of a single <period> character ('.').

**Note:** In the context of shell special built-in utilities, see *dot* in XCU [Section 2.14](#) (on page 2382).

Pathname Resolution is defined in detail in [Section 4.14](#) (on page 93).

### 3.107 Dot-Dot

The filename consisting solely of two <period> characters ("..").

**Note:** Pathname Resolution is defined in detail in [Section 4.14](#) (on page 93).

### 3.108 Dot-Po File

See *Portable Messages Object Source File* in [Section 3.257](#) (on page 65).

### 3.109 Double-Quote Character

The character '"', also known as <quotation-mark>.

**Note:** The “double” adjective in this term refers to the two strokes in the character glyph. POSIX.1-202x never uses the term “double-quote” to refer to two apostrophes or quotation-marks.

### 3.110 Downshifting

The conversion of an uppercase character that has a single-character lowercase representation into this lowercase representation.

### 3.111 Driver

A module that controls data transferred to and received from devices.

**Note:** Drivers are traditionally written to be a part of the system implementation, although they are frequently written separately from the writing of the implementation. A driver may contain processor-specific code, and therefore be non-portable.

### 3.112 Effective Group ID

An attribute of a process that is used in determining various permissions, including file access permissions; see also [Section 3.161](#) (on page 51).

### 3.113 Effective User ID

An attribute of a process that is used in determining various permissions, including file access permissions; see also [Section 3.394](#) (on page 84).

### 1786 3.197 Message Catalog

1787 In the context of providing natural language messages to the user, a file or storage area  
1788 containing program messages, command prompts, and responses to prompts for a particular  
1789 native language, territory, and codeset.

### 1790 3.198 Message Catalog Descriptor

1791 In the context of providing natural language messages to the user, a per-process unique value  
1792 used to identify an open message catalog. A message catalog descriptor may be implemented  
1793 using a file descriptor.

### 1794 3.199 Message Queue

1795 In the context of programmatic message passing, an object to which messages can be added and  
1796 removed. Messages may be removed in the order in which they were added or in priority order.

### 1797 3.200 Messages Object

1798 A file containing message identifiers and translations in an unspecified format. Used by the +  
1799 *gettext* family of functions and the *gettext* and *ngettext* utilities for internationalization and +  
1800 localization of programs and scripts. Messages objects have the filename suffix **.mo**, and can be +  
1801 created by the *msgfmt* utility. +

1802 See also [Section 3.374](#) (on page 81).

### 1803 3.201 Mode

1804 A collection of attributes that specifies a file's type and its access permissions.

1805 **Note:** File Access Permissions are defined in detail in [Section 4.6](#) (on page 90).

### 1806 3.202 Monotonic Clock

1807 A clock measuring real time, whose value cannot be set via *clock\_settime()* and which cannot  
1808 have negative clock jumps.

### 1809 3.203 Mount Point

1810 Either the system root directory or a directory for which the *st\_dev* field of structure **stat** differs  
1811 from that of its parent directory.

1812 **Note:** The **stat** structure is defined in detail in [<sys/stat.h>](#).

### 1813 3.204 Multi-Character Collating Element

1814 A sequence of two or more characters that collate as an entity. For example, in some coded

### 2023 3.252 Pipe

2024 An object identical to a FIFO which has no links in the file hierarchy.

2025 **Note:** The *pipe()* function is defined in detail in the System Interfaces volume of POSIX.1-202x.

### 2026 3.253 Polling

2027 A scheduling scheme whereby the local process periodically checks until the pre-specified  
2028 events (for example, read, write) have occurred.

### 2029 3.254 Portable Character Set

2030 The collection of characters that are required to be present in all locales supported by  
2031 conforming systems.

2032 **Note:** The Portable Character Set is defined in detail in [Section 6.1](#) (on page 105).

2033 This term is contrasted against the smaller portable filename character set; see also [Section 3.256](#).

### 2034 3.255 Portable Filename

2035 A filename consisting only of characters from the portable filename character set.

2036 **Note:** Applications should avoid using filenames that have the <hyphen-minus> character as the first  
2037 character since this may cause problems when filenames are passed as command line  
2038 arguments.

### 2039 3.256 Portable Filename Character Set

2040 The set of characters from which portable filenames are constructed.

2041 A B C D E F G H I J K L M N O P Q R S T U V W X Y Z  
2042 a b c d e f g h i j k l m n o p q r s t u v w x y z  
2043 0 1 2 3 4 5 6 7 8 9 . \_ -

2044 The last three characters are the <period>, <underscore>, and <hyphen-minus> characters,  
2045 respectively. See also [Section 3.245](#) (on page 63).

### 2046 3.257 Portable Messages Object Source File (or Dot-Po File)

2047 A text file containing messages and directives. A portable messages object source file can be  
2048 compiled into a messages object by the *msgfmt* utility.

2049 **Note:** By convention, portable messages object source files have filenames ending with the **.po** suffix. +  
2050 Utility descriptions in this standard frequently use dot-po file as a shorthand for portable +  
2051 messages object source file (even though the **.po** suffix need not be included in the filename). +  
2052 Template portable messages object source files can be created from C-language source files by +  
2053 the *xgettext* utility. +

### 2477 3.368 System Process

2478 An object other than a process executing an application, that is provided by the system and has a  
2479 process ID.

### 2480 3.369 System Reboot

2481 See *System Boot* defined in [Section 3.362](#) (on page 80).

### 2482 3.370 System-Wide

2483 Pertaining to events occurring in all processes existing in an implementation at a given point in  
2484 time.

### 2485 3.371 Tab Character (<tab>)

2486 A character that in the output stream indicates that printing or displaying should start at the  
2487 next horizontal tabulation position on the current line. It is the character designated by '`\t`' in  
2488 the C language. If the current position is at or past the last defined horizontal tabulation  
2489 position, the behavior is unspecified. It is unspecified whether this character is the exact  
2490 sequence transmitted to an output device by the system to accomplish the tabulation.

### 2491 3.372 Terminal (or Terminal Device)

2492 A character special file that obeys the specifications of the general terminal interface.

2493 **Note:** The General Terminal Interface is defined in detail in [Chapter 11](#) (on page 185).

### 2494 3.373 Text Column

2495 A roughly rectangular block of characters capable of being laid out side-by-side next to other  
2496 text columns on an output page or terminal screen. The widths of text columns are measured in  
2497 column positions.

### 2498 3.374 Text Domain

2499 A named collection of messages objects (one messages object per supported language) for +  
2500 internationalization and localization purposes. A text domain is often named after the +  
2501 application or library that provides the collection, but may have a more general name if it is +  
2502 intended to be shared by multiple applications or libraries. +

2503 **Note:** The use of text domains is defined in detail in the descriptions of the [bindtextdomain\(\)](#) and +  
2504 [gettext](#) family of functions in the System Interfaces volume of POSIX.1-202x. +

### 2505 3.375 Text File

2506 A file that contains characters organized into zero or more lines. The lines do not contain NUL

3607  
3608

3609 **7.1 General**

3610 A locale is the definition of the subset of a user's environment that depends on language and  
3611 cultural conventions. It is made up from one or more categories. Each category is identified by  
3612 its name and controls specific aspects of the behavior of components of the system. Category  
3613 names correspond to the following environment variable names:

- 3614 *LC\_CTYPE* Character classification and case conversion.  
3615 *LC\_COLLATE* Collation order.  
3616 *LC\_MONETARY* Monetary formatting.  
3617 *LC\_NUMERIC* Numeric, non-monetary formatting.  
3618 *LC\_TIME* Date and time formats.  
3619 *LC\_MESSAGES* Formats of informative and diagnostic messages and interactive responses.

3620 The standard utilities in the Shell and Utilities volume of POSIX.1-202x shall base their behavior  
3621 on the current locale, as defined in the ENVIRONMENT VARIABLES section for each utility.  
3622 The behavior of some of the C-language functions defined in the System Interfaces volume of  
3623 POSIX.1-202x shall also be modified based on a locale selection. The locale to be used by these  
3624 functions can be selected in the following ways:

- 3625 1. For functions such as *isalnum\_l()* that take a locale object as an argument, a locale object  
3626 can be obtained from *newlocale()* or *duplocale()* and passed to the function.
- 3627 2. For functions that do not take a locale object as an argument, the current locale for the  
3628 thread can be set by calling *uselocale()* or the global locale for the process can be set by  
3629 calling *setlocale()*. Such functions shall use the current locale of the calling thread if one  
3630 has been set for that thread; otherwise, they shall use the global locale.
- 3631 3. Some functions, such as *catopen()* and those related to text domains, may reference +  
3632 various environment variables and a locale category of a specific locale to access files they +  
3633 need to use.

3634 Locales other than those supplied by the implementation can be created via the *localedef* utility,  
3635 provided that the *\_POSIX2\_LOCALEDEF* symbol is defined on the system. Even if *localedef* is  
3636 not provided, all implementations conforming to the System Interfaces volume of POSIX.1-202x  
3637 shall provide one or more locales that behave as described in this chapter. The input to the  
3638 utility is described in [Section 7.3](#) (on page 116). The value that is used to specify a locale when  
3639 using environment variables shall be the string specified as the *name* operand to the *localedef*  
3640 utility when the locale was created. The strings "C" and "POSIX" are reserved as identifiers for  
3641 the POSIX locale (see [Section 7.2](#), on page 116). When the value of a locale environment variable  
3642 begins with a <slash> (' / '), it shall be interpreted as the pathname of the locale definition; the  
3643 type of file (regular, directory, and so on) used to store the locale definition is implementation-  
3644 defined. If the value does not begin with a <slash>, the mechanism used to locate the locale is  
3645 implementation-defined.

## 8.2 Internationalization Variables

This section describes environment variables that are relevant to the operation of internationalized interfaces described in POSIX.1-202x.

Users may use the following environment variables to announce specific localization requirements to applications. Applications can retrieve this information using the `setlocale()` function to initialize the correct behavior of the internationalized interfaces. The descriptions of the internationalization environment variables describe the resulting behavior only when the application locale is initialized in this way. The use of the internationalization variables by utilities described in the Shell and Utilities volume of POSIX.1-202x is described in the ENVIRONMENT VARIABLES section for those utilities in addition to the global effects described in this section.

**LANG** This variable shall determine the locale category for native language, local customs, and coded character set in the absence of the `LC_ALL` and other `LC_*` (`LC_COLLATE`, `LC_CTYPE`, `LC_MESSAGES`, `LC_MONETARY`, `LC_NUMERIC`, `LC_TIME`) environment variables. This can be used by applications to determine the language to use for error messages and instructions, collating sequences, date formats, and so on.

**LANGUAGE**

XSI

The `LANGUAGE` environment variable shall be examined to determine the messages object to be used for the `gettext` family of functions or the `gettext` and `ngettext` utilities if `NLSPATH` is not set or the evaluation of `NLSPATH` did not lead to a suitable messages object being found. The value of `LANGUAGE` shall be a list of locale names separated by a <colon> (':') character. If `LANGUAGE` is set to a non-empty string, each locale name shall be tried in the specified order and if a messages object is found, it shall be used for translation. If a locale name has the format `language[_territory][.codeset][@modifier]`, additional searches of locale names without `.codeset` (if present), without `_territory` (if present), and without `@modifier` (if present) may be performed; if `.codeset` is not present, additional searches of locale names with an added `.codeset` may be performed. If locale names contain a <slash> ('/') character, or consist entirely of a dot (".") or dot-dot ("..") character sequence, or are empty the behavior is implementation defined and they may be ignored for security reasons.

The locale names in `LANGUAGE` shall override the locale name associated with the "active category" of the current locale or, in the case of functions with an `_l` suffix, the provided locale object, and the language-specific part of the default search path for messages objects, unless the locale name that would be overridden is C or POSIX. For the `dcgettext()`, `dcgettext_l()`, `dcngettext()`, and `dcngettext_l()` functions, the active category is specified by the `category` argument; for all other `gettext` family functions and for the `gettext` and `ngettext` utilities, the active category is `LC_MESSAGES`.

For example, if:

- The `LC_MESSAGES` environment variable is `"de_DE"` (and `LC_ALL` is unset) and `setlocale(LC_ALL, "")` has been used to set the current locale
- The `LANGUAGE` environment variable is `"fr_FR:it"`
- Messages objects are by default searched for in `/gettextlib`

then the following pathnames are tried in this order by `gettext` family functions that have neither a category argument nor an `_l` suffix until a valid messages object is found:

5481		• <code>/gettextlib/fr_FR/LC_MESSAGES/textdomain.mo</code>	+
5482		• (Optionally) <code>/gettextlib/fr/LC_MESSAGES/textdomain.mo</code>	+
5483		• (Optionally) the above two pathnames with added <code>.codeset</code> elements	+
5484		• <code>/gettextlib/it/LC_MESSAGES/textdomain.mo</code>	+
5485		• (Optionally) the above pathname with added <code>.codeset</code> elements	+
5486		• <code>/gettextlib/de_DE/LC_MESSAGES/textdomain.mo</code>	+
5487	<code>LC_ALL</code>	This variable shall determine the values for all locale categories. The value of the <code>LC_ALL</code> environment variable has precedence over any of the other environment variables starting with <code>LC_</code> ( <code>LC_COLLATE</code> , <code>LC_CTYPE</code> , <code>LC_MESSAGES</code> , <code>LC_MONETARY</code> , <code>LC_NUMERIC</code> , <code>LC_TIME</code> ) and the <code>LANG</code> environment variable.	
5488			
5489			
5490			
5491	<code>LC_COLLATE</code>		
5492		This variable shall determine the locale category for character collation. It determines collation information for regular expressions and sorting, including equivalence classes and multi-character collating elements, in various utilities and the <code>strcoll()</code> and <code>strxfrm()</code> functions. Additional semantics of this variable, if any, are implementation-defined.	
5493			
5494			
5495			
5496			
5497	<code>LC_CTYPE</code>	This variable shall determine the locale category for character handling functions, such as <code>tolower()</code> , <code>toupper()</code> , and <code>isalpha()</code> . This environment variable determines the interpretation of sequences of bytes of text data as characters (for example, single as opposed to multi-byte characters), the classification of characters (for example, alpha, digit, graph), and the behavior of character classes. Additional semantics of this variable, if any, are implementation-defined.	
5498			
5499			
5500			
5501			
5502			
5503	<code>LC_MESSAGES</code>		
5504		This variable shall determine the locale category for processing affirmative and negative responses and the language and cultural conventions in which messages should be written. It also affects the behavior of the <code>catopen()</code> function in determining the message catalog. Additional semantics of this variable, if any, are implementation-defined. The language and cultural conventions of diagnostic and informative messages whose format is unspecified by POSIX.1-202x should be affected by the setting of <code>LC_MESSAGES</code> .	
5505			
5506			
5507			
5508			
5509			
5510			
5511	<code>LC_MONETARY</code>		
5512		This variable shall determine the locale category for monetary-related numeric formatting information. Additional semantics of this variable, if any, are implementation-defined.	
5513			
5514			
5515	<code>LC_NUMERIC</code>		
5516		This variable shall determine the locale category for numeric formatting (for example, thousands separator and radix character) information in various utilities as well as the formatted I/O operations in <code>printf()</code> and <code>scanf()</code> and the string conversion functions in <code>strtod()</code> . Additional semantics of this variable, if any, are implementation-defined.	
5517			
5518			
5519			
5520			
5521	<code>LC_TIME</code>	This variable shall determine the locale category for date and time formatting information. It affects the behavior of the time functions in <code>strptime()</code> . Additional semantics of this variable, if any, are implementation-defined.	
5522			
5523			
5524	XSI	<code>NLSPATH</code>	
5525		This variable shall contain a sequence of templates to be used by <code>catopen()</code> when attempting to locate message catalogs, and by the <code>gettext</code> family of functions when locating messages objects. Each template consists of an optional prefix, one or	
5526			

more conversion specifications, and an optional suffix.

The conversion specification descriptions below refer to a “currently active text domain”. The currently active text domain is, in decreasing order of precedence:

- The *domain* parameter of the *gettext* family of functions or the *gettext* and *ngettext* utilities
- The text domain bound by the last call to *textdomain()* when using a *gettext* family function, or the *TEXTDOMAIN* environment variable when using the *gettext* and *ngettext* utilities
- The default text domain

Conversion specifications consist of a '%' symbol, followed by a single-letter keyword. The following conversion specifications are currently defined:

%N The value of the *name* parameter passed to *catopen()* or the currently active text domain of the *gettext* family of functions and the *gettext* and *ngettext* utilities (see above).

%L The locale name given by the value of the active category (see *LANGUAGE* above) in either the current locale or, in the case of functions with an *\_l* suffix, the provided locale object.

%l The *language* element of the locale name that would result from a %L conversion.

%t The *territory* element of the locale name that would result from a %L conversion.

%c The *codeset* element of the locale name that would result from a %L conversion.

%% A single '%' character.

An empty string shall be substituted if the specified value is not currently defined. The separators <underscore> ('\_') and <period> ('.') shall not be included in the %t and %c conversion specifications.

Templates defined in *NLSPATH* are separated by <colon> characters (':'). A leading, trailing, or two adjacent <colon> characters ("::") shall be equivalent to specifying %N.

Since <colon> is a separator in this context, directory names that might be used in *NLSPATH* should not include a <colon> character.

Example 1, for an application that uses *catopen()* but does not use the *gettext* family of functions:

```
NLSPATH="/system/nlslib/%N.cat"
```

indicates that *catopen()* should look for all message catalogs in the directory */system/nlslib*, where the catalog name should be constructed from the *name* argument (replacing %N) passed to *catopen()*, with the suffix *.cat*.

Example 2, for an application that uses the *gettext* family of functions but does not use *catopen()*:

```
NLSPATH="/usr/lib/locale/fr/LC_MESSAGES/%N.mo"
```

indicates that the *gettext* family of functions (and the *gettext* and *ngettext* utilities) should look for all messages objects in the directory */usr/lib/locale/fr/LC\_MESSAGES*, where the messages object's name should be



5570	constructed from the currently active text domain (replacing %N), with the suffix	
5571	<b>.mo</b> .	
5572	Example 3, for an application that uses <i>catopen()</i> but does not use the <i>gettext</i> family	
5573	of functions:	
5574	<code>NLSPATH="%N.cat:/nlslib/%L/%N.cat"</code>	
5575	indicates that <i>catopen()</i> should look for the requested message catalog in <i>name</i> ,	
5576	<i>name.cat</i> , and <code>/nlslib/locale/name.cat</code> , where <i>locale</i> is the locale name given	
5577	by the value of the <code>LC_MESSAGES</code> category of the current locale.	
5578	Example 4, for an application that uses the <i>gettext</i> family of functions but does not	+
5579	use <i>catopen()</i> :	+
5580	<code>NLSPATH="/usr/lib/locale/%L/%N.mo:/usr/lib/locale/fr/%N.mo"</code>	+
5581	indicates that the <i>gettext</i> family of functions (and the <i>gettext</i> and <i>ngettext</i> utilities)	+
5582	should look for all messages objects first in	+
5583	<code>/usr/lib/locale/localename/textdomain.mo</code> , and if not found there, then try in	+
5584	<code>/usr/lib/locale/fr/textdomain.mo</code> , where <i>locale</i> is the locale name given by the	+
5585	value of the active category in the current locale or provided locale object.	+
5586	Example 5, for an application that uses <i>catopen()</i> and the <i>gettext</i> family of	+
5587	functions:	+
5588	<code>NLSPATH="/usr/lib/locale/%L/%N.mo:/system/nlslib/%L/%N.cat"</code>	+
5589	indicates that the <i>gettext</i> family of functions (and the <i>gettext</i> and <i>ngettext</i> utilities)	+
5590	should look for all messages objects in <code>/usr/lib/locale/localename/textdomain.mo</code> ,	+
5591	where <i>locale</i> is the locale name given by the value of the active category in the	+
5592	current locale or provided locale object. Also, <i>catopen()</i> should look for all message	+
5593	catalogs in the directory <code>/system/nlslib/locale/name.cat</code> , (assuming that	+
5594	<code>/usr/lib/locale/localename/name.mo</code> is not a message catalog). In this scenario,	+
5595	<i>catopen()</i> ignores all files that are not valid message catalogs while traversing	+
5596	<code>NLSPATH</code> . Furthermore, the <i>gettext</i> family of functions and the <i>gettext</i> and <i>ngettext</i>	+
5597	utilities ignore all files that are not valid messages objects found while traversing	+
5598	<code>NLSPATH</code> .	+
5599	Users should not set the <code>NLSPATH</code> variable unless they have a specific reason to	
5600	override the default system path. Setting <code>NLSPATH</code> to override the default system	
5601	path may produce undefined results in the standard utilities other than <i>gettext</i> and	
5602	<i>ngettext</i> , and in applications with appropriate privileges.	
5603	Specifying a relative pathname in the <code>NLSPATH</code> environment variable should be	
5604	avoided without a specific reason, including the use of a leading, trailing, or two	
5605	adjacent <colon> characters, since it may result in messages objects being searched	
5606	for in a directory relative to the current working directory of the calling process; if	
5607	the process calls the <i>chdir()</i> function, the directory searched for may also be	
5608	changed.	
5609	<b>TEXTDOMAIN</b>	
5610	Specify the text domain name that the <i>gettext</i> and <i>ngettext</i> utilities use during the	
5611	search for messages objects. This is identical to the messages object filename	
5612	without the <b>.mo</b> suffix.	
5613	<b>TEXTDOMAINDIR</b>	
5614	Specify the pathname to the root directory of the messages object hierarchy the	
5615	<i>gettext</i> and <i>ngettext</i> utilities use during the search for messages objects. If present, it	

5616 XSI shall replace the default root directory pathname. *NLSPATH* has precedence over  
5617 *TEXTDOMAINDIR*.

5618 The environment variables *LANG*, *LC\_ALL*, *LC\_COLLATE*, *LC\_CTYPE*, *LC\_MESSAGES*,  
5619 *LC\_MONETARY*, *LC\_NUMERIC*, *LC\_TIME*, and *NLSPATH* provide for the support of  
5620 internationalized applications. The standard utilities shall make use of these environment  
5621 variables as described in this section and the individual ENVIRONMENT VARIABLES sections  
5622 for the utilities. If these variables specify locale categories that are not based upon the same  
5623 underlying codeset, the results are unspecified.

5624 The values of locale categories shall be determined by a precedence order; the first condition met  
5625 below determines the value:

- 5626 1. If the *LC\_ALL* environment variable is defined and is not null, the value of *LC\_ALL* shall  
5627 be used.
- 5628 2. If the *LC\_\** environment variable (*LC\_COLLATE*, *LC\_CTYPE*, *LC\_MESSAGES*,  
5629 *LC\_MONETARY*, *LC\_NUMERIC*, *LC\_TIME*) is defined and is not null, the value of the  
5630 environment variable shall be used to initialize the category that corresponds to the  
5631 environment variable.
- 5632 3. If the *LANG* environment variable is defined and is not null, the value of the *LANG*  
5633 environment variable shall be used.
- 5634 4. If the *LANG* environment variable is not set or is set to the empty string, the  
5635 implementation-defined default locale shall be used.

5636 If the locale value is "C" or "POSIX", the POSIX locale shall be used and the standard utilities  
5637 behave in accordance with the rules in [Section 7.2](#) (on page 116) for the associated category.

5638 If the locale value begins with a <slash>, it shall be interpreted as the pathname of a file that was  
5639 created in the output format used by the *localedef* utility; see OUTPUT FILES under *localedef*.  
5640 Referencing such a pathname shall result in that locale being used for the indicated category.

5641 XSI If the locale value has the form:

5642 *language[\_territory] [.codeset]*

5643 it refers to an implementation-provided locale, where settings of language, territory, and codeset  
5644 are implementation-defined.

5645 *LC\_COLLATE*, *LC\_CTYPE*, *LC\_MESSAGES*, *LC\_MONETARY*, *LC\_NUMERIC*, and *LC\_TIME* are  
5646 defined to accept an additional field *@modifier*, which allows the user to select a specific instance  
5647 of localization data within a single category (for example, for selecting the dictionary as opposed  
5648 to the character ordering of data). The syntax for these environment variables is thus defined as:

5649 *[language[\_territory] [.codeset] [@modifier]]*

5650 For example, if a user wanted to interact with the system in French, but required to sort German  
5651 text files, *LANG* and *LC\_COLLATE* could be defined as:

5652 *LANG=Fr\_FR*

5653 *LC\_COLLATE=De\_DE*

5654 This could be extended to select dictionary collation (say) by use of the *@modifier* field; for  
5655 example:

5656 *LC\_COLLATE=De\_DE@dict*

5657 An implementation may support other formats.

9068	<b>NAME</b>		
9069	libintl.h — international messaging		+
9070	<b>SYNOPSIS</b>		+
9071	#include <libintl.h>		+
9072	<b>DESCRIPTION</b>		+
9073	The <libintl.h> header may define the macro TEXTDOMAINMAX. If defined, it shall have the		+
9074	same value as {TEXTDOMAIN_MAX} in <limits.h>.		+
9075	The following shall be declared as functions and may also be defined as macros. Function		+
9076	prototypes shall be provided.		+
9077	char *bindtextdomain(const char *, const char *);		+
9078	char *bind_textdomain_codeset(const char *, const char *);		+
9079	char *dcgettext(const char *, const char *, int);		+
9080	char *dcgettext_l(const char *, const char *, int, locale_t);		+
9081	char *dcngettext(const char *, const char *, const char *,		+
9082	unsigned long int, int);		+
9083	char *dcngettext_l(const char *, const char *, const char *,		+
9084	unsigned long int, int, locale_t);		+
9085	char *dgettext(const char *, const char *);		+
9086	char *dgettext_l(const char *, const char *, locale_t);		+
9087	char *dngettext(const char *, const char *, const char *,		+
9088	unsigned long int);		+
9089	char *dngettext_l(const char *, const char *, const char *,		+
9090	unsigned long int, locale_t);		+
9091	char *gettext(const char *);		+
9092	char *gettext_l(const char *, locale_t);		+
9093	char *ngettext(const char *, const char *, unsigned long int);		+
9094	char *ngettext_l(const char *, const char *,		+
9095	unsigned long int, locale_t);		+
9096	char *textdomain(const char *);		+
9097	<b>APPLICATION USAGE</b>		+
9098	None.		+
9099	<b>RATIONALE</b>		+
9100	Some historical implementations defined TEXTDOMAINMAX in this header. This standard		+
9101	instead defines {TEXTDOMAIN_MAX} in <limits.h>. This was done to allow the maximum		+
9102	length of a text domain name to vary depending on the filesystem type used to store message		+
9103	catalogs. Implementations are allowed to continue to define TEXTDOMAINMAX in this header		+
9104	as an extension to the standard (see XSH <a href="#">Section 2.2.2</a> , on page 467).		+
9105	<b>FUTURE DIRECTIONS</b>		+
9106	None.		+
9107	<b>SEE ALSO</b>		+
9108	XSH <a href="#">gettext</a> , <a href="#">bindtextdomain()</a>		+
9109	<b>CHANGE HISTORY</b>		+
9110	First released in Issue 8.		+
9111			+

9288		{PIPE_BUF}	
9289		Maximum number of bytes that is guaranteed to be atomic when writing to a pipe.	
9290		Minimum Acceptable Value: {_POSIX_PIPE_BUF}	
9291	ADV	{POSIX_ALLOC_SIZE_MIN}	
9292		Minimum number of bytes of storage actually allocated for any portion of a file.	
9293		Minimum Acceptable Value: Not specified.	
9294	ADV	{POSIX_REC_INCR_XFER_SIZE}	
9295		Recommended increment for file transfer sizes between the	
9296		{POSIX_REC_MIN_XFER_SIZE} and {POSIX_REC_MAX_XFER_SIZE} values.	
9297		Minimum Acceptable Value: Not specified.	
9298	ADV	{POSIX_REC_MAX_XFER_SIZE}	
9299		Maximum recommended file transfer size.	
9300		Minimum Acceptable Value: Not specified.	
9301	ADV	{POSIX_REC_MIN_XFER_SIZE}	
9302		Minimum recommended file transfer size.	
9303		Minimum Acceptable Value: Not specified.	
9304	ADV	{POSIX_REC_XFER_ALIGN}	
9305		Recommended file transfer buffer alignment.	
9306		Minimum Acceptable Value: Not specified.	
9307		{SYMLINK_MAX}	
9308		Maximum number of bytes in a symbolic link.	
9309		Minimum Acceptable Value: {_POSIX_SYMLINK_MAX}	
9310		{TEXTDOMAIN_MAX}	+
9311		Maximum length of a text domain name, not including the terminating null byte.	+
9312		Minimum Acceptable Value: {_POSIX_NAME_MAX} – 3	+
9313	XSI	Minimum Acceptable Value: {_XOPEN_NAME_MAX} – 3	

### Runtime Increaseable Values

9314			
9315		The magnitude limitations in the following list shall be fixed by specific implementations. An	
9316		application should assume that the value of the symbolic constant defined by <limits.h> in a	
9317		specific implementation is the minimum that pertains whenever the application is run under	
9318		that implementation. A specific instance of a specific implementation may increase the value	
9319		relative to that supplied by <limits.h> for that implementation. The actual value supported by a	
9320		specific instance shall be provided by the <i>sysconf</i> ( ) function.	
9321		{BC_BASE_MAX}	
9322		Maximum <i>obase</i> values allowed by the <i>bc</i> utility.	
9323		Minimum Acceptable Value: {_POSIX2_BC_BASE_MAX}	
9324		{BC_DIM_MAX}	
9325		Maximum number of elements permitted in an array by the <i>bc</i> utility.	
9326		Minimum Acceptable Value: {_POSIX2_BC_DIM_MAX}	
9327		{BC_SCALE_MAX}	
9328		Maximum <i>scale</i> value allowed by the <i>bc</i> utility.	
9329		Minimum Acceptable Value: {_POSIX2_BC_SCALE_MAX}	
9330		{BC_STRING_MAX}	
9331		Maximum length of a string constant accepted by the <i>bc</i> utility.	
9332		Minimum Acceptable Value: {_POSIX2_BC_STRING_MAX}	

```

15290     _PC_PIPE_BUF
15291     _PC_PRIO_IO
15292     _PC_REC_INCR_XFER_SIZE
15293     _PC_REC_MAX_XFER_SIZE
15294     _PC_REC_MIN_XFER_SIZE
15295     _PC_REC_XFER_ALIGN
15296     _PC_SYMLINK_MAX
15297     _PC_SYNC_IO
15298     _PC_TEXTDOMAIN_MAX
15299     _PC_TIMESTAMP_RESOLUTION
15300     _PC_VDISABLE

```

+

15301 The <unistd.h> header shall define the following symbolic constants for *sysconf()*:

```

15302     _SC_2_C_BIND
15303     _SC_2_C_DEV
15304     _SC_2_CHAR_TERM
15305     _SC_2_FORT_RUN
15306     _SC_2_LOCALEDEF
15307     _SC_2_SW_DEV
15308     _SC_2_UPE
15309     _SC_2_VERSION
15310     _SC_ADVISORY_INFO
15311     _SC_AIO_LISTIO_MAX
15312     _SC_AIO_MAX
15313     _SC_AIO_PRIO_DELTA_MAX
15314     _SC_ARG_MAX
15315     _SC_ASYNCCHRONOUS_IO
15316     _SC_ATEXIT_MAX
15317     _SC_BARRIERS
15318     _SC_BC_BASE_MAX
15319     _SC_BC_DIM_MAX
15320     _SC_BC_SCALE_MAX
15321     _SC_BC_STRING_MAX
15322     _SC_CHILD_MAX
15323     _SC_CLK_TCK
15324     _SC_CLOCK_SELECTION
15325     _SC_COLL_WEIGHTS_MAX
15326     _SC_CPU_TIME
15327     _SC_DELAYTIMER_MAX
15328     _SC_EXPR_NEST_MAX
15329     _SC_FSYNC
15330     _SC_GETGR_R_SIZE_MAX
15331     _SC_GETPW_R_SIZE_MAX
15332     _SC_HOST_NAME_MAX
15333     _SC_IOV_MAX
15334     _SC_IPV6
15335     _SC_JOB_CONTROL
15336     _SC_LINE_MAX
15337     _SC_LOGIN_NAME_MAX
15338     _SC_MAPPED_FILES
15339     _SC_MEMLOCK
15340     _SC_MEMLOCK_RANGE

```

```

15484 pid_t      getpid(void);
15485 pid_t      getppid(void);
15486 XSI int      getresgid(gid_t *, gid_t *, gid_t *); +
15487 int      getresuid(uid_t *, uid_t *, uid_t *); +
15488 pid_t      getsid(pid_t);
15489 uid_t      getuid(void);
15490 int      isatty(int);
15491 int      lchown(const char *, uid_t, gid_t);
15492 int      link(const char *, const char *);
15493 int      linkat(int, const char *, int, const char *, int);
15494 XSI int      lockf(int, int, off_t);
15495 off_t      lseek(int, off_t, int);
15496 XSI int      nice(int);
15497 long      pathconf(const char *, int);
15498 int      pause(void);
15499 int      pipe(int [2]);
15500 int      pipe2(int [2], int);
15501 int      posix_close(int, int);
15502 ssize_t    pread(int, void *, size_t, off_t);
15503 ssize_t    pwrite(int, const void *, size_t, off_t);
15504 ssize_t    read(int, void *, size_t);
15505 ssize_t    readlink(const char *restrict, char *restrict, size_t);
15506 ssize_t    readlinkat(int, const char *restrict, char *restrict, size_t);
15507 int      rmdir(const char *);
15508 int      setegid(gid_t);
15509 int      seteuid(uid_t);
15510 int      setgid(gid_t);
15511 int      setpgid(pid_t, pid_t);
15512 XSI int      setregid(gid_t, gid_t);
15513 int      setresgid(gid_t, gid_t, gid_t); +
15514 int      setresuid(uid_t, uid_t, uid_t); +
15515 int      setreuid(uid_t, uid_t);
15516 pid_t      setsid(void);
15517 int      setuid(uid_t);
15518 unsigned   sleep(unsigned);
15519 XSI void     swab(const void *restrict, void *restrict, ssize_t);
15520 int      symlink(const char *, const char *);
15521 int      symlinkat(const char *, int, const char *);
15522 XSI void     sync(void);
15523 long      sysconf(int);
15524 pid_t      tcgetpgrp(int);
15525 int      tcsetpgrp(int, pid_t);
15526 int      truncate(const char *, off_t);
15527 char      *ttyname(int);
15528 int      ttyname_r(int, char *, size_t);
15529 int      unlink(const char *);
15530 int      unlinkat(int, const char *, int);
15531 ssize_t    write(int, const void *, size_t);

15532 The <unistd.h> header shall declare the following external variables:

15533 extern char *optarg;
15534 extern int  opterr, optind, optopt;

```

	Header	Prefix	Suffix	Complete Name
16366	<aio.h>	aio_, lio_, AIO_, LIO_		
16367	<arpa/inet.h>	inet_		
16368	<ctype.h>	to[a-z], is[a-z]		
16369	<dlfcn.h>	RTLD_, dli_		
16370	<dirent.h>	d_, DT_		
16371	<fcntl.h>	l_		
16372	<fmtmsg.h>	MM_		
16373 XSI	<fnmatch.h>	FNM_		
16374	<ftw.h>	FTW		
16375	<glob.h>	gl_, GLOB_		
16376 XSI	<grp.h>	gr_		
16377	<libintl.h>			TEXTDOMAINMAX
16378	<limits.h>		_MAX, _MIN	
16379	<math.h>	M_		
16380	<mqueue.h>	mq_, MQ_		
16381 XSI	<ndbm.h>	dbm_, DBM_		
16382 MSG	<netdb.h>	ai_, h_, n_, p_, s_		
16383 XSI	<net/if.h>	if_, IF_		
16384	<netinet/in.h>	in_, ip_, s_, sin_, INADDR_, IPPROTO_		
16385		in6_, in6addr_, s6_, sin6_, IPV6_		
16386 IP6	<netinet/tcp.h>	TCP_		
16387	<nl_types.h>	NL_		
16388	<poll.h>	pd_, ph_, ps_, POLL		
16389	<pthread.h>	pthread_, PTHREAD_		
16390	<pwd.h>	pw_		
16391	<regex.h>	re_, rm_, REG_		
16392	<sched.h>	sched_, SCHED_		
16393	<semaphore.h>	sem_, SEM_		
16394	<signal.h>	sa_, si_, sigev_, sival_, uc_, BUS_, CLD_, FPE_, ILL_, SA_, SEGV_, SI_, SIGEV_		
16395 CX		ss_, sv_, SS_, TRAP_		
16396 XSI	<stdlib.h>	str[a-z]		
16397	<string.h>	str[a-z], mem[a-z], wcs[a-z]		
16398 XSI	<sys/ipc.h>	ipc_, IPC_		key, pad, seq
16399	<sys/mman.h>	shm_, MAP_, MCL_, MS_, PROT_		
16400	<sys/msg.h>	msg, MSG_[A-Z]		msg
16401 XSI	<sys/resource.h>	rlim_, ru_, PRIO_, RLIMIT_, RUSAGE_		
16402 XSI	<sys/select.h>	fd_, fds_, FD_		

## 18104 2.9.5.2 Cancellation Points

18105 Cancellation points shall occur when a thread is executing the following functions:

18106	<i>accept()</i>	<i>nanosleep()</i>	<i>recvmsg()</i>
18107	<i>accept4()</i>	<i>open()</i>	<i>select()</i>
18108	<i>aio_suspend()</i>	<i>openat()</i>	<i>send()</i>
18109	<i>clock_nanosleep()</i>	<i>pause()</i>	<i>sendmsg()</i>
18110	<i>close()</i>	<i>poll()</i>	<i>sendto()</i>
18111	<i>connect()</i>	<i>ppoll()</i>	<i>sigsuspend()</i>
18112	<i>creat()</i>	<i>pread()</i>	<i>sigtimedwait()</i>
18113	<i>fcntl()</i> †	<i>pselect()</i>	<i>sigwait()</i>
18114	<i>fdasync()</i>	<i>pthread_cond_clockwait()</i>	<i>sigwaitinfo()</i>
18115	<i>fsync()</i>	<i>pthread_cond_timedwait()</i>	<i>sleep()</i>
18116	<i>lockf()</i> ††	<i>pthread_cond_wait()</i>	<i>tcdrain()</i>
18117	<i>mq_receive()</i>	<i>pthread_join()</i>	<i>wait()</i>
18118	<i>mq_send()</i>	<i>pthread_testcancel()</i>	<i>waitid()</i>
18119	<i>mq_timedreceive()</i>	<i>pwrite()</i>	<i>waitpid()</i>
18120	<i>mq_timedsend()</i>	<i>read()</i>	<i>write()</i>
18121	<i>msgrcv()</i>	<i>readv()</i>	<i>writew()</i>
18122	<i>msgsnd()</i>	<i>recv()</i>	
18123	<i>msync()</i>	<i>recvfrom()</i>	

18124 A cancellation point may also occur when a thread is executing the following functions:

18125	<i>access()</i>	<i>dngettext_l()</i>	<i>fgetws()</i>
18126	<i>bindtextdomain()</i>	<i>dprintf()</i>	<i>fntmsg()</i>
18127	<i>catclose()</i>	<i>endhostent()</i>	<i>fopen()</i>
18128	<i>catopen()</i>	<i>endnetent()</i>	<i>fpathconf()</i>
18129	<i>chmod()</i>	<i>endprotoent()</i>	<i>fprintf()</i>
18130	<i>chown()</i>	<i>endservent()</i>	<i>fputc()</i>
18131	<i>closedir()</i>	<i>faccessat()</i>	<i>fputs()</i>
18132	<i>closelog()</i>	<i>fchmod()</i>	<i>fputwc()</i>
18133	<i>ctermid()</i>	<i>fchmodat()</i>	<i>fputws()</i>
18134	<i>dcgettext()</i>	<i>fchown()</i>	<i>fread()</i>
18135	<i>dcgettext_l()</i>	<i>fchownat()</i>	<i>freopen()</i>
18136	<i>dcngettext()</i>	<i>fclose()</i>	<i>fscanf()</i>
18137	<i>dcngettext_l()</i>	<i>fcntl()</i> †††	<i>fseek()</i>
18138	<i>dgettext()</i>	<i>fflush()</i>	<i>fseeko()</i>
18139	<i>dgettext_l()</i>	<i>fgetc()</i>	<i>fsetpos()</i>
18140	<i>dlclose()</i>	<i>fgetpos()</i>	<i>fstat()</i>
18141	<i>dlopen()</i>	<i>fgets()</i>	<i>fstatat()</i>
18142	<i>dngettext()</i>	<i>fgetwc()</i>	<i>ftell()</i>

18143 † When the *cmd* argument is F\_SETLK.

18144 †† When the *function* argument is F\_LOCK.

18145 ††† For any value of the *cmd* argument.



18146	<i>ftello()</i>	+	<i>mkstemp()</i>	<i>rewind()</i>
18147	<i>futimens()</i>		<i>mktime()</i>	<i>rewinddir()</i>
18148	<i>fwprintf()</i>		<i>ngettext()</i>	<i>scandir()</i>
18149	<i>fwrite()</i>		<i>ngettext_l()</i>	<i>scanf()</i>
18150	<i>fwscanf()</i>		<i>opendir()</i>	<i>seekdir()</i>
18151	<i>getaddrinfo()</i>		<i>openlog()</i>	<i>sem_clockwait()</i>
18152	<i>getc()</i>		<i>pathconf()</i>	<i>sem_timedwait()</i>
18153	<i>getc_unlocked()</i>		<i>perror()</i>	<i>sem_wait()</i>
18154	<i>getchar()</i>		<i>popen()</i>	<i>semop()</i>
18155	<i>getchar_unlocked()</i>		<i>posix_fadvise()</i>	<i>sethostent()</i>
18156	<i>getcwd()</i>		<i>posix_fallocate()</i>	<i>setnetent()</i>
18157	<i>getdelim()</i>		<i>posix_getdents()</i>	<i>setprotoent()</i>
18158	<i>getgrgid_r()</i>		<i>posix_madvise()</i>	<i>setservent()</i>
18159	<i>getgrnam_r()</i>		<i>posix_openpt()</i>	<i>stat()</i>
18160	<i>gethostid()</i>		<i>posix_spawn()</i>	<i>strerror_l()</i>
18161	<i>gethostname()</i>		<i>posix_spawnnp()</i>	<i>strerror_r()</i>
18162	<i>getline()</i>		<i>posix_typed_mem_open()</i>	<i>strftime()</i>
18163	<i>getlogin_r()</i>		<i>printf()</i>	<i>strftime_l()</i>
18164	<i>getnameinfo()</i>		<i>psiginfo()</i>	<i>symlink()</i>
18165	<i>getpwnam_r()</i>		<i>psignal()</i>	<i>symlinkat()</i>
18166	<i>getpwuid_r()</i>		<i>pthread_rwlock_clockrdlock()</i>	<i>sync()</i>
18167	<i>gettext()</i>		<i>pthread_rwlock_clockwrlock()</i>	<i>syslog()</i>
18168	<i>gettext_l()</i>		<i>pthread_rwlock_rdlock()</i>	<i>tmpfile()</i>
18169	<i>getwc()</i>		<i>pthread_rwlock_timedrdlock()</i>	<i>tmpnam()</i>
18170	<i>getwchar()</i>		<i>pthread_rwlock_timedwrlock()</i>	<i>ttyname_r()</i>
18171	<i>glob()</i>		<i>pthread_rwlock_wrlock()</i>	<i>tzset()</i>
18172	<i>iconv_close()</i>		<i>ptsname()</i>	<i>ungetc()</i>
18173	<i>iconv_open()</i>		<i>ptsname_r()</i>	<i>ungetwc()</i>
18174	<i>link()</i>		<i>putc()</i>	<i>unlink()</i>
18175	<i>linkat()</i>		<i>putc_unlocked()</i>	<i>unlinkat()</i>
18176	<i>lio_listio()</i>		<i>putchar()</i>	<i>utimensat()</i>
18177	<i>localtime_r()</i>		<i>putchar_unlocked()</i>	<i>utimes()</i>
18178	<i>lockf()</i>		<i>puts()</i>	<i>vdprintf()</i>
18179	<i>lseek()</i>		<i>putwc()</i>	<i>vfprintf()</i>
18180	<i>lstat()</i>		<i>putwchar()</i>	<i>vfswprintf()</i>
18181	<i>mkdir()</i>		<i>readdir_r()</i>	<i>vprintf()</i>
18182	<i>mkdirat()</i>		<i>readlink()</i>	<i>vswprintf()</i>
18183	<i>mkdtemp()</i>		<i>readlinkat()</i>	<i>wcsftime()</i>
18184	<i>mkfifo()</i>		<i>remove()</i>	<i>wordexp()</i>
18185	<i>mkfifoat()</i>		<i>rename()</i>	<i>wprintf()</i>
18186	<i>mknod()</i>		<i>renameat()</i>	<i>wscanf()</i>
18187	<i>mknodat()</i>			

In addition, a cancellation point may occur when a thread is executing any function that this standard does not require to be thread-safe but the implementation documents as being thread-safe. If a thread is cancelled while executing a non-thread-safe function, the behavior is undefined.

An implementation shall not introduce cancellation points into any other functions specified in this volume of POSIX.1-202x.

The side-effects of acting upon a cancellation request while suspended during a call of a function are the same as the side-effects that may be seen in a single-threaded program when a call to a function is interrupted by a signal and the given function returns [EINTR]. Any such side-

## 21422 NAME

21423 bindtextdomain, bind\_textdomain\_codeset, textdomain — text domain manipulation functions +

## 21424 SYNOPSIS

21425 #include <libintl.h> +

21426 char \*bindtextdomain(const char \*domainname, const char \*dirname); +

21427 char \*bind\_textdomain\_codeset(const char \*domainname, +

21428 const char \*codeset); +

21429 char \*textdomain(const char \*domainname); +

## 21430 DESCRIPTION

21431 The *textdomain()* function shall set or query the name of the current text domain of the calling +  
 21432 process. The application shall ensure that the *domainname* argument is either a null pointer +  
 21433 (when querying), an empty string, or a string that, when used by the *gettext* family of functions +  
 21434 to construct a pathname to a messages object, results in a valid pathname. For portable +  
 21435 applications, it should only contain characters from the portable filename character set. +

21436 The text domain setting made by the last successful call to *textdomain()* shall remain in effect +  
 21437 across subsequent calls to *setlocale()*, *uselocale()*, and the *gettext* family of functions. +

21438 Applications should not use text domains whose names begin with the strings "SYS\_" or +  
 21439 "libc". These prefixes are reserved for implementation use. +

21440 The current setting of the text domain can be queried without affecting the current state of the +  
 21441 domain by calling *textdomain()* with *domainname* set to a null pointer. Calling *textdomain()* with a +  
 21442 *domainname* argument of an empty string shall set the text domain to the default domain, +  
 21443 "messages". +

21444 The *bindtextdomain()* function shall set or query the binding of a text domain to a *dirname* that is +  
 21445 used by the *gettext* family of functions to construct a pathname to a messages object in the text +  
 21446 domain: +

- 21447 • If *domainname* is a null pointer or an empty string, *bindtextdomain()* shall make no changes +  
 21448 and return a null pointer without changing *errno*. +

- 21449 • Otherwise, if *dirname* is a non-empty string: +

- 21450 — If *domainname* is not already bound, *bindtextdomain()* shall bind the text domain +  
 21451 specified by *domainname* to the pathname pointed to by *dirname* and return the bound +  
 21452 directory pathname on success or a null pointer on failure. +

- 21453 — If *domainname* is already bound, *bindtextdomain()* shall replace the existing binding +  
 21454 with the pathname pointed to by *dirname* and return the bound directory pathname +  
 21455 on success or a null pointer on failure. On failure, the existing binding shall remain +  
 21456 unchanged. +

21457 It is unspecified whether the *bindtextdomain()* function performs pathname resolution on +  
 21458 *dirname*, or whether that is done by the *gettext* family of functions. +

- 21459 • Otherwise, if *dirname* is a null pointer: +

- 21460 — If *domainname* is bound, the function shall return the bound directory pathname. +

- 21461 — If *domainname* is not bound, the function shall return the implementation-defined +  
 21462 default directory pathname used by the *gettext* family of functions. +

- 21463 • Otherwise, *dirname* is an empty string and the behavior is unspecified. +

21464 If a text domain is bound to a relative pathname and the current working directory is changed +  
 21465 after the binding is established, the pathnames used by the *gettext* family of functions to locate +

21466	messages objects for that text domain are unspecified.	+
21467	The <i>bind_textdomain_codeset()</i> function shall set or query the binding of a text domain to the	+
21468	output codeset used by the <i>gettext</i> family of functions for message strings retrieved from	+
21469	messages objects for the text domain specified by <i>domainname</i> :	+
21470	• If <i>domainname</i> is a null pointer or an empty string, <i>bind_textdomain_codeset()</i> shall make no	+
21471	changes and return a null pointer without changing <i>errno</i> .	+
21472	• Otherwise, if <i>codeset</i> is a non-empty string:	+
21473	— If <i>domainname</i> is not already bound, <i>bind_textdomain_codeset()</i> shall bind the text	+
21474	domain specified by <i>domainname</i> to the codeset pointed to by <i>codeset</i> and return the	+
21475	newly bound codeset on success or a null pointer on failure.	+
21476	— If <i>domainname</i> is already bound, <i>bind_textdomain_codeset()</i> shall replace the existing	+
21477	binding with the codeset pointed to by <i>codeset</i> and return the newly bound codeset	+
21478	on success or a null pointer on failure. On failure, the existing binding shall remain	+
21479	unchanged.	+
21480	The application shall ensure that the <i>codeset</i> argument, if non-empty, is a valid codeset	+
21481	name that can be used as the <i>tocode</i> argument of the <i>iconv_open()</i> function, and that in the	+
21482	codeset it specifies, the <NUL> character corresponds to a single null byte.	+
21483	• Otherwise, if <i>codeset</i> is a null pointer:	+
21484	— If <i>domainname</i> is bound, the function shall return the bound codeset.	+
21485	— If <i>domainname</i> is not bound, the function shall return the implementation-defined	+
21486	default codeset used by the <i>gettext</i> family of functions.	+
21487	• Otherwise, <i>codeset</i> is an empty string and the behavior is unspecified.	+
21488	If <i>codeset</i> is a null pointer and <i>domainname</i> is a non-empty string, <i>bind_textdomain_codeset()</i> shall	+
21489	return the current codeset for the named domain, or a null pointer if a codeset has not yet been	+
21490	set. The <i>bind_textdomain_codeset()</i> function can be called multiple times. If successfully called	+
21491	multiple times with the same <i>domainname</i> argument, the last such call shall override the setting	+
21492	made by the previous such call.	+
21493	<b>RETURN VALUE</b>	+
21494	The return value from a successful <i>textdomain()</i> call shall be a pointer to a string containing the	+
21495	current setting of the text domain. If <i>domainname</i> is a null pointer, <i>textdomain()</i> shall return a	+
21496	pointer to the string containing the current text domain. If <i>textdomain()</i> was not previously	+
21497	called and <i>domainname</i> is a null string, the name of the default text domain shall be returned.	+
21498	The name of the default text domain shall be the string "messages". If <i>textdomain()</i> fails, a null	+
21499	pointer shall be returned and <i>errno</i> shall be set to indicate the error.	+
21500	For <i>bindtextdomain()</i> return values see the DESCRIPTION. When <i>bindtextdomain()</i> is called with	+
21501	a non-empty <i>domainname</i> and returns a null pointer, it shall set <i>errno</i> to indicate the error. When	+
21502	<i>bindtextdomain()</i> returns a pathname for a bound text domain, the return value shall be a pointer	+
21503	to a copy of the <i>dirname</i> string passed to the <i>bindtextdomain()</i> call that created the binding. The	+
21504	returned string shall remain valid until the next successful call to <i>bindtextdomain()</i> with a non-	+
21505	empty <i>dirname</i> and same <i>domainname</i> . The application shall ensure that it does not modify the	+
21506	returned string.	+
21507	A call to the <i>bind_textdomain_codeset()</i> function with a non-empty <i>domainname</i> argument shall	+
21508	return one of the following:	+

21509	• The currently bound codeset name for that text domain if <i>codeset</i> is a null pointer	+
21510	• The newly bound codeset if <i>codeset</i> is non-empty	+
21511	• A null pointer without changing <i>errno</i> if no codeset has yet been bound for that text	+
21512	domain	+
21513	The application shall ensure that it does not modify the returned string. A subsequent call to	+
21514	<i>bind_textdomain_codeset()</i> with a non-empty <i>domainname</i> argument might invalidate the returned	+
21515	pointer or overwrite the string content. The returned pointer might also be invalidated if the	+
21516	calling thread is terminated. If <i>bind_textdomain_codeset()</i> fails, a null pointer shall be returned	+
21517	and <i>errno</i> shall be set to indicate the error.	+
21518	<b>ERRORS</b>	+
21519	For the conditions under which <i>bindtextdomain()</i> —if it performs pathname resolution—fails and	+
21520	may fail, refer to <i>open()</i> .	+
21521	In addition, the <i>textdomain()</i> , <i>bindtextdomain()</i> , and <i>bind_textdomain_codeset()</i> functions may fail	+
21522	if:	+
21523	[ENOMEM]      Insufficient memory available.	+
21524	<b>EXAMPLES</b>	+
21525	See the examples for <i>gettext</i> .	+
21526	<b>APPLICATION USAGE</b>	+
21527	A text <i>domainname</i> is limited to {TEXTDOMAIN_MAX} bytes.	+
21528	Application developers are responsible for ensuring that the text domain used is not used by	+
21529	other applications. To minimize the chances of collision, developers can prefix text domains with	+
21530	their company or application name (or both) and an underscore. For example, if your	+
21531	application name was "foo" and you wanted to use the text domain "errors", you could	+
21532	instead use the text domain "foo_errors". Note that if an application can be installed with a	+
21533	configurable name, a text domain prefix based on the application name should change with the	+
21534	application name.	+
21535	Specifying a relative pathname to the <i>bindtextdomain()</i> function should be avoided, since it may	+
21536	result in messages objects being searched for in a directory relative to the current working	+
21537	directory of the calling process; if the process calls the <i>chdir()</i> function, the directory searched for	+
21538	may also be changed.	+
21539	Since pathname resolution of <i>dirname</i> might not be performed by <i>bindtextdomain()</i> , but could be	+
21540	performed later by the <i>gettext</i> family of functions, and since the latter have no way to report an	+
21541	error, applications should verify, using for example <i>stat()</i> , that the directory is accessible if this is	+
21542	desired.	+
21543	<b>RATIONALE</b>	+
21544	Although the return type of these functions ought to be <b>const char *</b> , it is <b>char *</b> to match	+
21545	historical practice.	+
21546	Pathname resolution of the <i>dirname</i> argument passed to <i>bindtextdomain()</i> may be performed by	+
21547	<i>bindtextdomain()</i> itself or by the <i>gettext</i> family of functions. If pathname resolution fails in one of	+
21548	the <i>gettext</i> family of functions, it is neither allowed to modify <i>errno</i> nor to return an error, but if	+
21549	pathname resolution fails in <i>bindtextdomain()</i> , it is required to report an error and set <i>errno</i> just	+
21550	like <i>open()</i> does.	+
21551	Historically, <i>bindtextdomain()</i> did not perform pathname resolution. However, the standard	+
21552	developers decided to allow this as an option so that future implementations can, if desired,	+
21553	open a file descriptor for that directory in <i>bindtextdomain()</i> and then use that file descriptor with	+

21554	<i>openat()</i> in the <i>gettext</i> family of functions.	+
21555	The <i>dirname</i> parameter to <i>bindtextdomain()</i> may need to be copied to avoid the possibility of the	+
21556	application releasing the memory used by the argument while the <i>gettext</i> family of functions	+
21557	may still need to reference it.	+
21558	When <i>bindtextdomain()</i> is called with a non-empty <i>domainname</i> and an empty <i>dirname</i> , historical	+
21559	implementations of the <i>gettext</i> family of functions use the empty string for the <i>dirname</i> part of	+
21560	the messages object pathname, resulting in an absolute pathname of the form	+
21561	<i>/localename/categoryname/textdomainname.mo</i> . The standard developers did not believe this	+
21562	behavior to be useful. Using the empty <i>dirname</i> case as a way to remove an existing binding	+
21563	seemed to be a more useful behavior, and would be consistent with the behavior of <i>textdomain()</i> .	+
21564	However, because no historical implementations behave this way, the behavior is left	+
21565	unspecified.	+
21566	Some implementations set <i>errno</i> to [EAGAIN] to signal memory allocation failures that might	+
21567	succeed if retried and [ENOMEM] for failures that are unlikely to ever succeed, for example due	+
21568	to configured limits. Section 2.3 (on page 475) permits this behavior; when multiple error	+
21569	conditions are simultaneously true there is no precedence between them.	+
21570	<b>FUTURE DIRECTIONS</b>	+
21571	A future version of this standard may require implementations to prefix implementation-	+
21572	provided text domains with either "SYS_" or a prefix related to the implementor's company	+
21573	name to avoid namespace collisions.	+
21574	A future version of this standard may require <i>bindtextdomain()</i> to remove any binding for	+
21575	<i>domainname</i> when called with a non-empty <i>domainname</i> and an empty <i>dirname</i> .	+
21576	<b>SEE ALSO</b>	+
21577	<i>gettext</i> , <i>iconv_open()</i> , <i>setlocale()</i> , <i>uselocale()</i>	+
21578	XBD <i>&lt;libintl.h&gt;</i> , <i>&lt;limits.h&gt;</i>	+
21579	XCU <i>msgfmt</i> , <i>xgettext</i>	+
21580	<b>CHANGE HISTORY</b>	+
21581	First released in Issue 8.	+
21582		+

22136 **NAME**

22137 catopen — open a message catalog

22138 **SYNOPSIS**

22139 #include &lt;nl\_types.h&gt;

22140 nl\_catd catopen(const char \*name, int oflag);

22141 **DESCRIPTION**

22142 The *catopen()* function shall open a message catalog and return a message catalog descriptor.  
 22143 The *name* argument specifies the name of the message catalog to be opened. If *name* contains a  
 22144 `'/'`, then *name* specifies a pathname for the message catalog. Otherwise, the environment  
 22145 variable *NLSPATH* is used with *name* substituted for the `%N` conversion specification (see XBD  
 22146 Chapter 8, on page 155); if *NLSPATH* exists in the environment when the process starts, then if  
 22147 the process has appropriate privileges, the behavior of *catopen()* is undefined. If *NLSPATH* does  
 22148 not exist in the environment, or if a message catalog cannot be found in any of the components  
 22149 specified by *NLSPATH*, then an implementation-defined default path shall be used. This default  
 22150 may be affected by the setting of *LC\_MESSAGES* if the value of *oflag* is *NL\_CAT\_LOCALE*, or  
 22151 XSI the *LANG* environment variable if *oflag* is 0. When searching *NLSPATH*, *catopen()* shall ignore +  
 22152 any files it finds that are not valid message catalog files.

22153 A message catalog descriptor shall remain valid in a process until that process closes it, or a  
 22154 successful call to one of the *exec* functions. A change in the setting of the *LC\_MESSAGES*  
 22155 category may invalidate existing open catalogs.

22156 If a file descriptor is used to implement message catalog descriptors, the *FD\_CLOEXEC* flag  
 22157 shall be set; see <fcntl.h>.

22158 If the value of the *oflag* argument is 0, the *LANG* environment variable is used to locate the  
 22159 catalog without regard to the *LC\_MESSAGES* category. If the *oflag* argument is  
 22160 *NL\_CAT\_LOCALE*, the *LC\_MESSAGES* category is used to locate the message catalog (see XBD  
 22161 Section 8.2, on page 157).

22162 **RETURN VALUE**

22163 Upon successful completion, *catopen()* shall return a message catalog descriptor for use on  
 22164 subsequent calls to *catgets()* and *catclose()*. Otherwise, *catopen()* shall return (*nl\_catd*) -1 and set  
 22165 *errno* to indicate the error.

22166 **ERRORS**22167 The *catopen()* function may fail if:

22168 [EACCES] Search permission is denied for the component of the path prefix of the  
 22169 message catalog or read permission is denied for the message catalog.

22170 [EMFILE] All file descriptors available to the process are currently open.

22171 [ENAMETOOLONG]  
 22172 The length of a component of a pathname is longer than {NAME\_MAX}.

22173 [ENAMETOOLONG]  
 22174 The length of a pathname exceeds {PATH\_MAX}, or pathname resolution of a  
 22175 symbolic link produced an intermediate result with a length that exceeds  
 22176 {PATH\_MAX}.

22177 [ENFILE] Too many files are currently open in the system.

22178 [ENOENT] The *name* argument contains a `'/'` and does not name an existing message  
 22179 XSI catalog, the *name* argument does not contain a `'/'` and searching *NLSPATH* (if  
 22180 set) and then the implementation-defined default path for a message catalog

22181 with that name failed, one or more files exist but all are of an invalid format, |  
 22182 or the *name* argument points to an empty string.

22183 [ENOMEM] Insufficient storage space is available.

22184 [ENOTDIR] A component of the path prefix of the message catalog names an existing file  
 22185 that is neither a directory nor a symbolic link to a directory, or the pathname  
 22186 of the message catalog contains at least one non-`<slash>` character and ends  
 22187 with one or more trailing `<slash>` characters and the last pathname  
 22188 component names an existing file that is neither a directory nor a symbolic  
 22189 link to a directory.

## 22190 EXAMPLES

22191 None.

## 22192 APPLICATION USAGE

22193 Some implementations of *catopen()* use *malloc()* to allocate space for internal buffer areas. The  
 22194 *catopen()* function may fail if there is insufficient storage space available to accommodate these  
 22195 buffers.

22196 Conforming applications must assume that message catalog descriptors are not valid after a call  
 22197 to one of the *exec* functions.

22198 Application developers should be aware that guidelines for the location of message catalogs  
 22199 have not yet been developed. Therefore they should take care to avoid conflicting with catalogs  
 22200 used by other applications and the standard utilities.

22201 To be sure that messages produced by an application running with appropriate privileges cannot  
 22202 be used by an attacker setting an unexpected value for *NLSPATH* in the environment to confuse  
 22203 a system administrator, such applications should use pathnames containing a `'/'` to get defined  
 22204 behavior when using *catopen()* to open a message catalog.

## 22205 RATIONALE

22206 None.

## 22207 FUTURE DIRECTIONS

22208 None.

## 22209 SEE ALSO

22210 *catclose()*, *catgets()*

22211 XBD Chapter 8 (on page 155), *<fcntl.h>*, *<nl\_types.h>*,

## 22212 CHANGE HISTORY

22213 First released in Issue 2.

## 22214 Issue 7

22215 Austin Group Interpretation 1003.1-2001 #143 is applied.

22216 SD5-XBD-ERN-4 is applied, changing the definition of the [EMFILE] error.

22217 The *catopen()* function is moved from the XSI option to the Base.

22218 POSIX.1-2008, Technical Corrigendum 1, XSH/TC1-2008/0045 [324] is applied.

22219 POSIX.1-2008, Technical Corrigendum 2, XSH/TC2-2008/0054 [645], XSH/TC2-2008/0055 [497],  
 22220 and XSH/TC2-2008/0056 [497] are applied.

25023 **NAME**

25024 dcgettext, dcgettext\_l, dcngettext, dcngettext\_l, dgettext, dgettext\_l — message handling  
25025 functions

25026 **SYNOPSIS**

```
25027 #include <libintl.h>

25028 char *dcgettext(const char *domainname, const char *msgid,
25029                int category);
25030 char *dcgettext_l(const char *domainname, const char *msgid,
25031                  int category, locale_t locale);
25032 char *dcngettext(const char *domainname, const char *msgid,
25033                  const char *msgid_plural, unsigned long int n,
25034                  int category);
25035 char *dcngettext_l(const char *domainname, const char *msgid,
25036                    const char *msgid_plural, unsigned long int n,
25037                    int category, locale_t locale);
25038 char *dgettext(const char *domainname, const char *msgid);
25039 char *dgettext_l(const char *domainname, const char *msgid,
25040                  locale_t locale);
```

25041 **DESCRIPTION**

25042 Refer to [gettext](#).



25654 **NAME**

25655 dngettext, dngettext\_l — message handling functions

25656 **SYNOPSIS**

25657 #include &lt;libintl.h&gt;

```
25658 char *dngettext(const char *domainname, const char *msgid,  
25659                const char *msgid_plural, unsigned long int n);  
25660 char *dngettext_l(const char *domainname, const char *msgid,  
25661                  const char *msgid_plural, unsigned long int n,  
25662                  locale_t locale);
```

25663 **DESCRIPTION**25664 Refer to *gettext*.

**NAME**

fpathconf, pathconf — get configurable pathname variables

**SYNOPSIS**

```
#include <unistd.h>
```

```
long fpathconf(int fildes, int name);
```

```
long pathconf(const char *path, int name);
```

**DESCRIPTION**

The *fpathconf()* and *pathconf()* functions shall determine the current value of a configurable limit or option (*variable*) that is associated with a file or directory.

For *pathconf()*, the *path* argument points to the pathname of a file or directory.

For *fpathconf()*, the *fildes* argument is an open file descriptor.

The *name* argument represents the variable to be queried relative to that file or directory. Implementations shall support all of the variables listed in the following table and may support others. The variables in the following table come from **<limits.h>** or **<unistd.h>** and the symbolic constants, defined in **<unistd.h>**, are the corresponding values used for *name*.

Variable	Value of <i>name</i>	Requirements
{FILESIZEBITS}	_PC_FILESIZEBITS	4,7
{LINK_MAX}	_PC_LINK_MAX	1
{MAX_CANON}	_PC_MAX_CANON	2
{MAX_INPUT}	_PC_MAX_INPUT	2
{NAME_MAX}	_PC_NAME_MAX	3,4
{PATH_MAX}	_PC_PATH_MAX	4,5
{PIPE_BUF}	_PC_PIPE_BUF	6
{POSIX2_SYMLINKS}	_PC_2_SYMLINKS	4
{POSIX_ALLOC_SIZE_MIN}	_PC_ALLOC_SIZE_MIN	10
{POSIX_REC_INCR_XFER_SIZE}	_PC_REC_INCR_XFER_SIZE	10
{POSIX_REC_MAX_XFER_SIZE}	_PC_REC_MAX_XFER_SIZE	10
{POSIX_REC_MIN_XFER_SIZE}	_PC_REC_MIN_XFER_SIZE	10
{POSIX_REC_XFER_ALIGN}	_PC_REC_XFER_ALIGN	10
{SYMLINK_MAX}	_PC_SYMLINK_MAX	4,9
{TEXTDOMAIN_MAX}	_PC_TEXTDOMAIN_MAX	3,4
_POSIX_CHOWN_RESTRICTED	_PC_CHOWN_RESTRICTED	7
_POSIX_NO_TRUNC	_PC_NO_TRUNC	3,4
_POSIX_VDISABLE	_PC_VDISABLE	2
_POSIX_ASYNC_IO	_PC_ASYNC_IO	8
_POSIX_FALLOC	_PC_FALLOC	8
_POSIX_PRIO_IO	_PC_PRIO_IO	8
_POSIX_SYNC_IO	_PC_SYNC_IO	8
_POSIX_TIMESTAMP_RESOLUTION	_PC_TIMESTAMP_RESOLUTION	1

**NAME**

getresgid — get real group ID, effective group ID, and saved set-group-ID

**SYNOPSIS**

XSI `#include <unistd.h>`

```
int getresgid(gid_t *rgid, gid_t *egid, gid_t *sgid);
```

**DESCRIPTION**

The `getresgid()` function shall store the real group ID, effective group ID, and saved set-group-ID of the calling process in the locations pointed to by the arguments *rgid*, *egid*, and *sgid*, respectively.

**RETURN VALUE**

Upon successful completion, 0 shall be returned. Otherwise, -1 shall be returned and *errno* set to indicate the error.

**ERRORS**

No errors are defined.

**EXAMPLES**

None.

**APPLICATION USAGE**

None.

**RATIONALE**

None.

**FUTURE DIRECTIONS**

None.

**SEE ALSO**

*exec*, *getegid()*, *geteuid()*, *getgid()*, *getresuid()*, *getuid()*, *setegid()*, *seteuid()*, *setgid()*, *setregid()*, *setresgid()*, *setresuid()*, *setreuid()*, *setuid()*

XBD `<unistd.h>`

**CHANGE HISTORY**

First released in Issue 8.

**NAME**

getresuid — get real user ID, effective user ID, and saved set-user-ID

**SYNOPSIS**

XSI `#include <unistd.h>`

```
int getresuid(uid_t *ruid, uid_t *euid, uid_t *suid);
```

**DESCRIPTION**

The *getresuid()* function shall store the real user ID, effective user ID, and saved set-user-ID of the calling process in the locations pointed to by the arguments *ruid*, *euid*, and *suid*, respectively.

**RETURN VALUE**

Upon successful completion, 0 shall be returned. Otherwise, -1 shall be returned and *errno* set to indicate the error.

**ERRORS**

No errors are defined.

**EXAMPLES**

None.

**APPLICATION USAGE**

None.

**RATIONALE**

None.

**FUTURE DIRECTIONS**

None.

**SEE ALSO**

*exec*, *getegid()*, *geteuid()*, *getgid()*, *getresgid()*, *getuid()*, *setegid()*, *seteuid()*, *setgid()*, *setregid()*, *setresgid()*, *setresuid()*, *setreuid()*, *setuid()*

XBD `<unistd.h>`

**CHANGE HISTORY**

First released in Issue 8.

**NAME**

`dgettext`, `dgettext_l`, `dcgettext`, `dcgettext_l`, `gettext`, `gettext_l`, `ngettext`, `ngettext_l`, `dngettext`, `dngettext_l`, `dcngettext`, `dcngettext_l` — message handling functions

**SYNOPSIS**

```
#include <libintl.h>

char *dgettext(const char *domainname, const char *msgid);
char *dgettext_l(const char *domainname, const char *msgid,
    locale_t locale);
char *dcgettext(const char *domainname, const char *msgid,
    int category);
char *dcgettext_l(const char *domainname, const char *msgid,
    int category, locale_t locale);
char *dngettext(const char *domainname, const char *msgid,
    const char *msgid_plural, unsigned long int n);
char *dngettext_l(const char *domainname, const char *msgid,
    const char *msgid_plural, unsigned long int n,
    locale_t locale);
char *dcngettext(const char *domainname, const char *msgid,
    const char *msgid_plural, unsigned long int n,
    int category);
char *dcngettext_l(const char *domainname, const char *msgid,
    const char *msgid_plural, unsigned long int n,
    int category, locale_t locale);
char *gettext(const char *msgid);
char *gettext_l(const char *msgid, locale_t locale);
char *ngettext(const char *msgid, const char *msgid_plural,
    unsigned long int n);
char *ngettext_l(const char *msgid, const char *msgid_plural,
    unsigned long int n, locale_t locale);
```

**DESCRIPTION**

The `gettext()` function shall:

- attempt to locate a suitable messages object (described in detail below) for the `LC_MESSAGES` category in the current locale, and for the current text domain (see [`bindtextdomain\(\)`](#)), containing the string identified by `msgid`,
- retrieve the string identified by `msgid` from the messages object,
- convert the string to the output codeset if necessary (described in detail below), and
- return the result.

If the locale name in effect is "POSIX" or "C" (i.e. the name associated with the `LC_MESSAGES` locale category in the current locale), or if no suitable messages object exists, or if no string identified by `msgid` exists in the messages object, or if an error occurs, `msgid` shall be returned.

The `dgettext()` function shall be equivalent to `gettext()`, except `domainname` shall be used instead of the current text domain to locate the messages object.

The `dcgettext()` function shall be equivalent to `dgettext()`, except the locale category identified by `category` shall be used instead of `LC_MESSAGES`.

The `ngettext()` function shall be equivalent to `gettext()`, except:

- The string to retrieve shall be identified by a combination of *msgid* and *n* (see *msgfmt*).
- If the locale name in effect is "POSIX" or "C", or if no suitable messages object exists, or if no string identified by the combination of *msgid* and *n* exists in the messages object, or if an error occurs, the return value shall be *msgid* if *n* is 1, otherwise *msgid\_plural*.

The *dngettext()* function shall be equivalent to *ngettext()*, except *domainname* shall be used instead of the current text domain to locate the messages object.

The *dcngettext()* function shall be equivalent to *dngettext()*, except the locale category identified by *category* shall be used instead of *LC\_MESSAGES*.

The *\*\_l()* functions shall be equivalent to their counterparts without the *\_l* suffix, except *locale* shall be used instead of the current locale. If *locale* is the special locale object *LC\_GLOBAL\_LOCALE* or is not a valid locale object handle, the behavior is undefined.

The application shall ensure that the *msgid* and *msgid\_plural* arguments are strings. If either *msgid* or *msgid\_plural* is an empty string, or contains characters not in the portable character set, the results are unspecified. If the *category* argument is *LC\_ALL*, the results are unspecified.

The location of the messages object shall be determined according to the following criteria, stopping when the first messages object is found:

1. If the *NLSPATH* environment variable is set to a non-empty string, an *NLSPATH* search shall be performed as described in XBD Section 8.2 (on page 157). If *NLSPATH* identifies more than one template to use, each template in turn shall be used until a valid messages object is found.
2. If the *LANGUAGE* environment variable is set to a non-empty string, a *LANGUAGE* search shall be performed as described below. If *LANGUAGE* identifies more than one directory to search, each directory shall be searched until a valid messages object is found.
3. A single-locale search shall be performed as described below.

For the *NLSPATH* search and the single-locale search, the single locale name used to locate the messages object shall be the locale name associated with the selected locale category from the current locale, or the provided locale object if calling one of the *\*\_l()* functions; additional searches of locale names without *.codeset* (if present), without *\_territory* (if present), and without *@modifier* (if present) may be performed.

For the *LANGUAGE* search, the value of the *LANGUAGE* environment variable shall be a list of one or more locale names separated by a <colon> (' : ') character. Each locale name shall be tried in the specified order. If a messages object for the locale does not exist, or cannot be opened, or is unsuitable for implementation-defined reasons (such as security), the next locale name (if any) shall be tried. If:

- a messages object for the locale can be opened but cannot be processed without error, or
- the messages object does not contain a string identified by *msgid*, or *msgid* and *n* for the *ngettext* functions,

it is unspecified whether the next locale name (if any) is tried. In all other cases, the messages object for the locale shall be used.

For each locale name in *LANGUAGE*, or if *LANGUAGE* is not set or is empty, or no suitable messages object is found in processing *LANGUAGE*, the pathname used to locate the messages object shall be *dirname/localename/categoryname/textdomainname.mo*, where:

- The *dirname* part is the *dirname* argument of the most recent successful call to *bindtextdomain()* that had *textdomainname* as the *domainname* argument; any trailing <slash> characters in *dirname* shall be discarded. If a successful call to *bindtextdomain()* has not been made for *textdomainname*, an implementation-defined default directory shall be used.
- For the *LANGUAGE* search, the *localename* part is each locale name from *LANGUAGE* in turn; if a locale name has the format *language[\_territory][.codeset][@modifier]*, additional searches of locale names without *.codeset* (if present), without *\_territory* (if present), and without *@modifier* (if present) may be performed; if *.codeset* is not present, additional searches of locale names with an added *.codeset* may be performed. For the single-locale search, the *localename* part is the name of the current locale, or the locale specified in an *\*\_l()* function call, for the category named by *categoryname*. Spellings of codeset names are not standardized, and implementations may attempt to use different commonly known spellings, for example "utf8" and "UTF-8".
- The *categoryname* part is the string "LC\_MESSAGES" if *gettext()*, *dgettext()*, *ngettext()*, or *dngettext()* is called, or the locale category name corresponding to the *category* argument to *dcgettext()* or *dcngettext()*. Likewise for the *\*\_l()* variants of all these functions.
- For *gettext()*, *gettext\_l()*, *ngettext()*, and *ngettext\_l()*, the *textdomainname* part is the text domain set by the last successful call to *textdomain()*. For *dgettext()*, *dcgettext()*, *dngettext()*, *dcngettext()*, and the *\*\_l()* variants of these functions, *textdomainname* is the text domain specified by the *domainname* argument. The *domainname* argument shall be equivalent in syntax and meaning to the *domainname* argument to *textdomain()*, except that the selection of the text domain shall affect only the *dgettext()*, *dcgettext()*, *dngettext()*, and *dcngettext()* function calls and their *\*\_l()* variants. If the *domainname* argument is a null pointer, the text domain set by the last successful call to *textdomain()* shall be used. For all of these functions, if a successful call to *textdomain()* has not been made the default text domain "messages" shall be used.

Resolution of the messages object pathname shall be performed the first time one of the *gettext* family of functions is called for a given combination of *dirname*, *localename*, *categoryname*, and *textdomainname*. It is unspecified whether the pathname is re-resolved if the combination has been used before in a call to one of the *gettext* family of functions. If *bindtextdomain()* performs pathname resolution of its *dirname* argument, only the part of the messages object pathname after *dirname* shall be resolved by the *gettext* family of functions.

When one of the *gettext* family of functions returns a message string that was found in a messages object, it shall convert the codeset of the message string to the output codeset if a codeset is specified in the messages object (see *msgfmt*) and the output codeset is not the same as that codeset. If a successful call to *bind\_textdomain\_codeset()* has been made with the text domain of the messages object as the *domainname* argument and a non-null *codeset* argument, the output codeset shall be the *codeset* argument from the most recent such call. Otherwise, the output codeset shall be the codeset of characters in the current locale, or the provided locale object if calling one of the *\*\_l()* functions, as specified by the *LC\_CTYPE* category of the locale. The conversion shall be performed as if by a call to *iconv()* using a conversion descriptor returned by *iconv\_open(<output codeset>, <messages object codeset>)*, except that if the return value of *iconv()* would be greater than zero, the non-identical conversions performed by the *gettext* family of functions need not be the same as those that such an *iconv()* call would perform. If an error prevents the codeset conversion from being performed, the *gettext* family of functions shall behave as if no message string was found in the messages object. If at least one non-identical conversion is performed that results in a fallback character (one that does not provide any information about the character it was converted from, for example, a <question-mark> or "replacement-character"), the *gettext* family of functions may behave as if no message string was

37764 found in the messages object.

## 37765 RETURN VALUE

37766 The `gettext()`, `gettext_l()`, `dgettext()`, `dgettext_l()`, `dcgettext()`, and `dcgettext_l()` functions shall  
37767 return the message string described in DESCRIPTION if successful. Otherwise, they shall return  
37768 `msgid`.

37769 The `ngettext()`, `ngettext_l()`, `dngettext()`, `dngettext_l()`, `dcngettext()`, and `dcngettext_l()` functions  
37770 shall return the message string described in DESCRIPTION if successful. Otherwise, `msgid` shall  
37771 be returned if *n* is equal to 1, or `msgid_plural` if *n* is not equal to 1.

37772 The application shall ensure that it does not modify the returned string. A subsequent call to a  
37773 `gettext` family function shall not overwrite or invalidate the returned string. The returned string  
37774 may be invalidated by a subsequent call to `bind_textdomain_codeset()`, `bindtextdomain()`,  
37775 `setlocale()`, or `textdomain()` in the same process, except for calls that only query values. The  
37776 returned string shall not be invalidated by a subsequent call to `uselocale()`.

## 37777 ERRORS

37778 The `gettext` family of functions shall not modify `errno`. If an error occurs these functions shall  
37779 return a string as described in RETURN VALUE.

## 37780 EXAMPLES

37781 The example code below assumes the following:

- 37782 • The implementation-defined default directory is `/system/gettextlib`.
- 37783 • The following locales are available on the target system: `en_US`, `en_GB`, `de_DE`. The  
37784 codeset used for all of these locales is UTF-8.
- 37785 • The `en_AU` locale is not available on the target system.
- 37786 • The target system supports conversion from ISO/IEC 8859-1 to UTF-8.
- 37787 • The codeset used for the POSIX locale is ASCII.
- 37788 • The target system does not support conversion from ISO/IEC 8859-1 to ASCII.

37789 Furthermore, the following `.mo` files (and only the following `.mo` files) are installed:

- 37790 • `/system/gettextlib/en_US/LC_MESSAGES/mail.mo`
- 37791 • `/messagecatalogs/example/en_US/LC_MESSAGES/mail.mo`

37792 These are compiled from a portable messages object source file (dot-po file) with the following  
37793 ISO/IEC 8859-1 encoded contents (see the EXTENDED DESCRIPTION of the *msgfmt* utility for a  
37794 description of the dot-po file format):

```
37795 msgid ""
37796 msgstr ""
37797 "Content-Type: text/plain; charset=ISO_8859-1\n"
37798 "Plural-Forms: nplurals=4; plural= n==1?0: (n>1&&n<10)?1: (n==0)?2:3;\n"
37799 msgid "recipient"
37800 msgid_plural "recipients"
37801 msgstr[0] "1 recipient"
37802 msgstr[1] "2 to 9 recipients"
37803 msgstr[2] "no recipients"
37804 msgstr[3] "more than 9 recipients"
```

37805 `/system/gettextlib/de_DE/LC_MESSAGES/mail.mo` is compiled from a dot-po file with the  
37806 following ISO/IEC 8859-1 encoded contents:



```

37807 msgid ""
37808 msgstr ""
37809 "Content-Type: text/plain; charset=ISO_8859-1\n"
37810 "Plural-Forms: nplurals=4; plural= n==1?0: (n>1&& n<5)?1: (n==0)?2:3;\n"
37811 msgid "recipient"
37812 msgid_plural "recipients"
37813 msgstr[0] "1 Empfänger"
37814 msgstr[1] "2 bis 4 Empfänger"
37815 msgstr[2] "keine Empfänger"
37816 msgstr[3] "mehr als 4 Empfänger"

37817 /messagecatalogs/example/en_GB/LC_MESSAGES/mail.mo is compiled from a dot-po file
37818 with the following ISO/IEC 8859-1 encoded contents:

37819 msgid ""
37820 msgstr ""
37821 "Content-Type: text/plain; charset=ISO_8859-1\n"
37822 "Plural-Forms: nplurals=4; plural= n==1?0: (n>1&& n<5)?1: (n==0)?2:3;\n"
37823 msgid "recipient"
37824 msgid_plural "recipients"
37825 msgstr[0] "1 recipient"
37826 msgstr[1] "2 to 4 recipients"
37827 msgstr[2] "no recipients"
37828 msgstr[3] "5 or more recipients"

37829 /messagecatalogs/example2/en_US/LC_MESSAGES/othermail.mo is not a suitable messages
37830 object file or is a suitable messages object file that does not contain the msgid "recipient".

37831 The following example demonstrates the interactions between bindtextdomain(),
37832 bind_textdomain_codeset(), textdomain(), and the gettext family of functions.

37833 unsigned long n_recipients;
37834 // strdup() is used to prevent default_domain from being invalidated by
37835 // a future call to bindtextdomain()
37836 const char *default_domain = strdup(bindtextdomain("mail", NULL));
37837 setlocale(LC_MESSAGES, "POSIX");
37838 setlocale(LC_CTYPE, "POSIX");

37839 n_recipients = 1;
37840 // The following outputs "recipient" with the same encoding as the
37841 // "recipient" argument to ngettext():
37842 printf("%s\n", ngettext("recipient", "recipients", n_recipients));

37843 n_recipients = 3;
37844 // The following outputs "recipients" with the same encoding as the
37845 // "recipients" argument to ngettext():
37846 printf("%s\n", ngettext("recipient", "recipients", n_recipients));

37847 setlocale(LC_MESSAGES, "en_US");
37848 setlocale(LC_CTYPE, "en_US");
37849 textdomain("mail");

37850 n_recipients = 1;
37851 // The following outputs "1 recipient", encoded in UTF-8:
37852 printf("%s\n", ngettext("recipient", "recipients", n_recipients));

37853 n_recipients = 3;

```

```
37854 // The following outputs "2 to 9 recipients", encoded in UTF-8:
37855 printf("%s\n", ngettext("recipient", "recipients", n_recipients));

37856 setlocale(LC_MESSAGES, "en_GB");
37857 setlocale(LC_CTYPE, "en_GB");
37858 bindtextdomain("mail", "/messagecatalogs/example/");

37859 n_recipients = 3;
37860 // The following outputs "2 to 4 recipients", encoded in UTF-8:
37861 printf("%s\n", ngettext("recipient", "recipients", n_recipients));

37862 setlocale(LC_MESSAGES, "en_US");
37863 setlocale(LC_CTYPE, "en_US");
37864 textdomain("othermail");
37865 bindtextdomain("othermail", "/messagecatalogs/example2/");

37866 n_recipients = 3;
37867 // The following outputs "recipients" with the same encoding as the
37868 // "recipients" argument to ngettext():
37869 printf("%s\n", ngettext("recipient", "recipients", n_recipients));

37870 // Because there is no locale named en_AU on the system, en_US is used:
37871 setenv("LANGUAGE", "en_AU:en_US:en_GB", 1);
37872 setlocale(LC_MESSAGES, "");
37873 setlocale(LC_CTYPE, "");
37874 bindtextdomain("mail", default_domain);

37875 // The following outputs "2 to 9 recipients", encoded in UTF-8:
37876 printf("%s\n", dngettext("mail", "recipient", "recipients", 3));

37877 textdomain("mail");
37878 bind_textdomain_codeset("mail", "UTF-8");
37879 setlocale(LC_MESSAGES, "de_DE");
37880 setlocale(LC_CTYPE, "de_DE");
37881 // Clear the LANGUAGE environment variable, otherwise it would take
37882 // precedence over the locale set above, and en_US would continue to
37883 // be used.
37884 setenv("LANGUAGE", "", 1);

37885 n_recipients = 1;
37886 // The following outputs "1 Empfänger", encoded in UTF-8:
37887 printf("%s\n", ngettext("recipient", "recipients", n_recipients));

37888 bind_textdomain_codeset("mail", "ASCII");
37889 setlocale(LC_CTYPE, "POSIX");

37890 n_recipients = 1;
37891 // The following outputs "recipient" with the same encoding as the
37892 // "recipient" argument to ngettext() - remember, the system is assumed
37893 // to not support conversion from ISO/IEC 8859-1 to ASCII:
37894 printf("%s\n", ngettext("recipient", "recipients", n_recipients));

37895 free(default_domain);
```

## APPLICATION USAGE

These functions do not impose a limit on message length. Note that translated strings typically have a different length than the input strings, possibly much longer, and applications using these translations in formatted text (for example, aligned columns for a table) should take that into account.

The *dcgettext()*, *dcgettext\_l()*, *dcngettext()*, and *dcngettext\_l()* functions are useful to retrieve locale-specific strings for a category other than *LC\_MESSAGES*. For example, they can be used to obtain a time format string from the *LC\_TIME* category; because the locale setting of *LC\_TIME* and *LC\_MESSAGES* can be different, using the other *gettext* family functions in such a case might cause an undesired result. All of the functions in the *gettext* family of functions, except *dcgettext()*, *dcgettext\_l()*, *dcngettext()*, and *dcngettext\_l()*, search for messages objects only in the *LC\_MESSAGES* category.

Implementations typically, but are not required to, *mmap()* the messages object file the first time one of the *gettext* family of functions is called, and keep that map in place until it is no longer expected to be used. For example, a successful call to *bindtextdomain()* will typically cause the next call to one of the *gettext* family of functions to *munmap()* the previous file and *mmap()* the new file. Applications should not rely on this behavior, however: the implementation is allowed to cache previously used maps, or not use *mmap()* at all and reopen the file each time one of the *gettext* family of functions is called.

The *msgid* and *msgid\_plural* arguments are typically in (US) English. The arguments are always used in the POSIX or C locale, and when a *gettext* family function encounters an error, so they should not be abstract message identifiers (for example, "message 123") and they should only use characters in the portable character set (to avoid outputting byte sequences that are not valid characters in the current output codeset). If the *xgettext* utility is used to extract the *msgid* and *msgid\_plural* arguments from C source files into a template dot-po file, the arguments must be string literals in order for the resulting file to be useful to translators.

The strings returned by the *gettext* family of functions are not guaranteed to contain only characters that are valid in the current output codeset. In particular, byte sequences that do not form valid characters can occur when:

- The *msgid* or *msgid\_plural* arguments use characters outside the portable character set.
- The messages object file does not specify a character set and uses characters outside the portable character set.

The strings returned by the *gettext* family of functions are guaranteed to remain valid until invalidated as described in the RETURN VALUE section. This includes strings that are created by codeset conversion; those strings are freed by the implementation, not the application. Thus, it is safe to call *gettext* family functions multiple times in situations such as:

```
printf("%s %s\n", gettext("foo"), gettext("bar"));
```

## RATIONALE

Although the return type of these functions ought to be **const char \***, it is **char \*** to match historical practice.

The *gettext* family of functions is frequently used in reporting errors. In fact, it is possible to have an application that attempts to create an error message that combines a translated string via *gettext()* with an error string provided by *strerror()*. The standard requires that the *gettext* family of functions does not modify *errno*, so that an application need not worry about complications of providing sequencing points to capture a stable value of *errno* prior to the translation of the error message, and so that the user will still get a somewhat useful string (even if it is the untranslated original string) on any failure.

37943 There are no wide character equivalents for these functions; historically no implementation is  
37944 known to exist, and the multi-byte message returned from these functions can, in most instances,  
37945 be converted to wide characters by the application if desired.

37946 Some historical *gettext* implementations returned the translated string from the messages object  
37947 without codeset conversion if *iconv\_open()* fails. This is considered to be a bug in those  
37948 implementations.

#### 37949 **FUTURE DIRECTIONS**

37950 None.

#### 37951 **SEE ALSO**

37952 *bindtextdomain()*, *catopen()*, *iconv()*, *setlocale()*, *uselocale()*

37953 XBD **<libintl.h>**, **<limits.h>**

37954 XCU *gettext*, *msgfmt*, *xgettext*

#### 37955 **CHANGE HISTORY**

37956 First released in Issue 8.

37957

47012 **NAME**

47013       ngettext, ngettext\_l — message handling functions

47014 **SYNOPSIS**

47015       #include &lt;libintl.h&gt;

47016       char \*ngettext(const char \*msgid, const char \*msgid\_plural,

47017                   unsigned long int n);

47018       char \*ngettext\_l(const char \*msgid, const char \*msgid\_plural,

47019                   unsigned long int n, locale\_t locale);

47020 **DESCRIPTION**47021       Refer to *gettext*.

**NAME**

setresgid — set real group ID, effective group ID, and saved set-group-ID

**SYNOPSIS**

XSI `#include <unistd.h>`

```
int setresgid(gid_t rgid, gid_t egid, gid_t sgid);
```

**DESCRIPTION**

The *setresgid()* function shall set the real group ID, effective group ID, and saved set-group-ID of the calling process to the values specified by *rgid*, *egid*, and *sgid*, respectively.

If an argument is  $-1$ , the corresponding ID shall not be changed.

Only a process with appropriate privileges can set the real group ID, effective group ID, and saved set-group-ID to any valid value.

A non-privileged process can set its real group ID, effective group ID, and saved set-group-ID, each to one of the values that it currently holds in its real group ID, effective group ID, or saved set-group-ID.

The real group ID, effective group ID, and saved set-group-ID can be set to different values in the same call.

Any supplementary group IDs of the calling process shall remain unchanged.

**RETURN VALUE**

Upon successful completion, 0 shall be returned. Otherwise,  $-1$  shall be returned and *errno* set to indicate the error, and none of the IDs shall be changed.

**ERRORS**

The *setresgid()* function shall fail if:

[EINVAL] The value of the *rgid*, *egid*, or *sgid* argument is invalid or out-of-range.

[EPERM] The calling process does not have appropriate privileges and an attempt was made to change the real group ID, effective group ID, or saved set-group-ID to a value that is not currently present in one of those IDs.

**EXAMPLES**

None.

**APPLICATION USAGE**

None.

**RATIONALE**

None.

**FUTURE DIRECTIONS**

None.

**SEE ALSO**

*exec*, *getegid()*, *geteuid()*, *getgid()*, *getresgid()*, *getresuid()*, *getuid()*, *setegid()*, *seteuid()*, *setgid()*, *setregid()*, *setresuid()*, *setreuid()*, *setuid()*

XBD `<unistd.h>`

61822	<b>CHANGE HISTORY</b>
61823	First released in Issue 8.
61824	

**NAME**

setresuid — set real user ID, effective user ID, and saved set-user-ID

**SYNOPSIS**

```
#include <unistd.h>
```

```
int setresuid(uid_t ruid, uid_t euid, uid_t suid);
```

**DESCRIPTION**

The *setresuid()* function shall set the real user ID, effective user ID, and saved set-user-ID of the calling process to the values specified by *ruid*, *euid*, and *suid*, respectively.

If an argument is  $-1$ , the corresponding ID shall not be changed.

Only a process with appropriate privileges can set the real user ID, effective user ID, and saved set-user-ID to any valid value.

A non-privileged process can set its real user ID, effective user ID, and saved set-user-ID, each to one of the values that it currently holds in its real user ID, effective user ID, or saved set-user-ID.

The real user ID, effective user ID, and saved set-user-ID can be set to different values in the same call.

**RETURN VALUE**

Upon successful completion, 0 shall be returned. Otherwise,  $-1$  shall be returned and *errno* set to indicate the error, and none of the IDs shall be changed.

**ERRORS**

The *setresuid()* function shall fail if:

- |          |   |
|----------|---|
| [EINVAL] | The value of the <i>ruid</i> , <i>euid</i> , or <i>suid</i> argument is invalid or out-of-range.  |
| [EPERM]  | The calling process does not have appropriate privileges and an attempt was made to change the real user ID, effective user ID, or saved set-user-ID to a value that is not currently present in one of those IDs or an attempt was made to change the real user ID to a value not permitted by the implementation. |

**EXAMPLES**

None.

**APPLICATION USAGE**

None.

**RATIONALE**

None.

**FUTURE DIRECTIONS**

None.

**SEE ALSO**

*exec*, *getegid()*, *geteuid()*, *getgid()*, *getresgid()*, *getresuid()*, *getuid()*, *setegid()*, *seteuid()*, *setgid()*, *setregid()*, *setresgid()*, *setreuid()*, *setuid()*

XBD [<unistd.h>](#)

**CHANGE HISTORY**

First released in Issue 8.



69197 **NAME**  
69198       textdomain — text domain manipulation function  
  
69199 **SYNOPSIS**  
69200       #include <libintl.h>  
  
69201       char \*textdomain(const char \*domainname);  
  
69202 **DESCRIPTION**  
69203       Refer to *bindtextdomain()*.

|  
|  
|  
|  
|  
|

**84828 STDOUT**

84829 When the `-v` option is specified, standard output shall be formatted as:

84830 `"%s\n", <pathname or command>`

84831 When the `-V` option is specified, standard output shall be formatted as:

84832 `"%s\n", <unspecified>`

**84833 STDERR**

84834 The standard error shall be used only for diagnostic messages.

**84835 OUTPUT FILES**

84836 None.

**84837 EXTENDED DESCRIPTION**

84838 None.

**84839 EXIT STATUS**

84840 When the `-v` or `-V` options are specified, the following exit values shall be returned:

84841 0 Successful completion.

84842 >0 The *command\_name* could not be found or an error occurred.

84843 Otherwise, the following exit values shall be returned:

84844 126 The utility specified by *command\_name* was found but could not be invoked.

84845 127 An error occurred in the *command* utility or the utility specified by *command\_name* could not  
84846 be found.

84847 Otherwise, the exit status of *command* shall be that of the simple command specified by the  
84848 arguments to *command*.

**84849 CONSEQUENCES OF ERRORS**

84850 Default.

**84851 APPLICATION USAGE**

84852 This utility is required to be intrinsic. See [Section 1.7](#) (on page 2336) for details.

84853 The order for command search allows functions to override regular built-ins and path searches.  
84854 This utility is necessary to allow functions that have the same name as a utility to call the utility  
84855 (instead of a recursive call to the function).

84856 The system default path is available using *getconf*; however, since *getconf* may need to have the  
84857 *PATH* set up before it can be called itself, the following can be used:

84858 `command -p getconf PATH`

84859 There are some advantages to suppressing the special characteristics of special built-ins on  
84860 occasion. For example:

84861 `command exec > unwritable-file`

84862 does not cause a non-interactive script to abort, so that the output status can be checked by the  
84863 script.

84864 The *command*, *env*, *nohup*, *time*, *timeout*, and *xargs* utilities have been specified to use exit code 127 +  
84865 if an error occurs so that applications can distinguish “failure to find a utility” from “invoked  
84866 utility exited with an error indication”. The value 127 was chosen because it is not commonly  
84867 used for other meanings; most utilities use small values for “normal error conditions” and the  
84868 values above 128 can be confused with termination due to receipt of a signal. The value 126 was

88560 XSI **NLSPATH** Determine the location of message catalogs for the processing of *LC\_MESSAGES*.  
 88561 **PATH** Determine the location of the *utility*, as described in XBD [Chapter 8](#) (on page 155).  
 88562 If *PATH* is specified as a *name=value* operand to *env*, the *value* given shall be used in  
 88563 the search for *utility*.

## 88564 ASYNCHRONOUS EVENTS

88565 Default.

## 88566 STDOUT

88567 If no *utility* operand is specified, each *name=value* pair in the resulting environment shall be  
 88568 written in the form:

88569 "*%s=%s\n*", *<name>*, *<value>*

88570 If the *utility* operand is specified, the *env* utility shall not write to standard output.

## 88571 STDERR

88572 The standard error shall be used only for diagnostic messages.

## 88573 OUTPUT FILES

88574 None.

## 88575 EXTENDED DESCRIPTION

88576 None.

## 88577 EXIT STATUS

88578 If *utility* is invoked, the exit status of *env* shall be the exit status of *utility*; otherwise, the *env*  
 88579 utility shall exit with one of the following values:

88580 0 The *env* utility completed successfully.

88581 1–125 An error occurred in the *env* utility.

88582 126 The utility specified by *utility* was found but could not be invoked.

88583 127 The utility specified by *utility* could not be found.

## 88584 CONSEQUENCES OF ERRORS

88585 Default.

## 88586 APPLICATION USAGE

88587 The *command*, *env*, *nice*, *nohup*, *time*, *timeout*, and *xargs* utilities have been specified to use exit +  
 88588 code 127 if an error occurs so that applications can distinguish “failure to find a utility” from  
 88589 “invoked utility exited with an error indication”. The value 127 was chosen because it is not  
 88590 commonly used for other meanings; most utilities use small values for “normal error  
 88591 conditions” and the values above 128 can be confused with termination due to receipt of a  
 88592 signal. The value 126 was chosen in a similar manner to indicate that the utility could be found,  
 88593 but not invoked. Some scripts produce meaningful error messages differentiating the 126 and  
 88594 127 cases. The distinction between exit codes 126 and 127 is based on KornShell practice that  
 88595 uses 127 when all attempts to *exec* the utility fail with [ENOENT], and uses 126 when any  
 88596 attempt to *exec* the utility fails for any other reason.

88597 Historical implementations of the *env* utility use the *execvp()* or *execlp()* functions defined in the  
 88598 System Interfaces volume of POSIX.1-202x to invoke the specified utility; this provides better  
 88599 performance and keeps users from having to escape characters with special meaning to the shell.  
 88600 Therefore, shell functions, special built-ins, and built-ins that are only provided by the shell are  
 88601 not found by this type of *env* implementation. However, *env* can be implemented as a shell built-  
 88602 in, in which case it may be able to execute shell functions and built-ins. An application wishing  
 88603 to ensure execution of a non-built-in utility can use *exec* in a subshell for this purpose.

94098 **NAME**

94099        `gettext, ngettext` — retrieve text string from messages object +

94100 **SYNOPSIS** +

94101        `gettext [-e|-E] [-d textdomain] [textdomain] msgid` +

94102        `gettext [-e|-E] [-n] -s [-d textdomain] msgid...` +

94103        `ngettext [-e|-E] [-d textdomain] [textdomain] msgid msgid_plural n` +

94104 **DESCRIPTION** +

94105        The *gettext* and *ngettext* utilities shall write to standard output the message string(s) that would +  
 94106        result from the following calls to functions defined in the System Interfaces volume of +  
 94107        POSIX.1-202x: +

```
94108        if (textdomainname == NULL || textdomainname[0] == '\0') +
94109            message_string = msgid; +
94110        else { +
94111            setlocale(LC_ALL, ""); +
94112            if (textdomainindir != NULL) +
94113                bindtextdomain(textdomainname, textdomainindir); +
94114            if (msgid_plural == NULL) +
94115                message_string = dgettext(textdomainname, msgid); +
94116            else +
94117                message_string = dngettext(textdomainname, msgid, msgid_plural, n); +
94118        }
```

94119        where: +

94120        • The *textdomainindir* variable is a string containing the value of the *TEXTDOMAINDIR* +  
 94121        environment variable, if set and not empty, or is NULL otherwise. +

94122        • The *textdomainname* variable is a string containing the text domain name obtained from, in +  
 94123        decreasing order of precedence: +

94124            — The optional operand *textdomain*, if present +

94125            — The *-d textdomain* option, if specified +

94126            — The *TEXTDOMAIN* environment variable, if set and not empty +

94127        If the text domain name cannot be obtained from these sources, the *textdomainname* +  
 94128        variable is NULL. +

94129        • If the *-s* option of *gettext* is not specified and for the *ngettext* utility, the *msgid* variable is a +  
 94130        string containing: +

94131            — The value of the *msgid* operand, if the *-E* option is specified +

94132            — The value of the *msgid* operand with C-language escape sequences processed (see +  
 94133            below), if the *-e* option is specified +

94134            — The value of the *msgid* operand with C-language escape sequences optionally +  
 94135            processed (see below), otherwise +

94136        • If the *-s* option of *gettext* is specified, the *msgid* variable is a string containing: +

94137            — The value of each *msgid* operand in turn, if the *-E* option is specified or neither the *-e* +  
 94138            nor the *-E* option is specified +

94139	— The value of each <i>msgid</i> operand in turn with C-language escape sequences	+
94140	processed (see below), if the <b>-e</b> option is specified	+
94141	• For the <i>gettext</i> utility, the <i>msgid_plural</i> variable is NULL. For the <i>ngettext</i> utility, the	+
94142	<i>msgid_plural</i> variable is a string containing:	+
94143	— The value of the <i>msgid_plural</i> operand, if the <b>-E</b> option is specified	+
94144	— The value of the <i>msgid_plural</i> operand with C-language escape sequences processed	+
94145	(see below), if the <b>-e</b> option is specified	+
94146	— The value of the <i>msgid_plural</i> operand with C-language escape sequences optionally	+
94147	processed (see below), otherwise	+
94148	• For the <i>gettext</i> utility, the <i>n</i> variable is 1 (one). For the <i>ngettext</i> utility the <i>n</i> variable is the <i>n</i>	+
94149	operand, parsed as an integer as if by using the <i>strtoul()</i> function with a <i>base</i> argument of	+
94150	10.	+
94151	When C-language escape sequences are processed, they shall be processed as specified for	+
94152	character string literals in the ISO C standard, except that <i>universal-character-name</i> escape	+
94153	sequences need not be supported. Implementations may also support a <backslash> 'c' escape	+
94154	sequence; if supported, the '\c' and all characters following it shall be removed and, if the <b>-s</b>	+
94155	option is specified, the behavior shall be as if the <b>-n</b> option is also specified.	+
94156	For the <i>ngettext</i> utility, and for the <i>gettext</i> utility if the <b>-s</b> option is not specified, the resulting	+
94157	message string shall be written to standard output. If the <b>-s</b> option of <i>gettext</i> is specified, the	+
94158	resulting message string for each <i>msgid</i> shall be written to standard output with consecutive	+
94159	message strings separated by a single <space> character and, if the <b>-n</b> option is not specified, a	+
94160	<newline> shall be written after the last message string. If the <b>-s</b> and <b>-n</b> options are specified,	+
94161	the trailing <newline> shall be omitted.	+
94162	Under conditions where the <i>textdomainname</i> variable in the above code would be NULL, these	+
94163	utilities may write a diagnostic message to standard error and exit with non-zero status.	+
94164	<b>OPTIONS</b>	+
94165	These utilities shall conform to XBD <a href="#">Section 12.2</a> (on page 201).	+
94166	The following options shall be supported:	+
94167	<b>-d</b> <i>textdomain</i>	+
94168	Retrieve the translated message from the domain <i>textdomain</i> , if <i>textdomain</i> is not	+
94169	specified as an operand.	+
94170	<b>-e</b> Process C-language escape sequences in <i>msgid</i> and <i>msgid_plural</i> operands.	+
94171	<b>-E</b> Do not process C-language escape sequences in <i>msgid</i> and <i>msgid_plural</i> operands.	+
94172	The <i>gettext</i> utility shall also support the following options:	+
94173	<b>-n</b> Modify the behavior of the <b>-s</b> option such that a <newline> is not appended to the	+
94174	output.	+
94175	<b>-s</b> Separate the message strings obtained from each <i>msgid</i> operand with <space>	+
94176	characters in the output, and (if <b>-n</b> is not also specified) append a <newline> to the	+
94177	output.	+
94178	If neither of the mutually exclusive <b>-e</b> and <b>-E</b> options is specified, it is unspecified which is the	+
94179	default, except that if the <b>-s</b> option of <i>gettext</i> is specified then <b>-E</b> shall be the default.	+

94180	<b>OPERANDS</b>		+
94181		The following operands shall be supported:	+
94182	<i>textdomain</i>	A text domain name used to retrieve the translated message. This shall override	+
94183		the specification by the <b>-d</b> option, if present.	+
94184	<i>msgid</i>	A key to retrieve the translated message.	+
94185	<i>msgid_plural</i>	A default plural if no corresponding plural message can be found.	+
94186	<i>n</i>	A non-negative decimal integer to be used as the <i>n</i> argument to <i>dngettext()</i> (see the	+
94187		DESCRIPTION).	+
94188	<b>STDIN</b>		+
94189		Not used.	+
94190	<b>INPUT FILES</b>		+
94191		The input files are messages object files (see <i>msgfmt</i> ).	+
94192	<b>ENVIRONMENT VARIABLES</b>		+
94193		The following environment variables shall affect the execution of <i>gettext</i> and <i>ngettext</i> :	+
94194	<i>LANG</i>	Provide a default value for the internationalization variables that are unset or null.	+
94195		(See XBD <a href="#">Section 8.2</a> (on page 157) for the precedence of internationalization	+
94196		variables used to determine the values of locale categories.)	+
94197	XSI <i>LANGUAGE</i>	Determine the location of messages objects if <i>NLSPATH</i> is not set or the evaluation	+
94198		of <i>NLSPATH</i> did not lead to a suitable messages object being found.	+
94199	<i>LC_ALL</i>	If set to a non-empty string value, override the values of all the other	+
94200		internationalization variables.	+
94201	<i>LC_MESSAGES</i>		+
94202		Determine the locale name used to locate messages objects, and the locale that	+
94203		should be used to affect the format and contents of diagnostic messages written to	+
94204		standard error.	+
94205	XSI <i>NLSPATH</i>	Determine the location of messages objects and message catalogs.	+
94206	<i>TEXTDOMAIN</i>		+
94207		Specify the text domain name. (See XBD <a href="#">Section 3.374</a> (on page 81).)	+
94208	<i>TEXTDOMAINDIR</i>		+
94209	XSI	Specify the pathname to the messages object hierarchy. <i>NLSPATH</i> shall have	+
94210		precedence over <i>TEXTDOMAINDIR</i> .	+
94211	<b>ASYNCHRONOUS EVENTS</b>		+
94212		Default.	+
94213	<b>STDOUT</b>		+
94214		See DESCRIPTION.	+
94215	<b>STDERR</b>		+
94216		The standard error shall be used only for diagnostic messages.	+
94217	<b>OUTPUT FILES</b>		+
94218		None.	+

94219	<b>EXTENDED DESCRIPTION</b>	+
94220	None.	+
94221	<b>EXIT STATUS</b>	+
94222	The following exit values shall be returned:	+
94223	0 Successful completion.	+
94224	>0 An error occurred.	+
94225	<b>CONSEQUENCES OF ERRORS</b>	+
94226	Default.	+
94227	<b>APPLICATION USAGE</b>	+
94228	Since it is unspecified which of the <b>-e</b> or <b>-E</b> options is the default, except when the <b>-s</b> option of	+
94229	<i>gettext</i> is specified, portable applications need to ensure that <b>-e</b> , <b>-E</b> , or (for <i>gettext</i> ) <b>-s</b> is specified	+
94230	whenever a <i>msgid</i> or <i>msgid_plural</i> operand contains, or might contain, a <backslash> character.	+
94231	Note that, unless the <b>-s</b> option of <i>gettext</i> is specified without <b>-n</b> , the message(s) written to	+
94232	standard output are not followed by a <newline>. (Therefore the output only ends with a	+
94233	<newline> if the last message ends with one.)	+
94234	Both <i>msgid</i> and <i>msgid_plural</i> should be properly quoted for the shell.	+
94235	<b>EXAMPLES</b>	+
94236	The following examples assume that the following portable messages object source file (dot-po	+
94237	file) has been compiled to a valid file <b>mail.mo</b> by the <i>msgfmt</i> utility. See the EXTENDED	+
94238	DESCRIPTION section of the <i>msgfmt</i> utility for a description of the dot-po file format.	+
94239	<code>msgid ""</code>	+
94240	<code>msgstr ""</code>	+
94241	<code>"Content-Type: text/plain; charset=utf-8\n"</code>	+
94242	<code>"Plural-Forms: nplurals=4; plural=n==1?0: (n&gt;1&amp;&amp;n&lt;=10)?1: (n==0)?2:3;\n"</code>	+
94243	<code>msgid "recipient"</code>	+
94244	<code>msgid_plural "recipients"</code>	+
94245	<code>msgstr[0] "1 recipient"</code>	+
94246	<code>msgstr[1] "2 to 10 recipients"</code>	+
94247	<code>msgstr[2] "no recipients"</code>	+
94248	<code>msgstr[3] "more than 10 recipients"</code>	+
94249	<code>msgid "%d attachment\n"</code>	+
94250	<code>msgid_plural "%d attachments\n"</code>	+
94251	<code>msgstr[0] "1 (%d) attachment\n"</code>	+
94252	<code>msgstr[1] "2 to 10 (%d) attachments\n"</code>	+
94253	<code>msgstr[2] "no (%d) attachments\n"</code>	+
94254	<code>msgstr[3] "more than 10 (%d) attachments\n"</code>	+
94255	They also assume that <b>mail.mo</b> is installed in the directory that <i>gettext</i> and <i>ngettext</i> search for the	+
94256	current locale. See the OPTIONS and ENVIRONMENT VARIABLES sections above and the	+
94257	description of <i>gettext()</i> for details on how this search is performed.	+
94258	The command	+
94259	<code>ngettext -d mail recipient recipients 0</code>	+
94260	will write "no recipients".	+
94261	The command	+

```

94262      ngettext -d mail recipient recipients 1      +
94263      will write "1 recipient".                      +
94264      The command                                    +
94265      ngettext -d mail recipient recipients 5        +
94266      will write "2 to 10 recipients".              +
94267      The command                                    +
94268      ngettext -d mail recipient recipients 11       +
94269      will write "more than 10 recipients".          +
94270      The command                                    +
94271      ngettext -d mail Call Calls 1                  +
94272      will write "Call". Note that "Call" is not in the messages object. +
94273      The command                                    +
94274      ngettext -d mail Call Calls 0                  +
94275      will write "Calls".                            +
94276      The command                                    +
94277      ngettext -d mail Call Calls 10                 +
94278      will write "Calls".                            +
94279      The command                                    +
94280      ngettext -ed mail "%d attachment\n" "%d attachments\n" 1      +
94281      will write the same as                          +
94282      printf "1 (%d) attachment\n"                    +
94283      (i.e. "1 (%d) attachment" followed by a <newline> character). The output of ngettext can be +
94284      used as a format string for printf.             +
94285      The command                                    +
94286      printf "$(ngettext -ed mail "%d attachment\n" "%d attachments\n" 1)" 10 +
94287      will write the same as                          +
94288      printf "1 (%d) attachment\n" 10                 +
94289      (i.e. "1 (10) attachment" followed by a <newline> character).    +
94290      The command                                    +
94291      ngettext -e -d mail "\tsubject\n" "\tsubjects\n" 0          +
94292      will write the same as                          +
94293      printf "\tsubjects\n"                            +
94294      (i.e. a <tab> character, followed by "subjects" followed by a <newline> character). Note that +
94295      "\tsubject\n" is not in the messages object.     +
94296      The command                                    +

```



94297	<code>printf "%s\n" "\$(gettext -E -d mail "subject" "subjects" 0)"</code>	+
94298	will write the same as	+
94299	<code>printf "subjects\n"</code>	+
94300	(i.e. "subjects" followed by a <newline> character). Note that "subject" is not in the	+
94301	messages object.	+
94302	The command	+
94303	<code>gettext -s -d mail "recipient"</code>	+
94304	will write "1 recipient" followed by a <newline> character.	+
94305	The command	+
94306	<code>gettext -s -n -d mail "recipient"</code>	+
94307	will write "1 recipient" without a <newline> character.	+
94308	<b>RATIONALE</b>	+
94309	Historical implementations did not support the '\a' C-language escape sequence. This	+
94310	standard requires it to be supported for consistency with other utilities that support the table in	+
94311	XBD <a href="#">Chapter 5</a> (on page 101).	+
94312	Unlike other standard utilities, the behavior of <i>gettext</i> and <i>ngettext</i> is not undefined when	+
94313	<i>NLSPATH</i> overrides the system default path; see XBD <a href="#">Section 8.2</a> (on page 157). This is so that	+
94314	applications can use these utilities to obtain message strings from messages objects in other	+
94315	locations. However, it also means that they need to be implemented in such a way that they do	+
94316	not do anything that would result in undefined behavior when they need to write a diagnostic	+
94317	message. In particular, they should not use a string obtained from a message catalog or a	+
94318	messages object as a format string (or should only do so after checking that the string contains	+
94319	the correct conversions).	+
94320	<b>FUTURE DIRECTIONS</b>	+
94321	None.	+
94322	<b>SEE ALSO</b>	+
94323	<a href="#"><i>msgfmt</i></a> , <a href="#"><i>printf</i></a>	+
94324	XBD <a href="#">Chapter 7</a> (on page 115), <a href="#">Chapter 8</a> (on page 155), <a href="#">Section 12.2</a> (on page 201)	+
94325	XSH <a href="#"><i>gettext</i></a> , <a href="#"><i>iconv()</i></a> , <a href="#"><i>setlocale()</i></a>	+
94326	<b>CHANGE HISTORY</b>	+
94327	First released in Issue 8.	+
94328		+

101620 **NAME**

101621 msgfmt — create messages objects from portable messages object source files

101622 **SYNOPSIS**101623 msgfmt [-cfSv] [-D *dir*] [-o *outputfile*] *pathname...*101624 **DESCRIPTION**101625 The *msgfmt* utility shall create messages object files from portable messages object source files  
101626 (dot-po files).101627 A dot-po file contains messages to be output by system commands or by applications. The  
101628 messages in these files should be able to be translated to any language supported by the system.101629 The *msgfmt* utility shall interpret message strings for output as characters according to the  
101630 codeset specified in the dot-po file or, if not present, the current setting of the *LC\_CTYPE* locale  
101631 category.101632 **OPTIONS**101633 The *msgfmt* utility shall conform to XBD [Section 12.2](#) (on page 201).

101634 The following options shall be supported:

101635 **-c** If this option and **-v** are both specified, *msgfmt* shall detect and diagnose input file  
101636 abnormalities which might represent translation errors. The **msgid** and **msgstr**  
101637 strings shall be compared. It shall be considered abnormal if one string starts or  
101638 ends with a <newline> while the other does not. Also, if the flag **c-format** appears  
101639 in a "#, " comment for a **msgid** directive (see EXTENDED DESCRIPTION), it shall  
101640 be considered abnormal if the strings do not have the same number of '%'   
101641 conversion specifiers, or if corresponding conversion specifiers take different  
101642 argument types (see XSH *fprintf()*, on page 909). If an abnormality is detected, the  
101643 exit status shall be non-zero and a diagnostic message shall be output. Additional  
101644 checks beyond those described here may also be performed. These checks may  
101645 produce diagnostics or informational messages and need not affect the exit status.  
101646 If **-c** is specified without **-v** or **-v** is specified without **-c**, the behavior is  
101647 unspecified.

101648 **-D *dir*** Add *dir* to the list of directories to search for input files.101649 **-f** Use fuzzy entries in output. If this option is not specified, fuzzy entries shall not be  
101650 included in the output.

101651 **-o *outputfile***  
101652 Specify the name of an output file to be used instead of the default filename(s)  
101653 specified in EXTENDED DESCRIPTION. All **domain** *domainname* directives in the  
101654 dot-po file(s) shall be ignored.

101655 **-S** Append the suffix **.mo** to each generated messages object filename if it does not  
101656 have this suffix.101657 **-v** See **-c**.101658 **OPERANDS**

101659 The following operand shall be supported:

101660 *pathname* A pathname of a dot-po file.

101661	<b>STDIN</b>	
101662		Not used.
101663	<b>INPUT FILES</b>	
101664		The input files shall be text files in the format described in EXTENDED DESCRIPTION.
101665	<b>ENVIRONMENT VARIABLES</b>	
101666		The following environment variables shall affect the execution of <i>msgfmt</i> :
101667	<i>LANG</i>	Provide a default value for the internationalization variables that are unset or null.
101668		(See XBD <a href="#">Section 8.2</a> (on page 157) for the precedence of internationalization
101669		variables used to determine the values of locale categories.)
101670	XSI <i>LANGUAGE</i>	Determine the location of messages objects if <i>NLSPATH</i> is not set or the evaluation
101671		of <i>NLSPATH</i> did not lead to a suitable messages object being found.
101672	<i>LC_ALL</i>	If set to a non-empty string value, override the values of all the other
101673		internationalization variables.
101674	<i>LC_CTYPE</i>	Determine the locale for the interpretation of sequences of bytes of text data as
101675		characters (for example, single-byte as opposed to multi-byte characters in
101676		arguments and input files).
101677	<i>LC_MESSAGES</i>	
101678		Determine the locale name used to locate messages objects, and the locale that
101679		should be used to affect the format and contents of diagnostic messages written to
101680		standard error.
101681	XSI <i>NLSPATH</i>	Determine the location of messages objects and message catalogs.
101682	<b>ASYNCHRONOUS EVENTS</b>	
101683		Default.
101684	<b>STDOUT</b>	
101685		Not Used.
101686	<b>STDERR</b>	
101687		The standard error shall be used for diagnostic messages and may also be used for warning
101688		messages. If the <i>-c</i> and <i>-v</i> options are specified, additional unspecified informational messages
101689		may be written to standard error.
101690	<b>OUTPUT FILES</b>	
101691		The format of the created messages object files is unspecified.
101692	<b>EXTENDED DESCRIPTION</b>	
101693		The <i>msgfmt</i> utility shall accept portable messages object source files (dot-po files) in the
101694		following format.
101695		A dot-po file contains zero or more lines, with each non-blank line containing a comment, a
101696		statement, or a statement continuation. A comment has an unquoted <number-sign> ('#') as
101697		the first non-<blank> character and ends with the next <newline> character. A statement
101698		continuation is a double-quoted string on a line by itself, optionally preceded and/or followed
101699		by <blank> characters, and the string shall be concatenated with the string on the previous
101700		statement line. If a comment occurs between a statement and a statement continuation, the
101701		behavior is unspecified. All other comments, except for comments beginning with <number-
101702		sign><comma> ("#, "), and blank lines shall be ignored.
101703		The format of a statement is:
101704		<i>directive value</i>

The *directive* starts at the first non-`<blank>` character of the line and is separated from the *value* by one or more `<blank>` characters. The *value* consists of a double-quoted string optionally followed by `<blank>` characters. Zero or more statement continuation lines (see above) can follow the statement. The following directives shall be supported:

```
domain domainname
msgid message_identifier
msgid_plural untranslated_string_plural
msgstr message_string
msgstr[index] message_string
```

A dot-po file consists of zero or more sections. Each section specifies the messages to be processed in a domain. The first directive in each section shall be a **domain** directive (except for the first section which shall behave as if

```
domain "messages"
```

had been specified if the first directive is not a **domain** directive).

The behavior of the **domain** directive is affected by the options used. See OPTIONS for the behavior when the `-o` option is specified. If the `-o` option is not specified, all data obtained from the non-**domain** directives in a dot-po section shall be output to the messages object file named *domainname.mo* when the `-S` option is specified. When the `-S` option is not specified, it is implementation-defined whether *domainname* or *domainname.mo* is used.

If multiple **domain** directives specify the same *domainname*, the sections shall be processed as if there was only one section that starts with a **domain** *domainname* statement which contained the statements of the sections, in the same order, excluding all but the first **domain** *domainname* statement.

Within each section, there can be a header. A header is identified by having a **msgid** directive with the empty string (`"`) as the *message\_identifier* immediately followed by a statement containing a **msgstr** directive. The *message\_string* in this **msgstr** statement in a header shall be treated specially. If *message\_string* contains a specification of the form:

```
"nplurals=count; plural=expression"
```

then *count* indicates the number of plural forms for messages in that domain, and *expression* is a C-language expression that evaluates to an unsigned integer value which determines the **msgstr**[*index*] directive to be used. The value of *expression* is used as the index value. The variable *n* in *expression* is assigned the value of the *n* argument to the *ngettext()*, *ngettext\_l()*, *dngettext()*, *dngettext\_l()*, *dcngettext()*, and *dcngettext\_l()* functions or of the *n* operand of the *ngettext* utility before *expression* is evaluated. The application shall ensure that *expression* evaluates to a non-negative value less than *count* for all *n* that can be supplied by the aforementioned functions and utility.

If *message\_string* in the header contains a specification of the form:

```
"charset=codeset"
```

then *codeset* indicates the codeset to be used to encode the message strings in this section's domain (overriding `LC_CTYPE`). If the output string's codeset is different from the message string's codeset, codeset conversion from the message string's codeset to the output string's codeset shall be performed by the *gettext* family of functions and by the *gettext* and *ngettext* utilities. See XSH *gettext* and *gettext*. The output string's codeset shall be determined by the current or specified locale's codeset.

101749 **Note:** It is the responsibility of translators to ensure that the characters they enter into message strings  
 101750 in a dot-po file are encoded in the codeset specified in the header.

101751 If a header is present in a section, the application shall ensure that the header is provided by the  
 101752 first **msgid** directive in that section.

101753 After the header, if present, zero or more messages are identified by a **msgid** directive with a  
 101754 *message\_identifier* that is not an empty string. Each of these directives start a subsection that is  
 101755 used to get a translated message from the *gettext* family of functions and from the *gettext* and  
 101756 *ngettext* utilities. If the *message\_identifier* string is the string identified by the *gettext* family of  
 101757 functions *msgid* argument or by the *gettext* and *ngettext* utility *msgid* operand, this subsection  
 101758 specifies how that translation is to be processed.

101759 If there is only a singular form for the given *message\_identifier*, the application shall ensure that  
 101760 the statement containing the **msgid** directive is immediately followed by a **msgstr** directive.

101761 If there are plural forms for the given *message\_identifier* and the header for this section exists and  
 101762 contains an

101763 `"nplurals=count; plural=expression"`

101764 specification, the application shall ensure that the statement containing the **msgid** directive is  
 101765 immediately followed by a **msgid\_plural** directive and that each statement containing a  
 101766 **msgid\_plural** directive is followed by *count* statements containing **msgstr[index]** directives,  
 101767 starting with **msgstr[0]** and ending with **msgstr[count-1]** in monotonically increasing order. If a  
 101768 header for this section does not exist or does not contain an

101769 `"nplurals=count; plural=expression"`

101770 specification, the application shall ensure that no **msgid\_plural** or **msgstr[index]** directives are  
 101771 used in this section.

101772 For example, if the header's *message\_string* contains the specification:

101773 `"nplurals=2; plural= n == 1 ? 0 : 1"`

101774 there are two forms in the domain; **msgstr[0]** is used if *n* is equal to 1, otherwise **msgstr[1]** is  
 101775 used. For another example, if the header's *message\_string* contains:

101776 `"nplurals=3; plural= n == 1 ? 0 : n == 2 ? 1 : 2"`

101777 there are three forms in the domain; **msgstr[0]** is used if *n* is equal to 1, **msgstr[1]** is used if *n* is  
 101778 equal to 2, otherwise **msgstr[2]** is used.

101779 C-language escape sequences in strings shall be processed as specified for character string  
 101780 literals in the ISO C standard, except that *universal-character-name* escape sequences need not be  
 101781 supported.

101782 Comments in a dot-po file can be in one of the following formats:

101783 `#: reference`  
 101784 `#. utility-added-comments`  
 101785 `#, flag`  
 101786 `#translator-comments` (where *translator-comments* does not begin with '.', ': ' or ', ')

101787 A `#: reference` comment indicates the location(s) of the **msgid** string in the source files, in  
 101788 `pathname1:linenumber1 [pathname2:linenumber2 ... ]`

101789 format. They can be added, as might `"#. "` prefixed additional comments of unspecified format,  
 101790 by the *xgettext* utility. All comments that do not begin with `"#, "` are informative only and shall  
 101791 be silently ignored by the *msgfmt* utility. In `"#, "` comments the following values for *flag* can be

101792	specified:	
101793	<b>fuzzy</b>	This flag indicates that the <b>msgstr</b> string might not be a correct translation at this point in time. Only the translator can judge if the translation requires further modification or is acceptable as is. Once satisfied with the translation, the translator should remove this <b>fuzzy</b> flag. If this flag is specified, the <i>msgfmt</i> utility shall not generate the entry for the next following <b>msgid</b> in the output message catalog, unless the <b>-f</b> option is specified. If other flag comments are specified between <b>fuzzy</b> and the <b>msgid</b> , the behavior is unspecified.
101794		
101795		
101796		
101797		
101798		
101799		
101800	<b>c-format</b>	
101801	<b>no-c-format</b>	The <b>c-format</b> flag indicates that the next following <b>msgid</b> string contains a <i>printf()</i> format string. When the <b>c-format</b> flag is given and the <b>-c</b> and <b>-v</b> options are specified, the <i>msgfmt</i> utility shall perform additional tests to check the validity of the translation (see <b>OPTIONS</b> ); these additional tests may also be performed if neither <b>c-format</b> nor <b>no-c-format</b> is given. When the <b>no-c-format</b> flag is given for a string, no additional checks shall be performed for the string. When both the <b>c-format</b> and the <b>no-c-format</b> flags are given, the last flag specified takes precedence.
101802		
101803		
101804		
101805		
101806		
101807		
101808	<b>EXIT STATUS</b>	
101809	The following exit values shall be returned:	
101810	0	Successful completion.
101811	>0	An error occurred.
101812	<b>CONSEQUENCES OF ERRORS</b>	
101813	The <i>msgfmt</i> utility need not continue processing later <i>pathname</i> operands when an error condition that affects the exit status is detected. It is unspecified whether a messages object file is written when checks performed for the <b>-c</b> and <b>-v</b> options fail.	
101814		
101815		
101816	<b>APPLICATION USAGE</b>	
101817	The <i>xgettext</i> utility can be used to create template dot-po files from C-language source files.	
101818	Installing messages object files for the POSIX or C locale is not recommended, since they may be ignored for the sake of efficiency.	
101819		
101820	The first section for each domain in a dot-po file should include a header containing a	
101821	<code>"charset=codeset"</code>	
101822	specification. If this specification is omitted, message conversions in the <i>gettext</i> family of functions and in the <i>gettext</i> and <i>ngettext</i> utilities may fail.	
101823		
101824	The <b>msgid_plural</b> directive's <i>untranslated_string_plural</i> string comes from the <i>msgid_plural</i> arguments in calls to the <i>ngettext()</i> , <i>ngettext_l()</i> , <i>dngettext()</i> , <i>dngettext_l()</i> , <i>dcngettext()</i> , and <i>dcngettext_l()</i> functions when a prototype dot-po file is created by the <i>xgettext</i> utility. These strings (and the <i>msgid_plural</i> operands in calls to the <i>ngettext</i> utility) can provide context when a translator is modifying a template dot-po file into a dot-po file for a specific language. These functions and the <i>ngettext</i> utility do not try to match the <i>msgid_plural</i> arguments or operands with anything in a messages object file; they only match the <i>msgid</i> arguments and operands.	
101825		
101826		
101827		
101828		
101829		
101830		
101831	Unlike shell command language strings, double-quoted strings in dot-po files cannot contain a literal <newline> character.	
101832		

101833 **EXAMPLES**

101834 In this example, **module1.po** and **module2.po** are portable messages object source files.

```
101835 $ cat module1.po
101836 # default domain "messages"
101837 msgid ""
101838 msgstr "charset=utf-8"
101839 msgid "msg 1"
101840 msgstr "msg 1 translation"
101841 #
101842 domain "help_domain"
101843 msgid ""
101844 msgstr "charset=utf-8"
101845 msgid "help 2"
101846 msgstr "help 2 translation"
101847 #
101848 domain "error_domain"
101849 msgid ""
101850 msgstr "charset=utf-8"
101851 msgid "error 3"
101852 msgstr "error 3 translation"

101853 $ cat module2.po
101854 # default domain "messages"
101855 msgid ""
101856 msgstr "charset=utf-8"
101857 msgid "mesg 4"
101858 msgstr "mesg 4 translation"
101859 #
101860 domain "error_domain"
101861 msgid ""
101862 msgstr "charset=utf-8"
101863 #, c-format
101864 msgid "error 5 %s"
101865 msgstr "error 5 translation %s"
101866 #
101867 domain "window_domain"
101868 msgid ""
101869 msgstr "charset=utf-8"
101870 msgid "window 6"
101871 msgstr "window 6 translation"

101872 $ cat module3.po
101873 # default domain "messages"
101874 # header will be used for the whole output file in the third example
101875 msgid ""
101876 msgstr "charset=utf-8"
101877 msgid "info 0"
101878 msgstr "info 0 translation"

101879 $ cat opt_debug.po
101880 #
101881 domain "debug_domain"
101882 msgid "debug 8"
```

101883       **msgstr "debug 8 translation"**

101884       The following command will produce the output files **messages.mo**, **help\_domain.mo**, and  
101885       **error\_domain.mo**:

101886       **\$ msgfmt -S module1.po**

101887       The following command will produce the output files **messages.mo**, **help\_domain.mo**,  
101888       **error\_domain.mo**, and **window\_domain.mo**:

101889       **\$ msgfmt -S module1.po module2.po**

101890       The following command will produce the output file **hello.mo**:

101891       **\$ msgfmt -o hello.mo module3.po opt\_debug.po**

#### 101892 **RATIONALE**

101893       Some implementations are less strict about the format of dot-po files and simply treat all  
101894       occurrences of one or more white space characters as a separator. The format described in this  
101895       standard is accepted by all known implementations.

101896       In some implementations, duplicate **msgid** directives within a domain are ignored, and only an  
101897       entry for the first **msgid** directive and the following **msgid**, **msgid\_plural**, **msgstr**, or  
101898       **msgstr[index]** directives is created. However, some implementations consider duplicate **msgid**  
101899       directives within a domain to be an error and do not produce output at all. Consequently this  
101900       standard does not specify the behavior of *msgfmt* if duplicate **msgid** directives are encountered  
101901       within one domain.

#### 101902 **FUTURE DIRECTIONS**

101903       None.

#### 101904 **SEE ALSO**

101905       *gettext*, *xgettext*

101906       XSH *fprintf()*, *gettext*

#### 101907 **CHANGE HISTORY**

101908       First released in Issue 8.

101909



102281 **NAME**

102282       ngettext — retrieve text string from messages object

102283 **SYNOPSIS**102284       ngettext [-e|-E] [-d *textdomain*] [*textdomain*] *msgid msgid\_plural n*102285 **DESCRIPTION**102286       Refer to *gettext*.

102329 *PATH* Determine the search path used to locate the utility to be invoked. See XBD  
 102330 Chapter 8 (on page 155).

#### 102331 ASYNCHRONOUS EVENTS

102332 Default.

#### 102333 STDOUT

102334 Not used.

#### 102335 STDERR

102336 The standard error shall be used only for diagnostic messages.

#### 102337 OUTPUT FILES

102338 None.

#### 102339 EXTENDED DESCRIPTION

102340 None.

#### 102341 EXIT STATUS

102342 If *utility* is invoked, the exit status of *nice* shall be the exit status of *utility*; otherwise, the *nice*  
 102343 utility shall exit with one of the following values:

102344 1-125 An error occurred in the *nice* utility.

102345 126 The utility specified by *utility* was found but could not be invoked.

102346 127 The utility specified by *utility* could not be found.

#### 102347 CONSEQUENCES OF ERRORS

102348 Default.

#### 102349 APPLICATION USAGE

102350 The only guaranteed portable uses of this utility are:

102351 *nice utility*

102352 Run *utility* with the default higher or equal nice value.

102353 *nice -n <positive integer> utility*

102354 Run *utility* with a higher nice value.

102355 On some implementations they have no discernible effect on the invoked utility and on some  
 102356 others they are exactly equivalent.

102357 Historical systems have frequently supported the *<positive integer>* up to 20. Since there is no  
 102358 error penalty associated with guessing a number that is too high, users without access to the  
 102359 system conformance document (to see what limits are actually in place) could use the historical 1  
 102360 to 20 range or attempt to use very large numbers if the job should be truly low priority.

102361 The nice value of a process can be displayed using the command:

102362 `ps -o nice`

102363 The *command*, *env*, *nice*, *nohup*, *time*, *timeout*, and *xargs* utilities have been specified to use exit +  
 102364 code 127 if an error occurs so that applications can distinguish “failure to find a utility” from  
 102365 “invoked utility exited with an error indication”. The value 127 was chosen because it is not  
 102366 commonly used for other meanings; most utilities use small values for “normal error  
 102367 conditions” and the values above 128 can be confused with termination due to receipt of a  
 102368 signal. The value 126 was chosen in a similar manner to indicate that the utility could be found,  
 102369 but not invoked. Some scripts produce meaningful error messages differentiating the 126 and  
 102370 127 cases. The distinction between exit codes 126 and 127 is based on KornShell practice that  
 102371 uses 127 when all attempts to *exec* the utility fail with [ENOENT], and uses 126 when any

102772 *LC\_CTYPE* Determine the locale for the interpretation of sequences of bytes of text data as  
 102773 characters (for example, single-byte as opposed to multi-byte characters in  
 102774 arguments).

102775 *LC\_MESSAGES*  
 102776 Determine the locale that should be used to affect the format and contents of  
 102777 diagnostic messages written to standard error.

102778 XSI *NLSPATH* Determine the location of message catalogs for the processing of *LC\_MESSAGES*.

102779 *PATH* Determine the search path that is used to locate the utility to be invoked. See XBD  
 102780 Chapter 8 (on page 155).

#### 102781 ASYNCHRONOUS EVENTS

102782 The *nohup* utility shall take the standard action for all signals except that SIGHUP shall be  
 102783 ignored.

#### 102784 STDOUT

102785 If the standard output is not a terminal, the standard output of *nohup* shall be the standard  
 102786 output generated by the execution of the *utility* specified by the operands. Otherwise, nothing  
 102787 shall be written to the standard output.

#### 102788 STDERR

102789 If the standard output is a terminal, a message shall be written to the standard error, indicating  
 102790 the name of the file to which the output is being appended. The name of the file shall be either  
 102791 **nohup.out** or **\$HOME/nohup.out**.

#### 102792 OUTPUT FILES

102793 Output written by the named utility is appended to the file **nohup.out** (or **\$HOME/nohup.out**),  
 102794 if the conditions hold as described in the DESCRIPTION.

#### 102795 EXTENDED DESCRIPTION

102796 None.

#### 102797 EXIT STATUS

102798 The following exit values shall be returned:

102799 126 The utility specified by *utility* was found but could not be invoked.

102800 127 An error occurred in the *nohup* utility or the utility specified by *utility* could not be  
 102801 found.

102802 Otherwise, the exit status of *nohup* shall be that of the utility specified by the *utility* operand.

#### 102803 CONSEQUENCES OF ERRORS

102804 Default.

#### 102805 APPLICATION USAGE

102806 The *command*, *env*, *nice*, *nohup*, *time*, *timeout*, and *xargs* utilities have been specified to use exit +  
 102807 code 127 if an error occurs so that applications can distinguish “failure to find a utility” from  
 102808 “invoked utility exited with an error indication”. The value 127 was chosen because it is not  
 102809 commonly used for other meanings; most utilities use small values for “normal error  
 102810 conditions” and the values above 128 can be confused with termination due to receipt of a  
 102811 signal. The value 126 was chosen in a similar manner to indicate that the utility could be found,  
 102812 but not invoked. Some scripts produce meaningful error messages differentiating the 126 and  
 102813 127 cases. The distinction between exit codes 126 and 127 is based on KornShell practice that  
 102814 uses 127 when all attempts to *exec* the utility fail with [ENOENT], and uses 126 when any  
 102815 attempt to *exec* the utility fails for any other reason.

106666 **NAME**

106667 readlink — display the contents of a symbolic link

106668 **SYNOPSIS**106669 readlink [-n] *file*106670 **DESCRIPTION**

106671 If the *file* operand names a symbolic link, the *readlink* utility shall not follow the symbolic link  
 106672 when resolving *file* and shall write the contents of the symbolic link to standard output. If the **-n**  
 106673 option is not specified, the output to standard output shall be followed by a <newline>  
 106674 character.

106675 If *file* does not name a symbolic link, *readlink* shall write a diagnostic message to standard error  
 106676 and exit with non-zero status.

106677 **OPTIONS**106678 The *readlink* utility shall conform to XBD [Section 12.2](#) (on page 201).

106679 The following option shall be supported:

106680 **-n** Do not output a trailing <newline> character.106681 **OPERANDS**

106682 The following operand shall be supported:

106683 *file* A pathname of a symbolic link to be read.106684 **STDIN**

106685 Not used.

106686 **INPUT FILES**

106687 None.

106688 **ENVIRONMENT VARIABLES**106689 The following environment variables shall affect the execution of *readlink*:

106690 **LANG** Provide a default value for the internationalization variables that are unset or null.  
 106691 (See XBD [Section 8.2](#) (on page 157) for the precedence of internationalization  
 106692 variables used to determine the values of locale categories.)

106693 **LC\_ALL** If set to a non-empty string value, override the values of all the other  
 106694 internationalization variables.

106695 **LC\_CTYPE** Determine the locale for the interpretation of sequences of bytes of text data as  
 106696 characters (for example, single-byte as opposed to multi-byte characters in  
 106697 arguments and input files).

106698 **LC\_MESSAGES**

106699 Determine the locale that should be used to affect the format and contents of  
 106700 diagnostic messages written to standard error.

106701 **XSI** **NLSPATH** Determine the location of messages objects and message catalogs.106702 **ASYNCHRONOUS EVENTS**

106703 Default.

106704 **STDOUT**

106705 See DESCRIPTION.

**106706 STDERR**

106707 The standard error shall be used only for diagnostic messages.

**106708 OUTPUT FILES**

106709 None.

**106710 EXTENDED DESCRIPTION**

106711 None.

**106712 EXIT STATUS**

106713 The following exit values shall be returned:

106714 0 Successful completion.

106715 >0 An error occurred.

**106716 CONSEQUENCES OF ERRORS**

106717 Default.

**106718 APPLICATION USAGE**

106719 None.

**106720 EXAMPLES**

106721 None.

**106722 RATIONALE**

106723 The *readlink* utility was added because using *ls -l* to obtain the contents of a symbolic link is  
106724 difficult if the output includes more than one occurrence of the string " -> ".

106725 The *-f* option found in many implementations was not included, as the *realpath* utility provides  
106726 equivalent functionality with a choice of behaviors.

**106727 FUTURE DIRECTIONS**

106728 None.

**106729 SEE ALSO**

106730 *ln*, *ls*, *realpath*

106731 XBD [Chapter 8](#) (on page 155), [Section 12.2](#) (on page 201)

106732 XSH [readlink\(\)](#)

**106733 CHANGE HISTORY**

106734 First released in Issue 8.

106735

**NAME**

realpath — resolve a pathname

**SYNOPSIS**

realpath [-E|-e] file

**DESCRIPTION**

The *realpath* utility shall canonicalize the pathname specified by the *file* operand as follows:

If a call to the *realpath*() function with the specified pathname as its first argument would succeed, the canonicalized pathname shall be the pathname that would be returned by that *realpath*() call. Otherwise:

- If the **-e** option is specified, the canonicalization shall fail.
- If the **-E** option is specified, then if a call to the *realpath*() function with the specified pathname as its first argument would encounter an error condition other than [ENOENT], the canonicalization shall fail; if the call would encounter an [ENOENT] error, *realpath* shall expand all symbolic links that would be encountered in an attempt to resolve the specified pathname using the algorithm specified in XBD [Section 4.14](#) (on page 93), except that any trailing <slash> characters that are not also leading <slash> characters shall be ignored. If this expansion succeeds and the path prefix of the expanded pathname resolves to an existing directory, the canonicalized pathname shall be the expanded pathname. In all other cases, the canonicalization shall fail. If the expanded pathname is not empty, does not begin with a <slash>, and has exactly one pathname component, it shall be treated as if it had a path prefix of ". /".
- If no options are specified, *realpath* shall canonicalize the specified pathname in an unspecified manner such that the resulting absolute pathname does not contain any components that refer to files of type symbolic link and does not contain any components that are dot or dot-dot.

Upon successful canonicalization, *realpath* shall write the canonicalized pathname, followed by a <newline> character, to standard output.

If canonicalization fails, or the canonicalized pathname is empty, nothing shall be written to standard output, a diagnostic message shall be written to standard error, and *realpath* shall exit with non-zero status.

**OPTIONS**

The *realpath* utility shall conform to XBD [Section 12.2](#) (on page 201).

The following options shall be supported:

- E** Do not treat it as an error if attempting to resolve the last component of the canonicalized form of the *file* operand results in an [ENOENT] error condition.
- e** Treat it as an error if attempting to resolve the last component of the canonicalized form of the *file* operand results in an [ENOENT] error condition.

Specifying more than one of the mutually-exclusive options **-E** and **-e** shall not be considered an error. The last option specified shall determine the behavior of the utility.

**OPERANDS**

The following operand shall be supported:

*file* A pathname to be canonicalized.

106778	<b>STDIN</b>	
106779		Not used.
106780	<b>INPUT FILES</b>	
106781		None.
106782	<b>ENVIRONMENT VARIABLES</b>	
106783		The following environment variables shall affect the execution of <i>realpath</i> :
106784	<i>LANG</i>	Provide a default value for the internationalization variables that are unset or null.
106785		(See XBD <a href="#">Section 8.2</a> (on page 157) for the precedence of internationalization
106786		variables used to determine the values of locale categories.)
106787	<i>LC_ALL</i>	If set to a non-empty string value, override the values of all the other
106788		internationalization variables.
106789	<i>LC_CTYPE</i>	Determine the locale for the interpretation of sequences of bytes of text data as
106790		characters (for example, single-byte as opposed to multi-byte characters in
106791		arguments and input files).
106792	<i>LC_MESSAGES</i>	
106793		Determine the locale that should be used to affect the format and contents of
106794		diagnostic messages written to standard error.
106795	XSI <i>NLSPATH</i>	Determine the location of messages objects and message catalogs.
106796	<b>ASYNCHRONOUS EVENTS</b>	
106797		Default.
106798	<b>STDOUT</b>	
106799		See DESCRIPTION.
106800	<b>STDERR</b>	
106801		The standard error shall be used only for diagnostic messages.
106802	<b>OUTPUT FILES</b>	
106803		None.
106804	<b>EXTENDED DESCRIPTION</b>	
106805		None.
106806	<b>EXIT STATUS</b>	
106807		The following exit values shall be returned:
106808	0	Successful completion.
106809	>0	An error occurred.
106810	<b>CONSEQUENCES OF ERRORS</b>	
106811		Default.

**APPLICATION USAGE**

If neither the `-e` nor the `-E` option is specified, some implementations behave as if `-e` had been specified and others as if `-E` had been specified, but there are also implementations where the behavior differs from both of these. For example, the *mksh* shell has an internal implementation of *realpath* that canonicalizes */dir/regular\_file/..* to */dir*, whereas the *realpath()* function would return an [ENOTDIR] error in this case. Portable applications should always specify either `-e` or `-E`.

**EXAMPLES**

None.

**RATIONALE**

The *realpath* utility was added in preference to a `-f` option found in some implementations of the *readlink* utility because it allows the application to specify whether or not a missing final component is to be treated as an error.

The behavior with the `-E` option when *file* does not resolve (with symbolic links followed) to an existing file is not the same as simply calling *realpath()* with the path prefix of the *file* operand and writing the resulting pathname, a `<slash>`, and the last component of *file* to standard output. For example, if */tmp/nofile* does not exist, and *file* is *A/B* where *A* is an existing directory and *B* is a symbolic link to */tmp/nofile*, *realpath* with `-E` will output */tmp/nofile*, but if *B* is a symbolic link to */tmp/nofile/foo*, *realpath* with `-E` will treat this as an error. In both cases *realpath("A/B")* would fail with *errno* set to [ENOENT]. Even though *realpath("A")* would succeed, in neither case is anything ending */B* the result.

Trailing `<slash>` characters (that follow a non-`<slash>`) are handled differently with `-E` than with `-e`. With `-e` they are handled as for the *realpath()* function. With `-E` they are sometimes effectively ignored, and they are never included in the output. For example, if */tmp/nofile* does not exist and */tmp/regfile* is an existing regular file:

```
$ realpath -E /tmp/nofile/
/tmp/nofile
$ realpath -E /tmp/regfile/
realpath: /tmp/regfile/: Not a directory
```

Although the behavior of the *realpath* utility is specified by reference to the *realpath()* function, which is part of the XSI option, non-XSI implementations that do not support *realpath()* are nevertheless required to implement *realpath* in accordance with the requirements described in this standard for *realpath()*.

**FUTURE DIRECTIONS**

None.

**SEE ALSO**

*ln*, *ls*, *pwd*, *readlink*

XBD [Chapter 8](#) (on page 155), [Section 12.2](#) (on page 201)

XSH [Section 2.3](#) (on page 475), *realpath()*

**CHANGE HISTORY**

First released in Issue 8.



110894 127 The utility specified by *utility* could not be found.

## 110895 CONSEQUENCES OF ERRORS

110896 Default.

## 110897 APPLICATION USAGE

110898 The *command*, *env*, *nice*, *nohup*, *time*, *timeout*, and *xargs* utilities have been specified to use exit +  
 110899 code 127 if an error occurs so that applications can distinguish “failure to find a utility” from  
 110900 “invoked utility exited with an error indication”. The value 127 was chosen because it is not  
 110901 commonly used for other meanings; most utilities use small values for “normal error  
 110902 conditions” and the values above 128 can be confused with termination due to receipt of a  
 110903 signal. The value 126 was chosen in a similar manner to indicate that the utility could be found,  
 110904 but not invoked. Some scripts produce meaningful error messages differentiating the 126 and  
 110905 127 cases. The distinction between exit codes 126 and 127 is based on KornShell practice that  
 110906 uses 127 when all attempts to *exec* the utility fail with [ENOENT], and uses 126 when any  
 110907 attempt to *exec* the utility fails for any other reason.

## 110908 EXAMPLES

110909 It is frequently desirable to apply *time* to pipelines or lists of commands. This can be done by  
 110910 placing pipelines and command lists in a single file; this file can then be invoked as a utility, and  
 110911 the *time* applies to everything in the file.

110912 Alternatively, the following command can be used to apply *time* to a complex command:

```
110913 time sh -c 'complex-command-line'
```

## 110914 RATIONALE

110915 When the *time* utility was originally proposed to be included in the ISO POSIX-2:1993 standard,  
 110916 questions were raised about its suitability for inclusion on the grounds that it was not useful for  
 110917 conforming applications, specifically:

- 110918 • The underlying CPU definitions from the System Interfaces volume of POSIX.1-202x are  
 110919 vague, so the numeric output could not be compared accurately between systems or even  
 110920 between invocations.
- 110921 • The creation of portable benchmark programs was outside the scope this volume of  
 110922 POSIX.1-202x.

110923 However, *time* does fit in the scope of user portability. Human judgement can be applied to the  
 110924 analysis of the output, and it could be very useful in hands-on debugging of applications or in  
 110925 providing subjective measures of system performance. Hence it has been included in this  
 110926 volume of POSIX.1-202x.

110927 The default output format has been left unspecified because historical implementations differ  
 110928 greatly in their style of depicting this numeric output. The *-p* option was invented to provide  
 110929 scripts with a common means of obtaining this information.

110930 In the KornShell, *time* is a shell reserved word that can be used to time an entire pipeline, rather  
 110931 than just a simple command. The POSIX definition has been worded to allow this  
 110932 implementation. Consideration was given to invalidating this approach because of the historical  
 110933 model from the C shell and System V shell. However, since the System V *time* utility historically  
 110934 has not produced accurate results in pipeline timing (because the constituent processes are not  
 110935 all owned by the same parent process, as allowed by POSIX), it did not seem worthwhile to  
 110936 break historical KornShell usage.

110937 The term *utility* is used, rather than *command*, to highlight the fact that shell compound  
 110938 commands, pipelines, special built-ins, and so on, cannot be used directly. However, *utility*  
 110939 includes user application programs and shell scripts, not just the standard utilities.

110958 **NAME**110959        `timeout` — execute a utility with a time limit110960 **SYNOPSIS**110961        `timeout [-fp] [-k time] [-s signal_name] duration utility [argument...]`110962 **DESCRIPTION**

110963        The *timeout* utility shall execute the utility named by the *utility* operand, with arguments  
 110964        supplied as the *argument* operands (if any), in a child process. If the value of the *duration*  
 110965        operand is non-zero and the child process has not terminated after the specified time period,  
 110966        *timeout* shall send the signal specified by the *-s* option, or the SIGTERM signal if *-s* is not given.

110967        If the *-f* option is specified, the signal shall be sent only to the child process. Otherwise, it is  
 110968        implementation defined which one of the following methods is used to signal additional  
 110969        processes:

- 110970            • The *timeout* utility ensures it is a process group leader before creating the child process  
 110971            which executes the utility, in which case it shall send the signal to its process group.
- 110972            • The *timeout* utility arranges for any descendents of the child process that are orphaned to  
 110973            have their parent process changed to the *timeout* utility, in which case the signal shall be  
 110974            sent to the child process and all of its descendents.

110975        If the subsequent wait status of the child process shows that it was stopped by a signal, a  
 110976        SIGCONT signal shall also be sent in the same manner as the first signal; otherwise, a SIGCONT  
 110977        signal may be sent in the same manner.

110978        If the *-k* option is specified, and the child process created to execute the utility still has not  
 110979        terminated after the time period specified by the *time* option-argument has elapsed since the first  
 110980        signal was sent, *timeout* shall send a SIGKILL signal in the same manner as the first signal. If  
 110981        *timeout* receives a signal and propagates it to the child process (see ASYNCHRONOUS EVENTS  
 110982        below), this shall be treated as the first signal.

110983 **OPTIONS**110984        The *timeout* utility shall conform to XBD [Section 12.2](#) (on page 201).

110985        The following options shall be supported:

- 110986        *-f*            Only time out the utility itself, not its descendents.
- 110987        *-k time*        Send a SIGKILL signal if the child process created to execute the utility has not  
 110988        terminated after the time period specified by *time* has elapsed since the first signal  
 110989        was sent. The value of *time* shall be interpreted as specified for the *duration*  
 110990        operand (see OPERANDS below).
- 110991        *-p*            Always preserve (mimic) the wait status of the executed utility, even if the time  
 110992        limit was reached.
- 110993        *-s signal\_name*    Specify the signal to send when the time limit is reached, using one of the symbolic  
 110994        names defined in the `<signal.h>` header. Values of *signal\_name* shall be recognized  
 110995        in a case-independent fashion, without the SIG prefix. By default, SIGTERM shall  
 110996        be sent.  
 110997

110998 **OPERANDS**

110999        The following operands shall be supported:

- 111000        *duration*        The maximum amount of time to allow the utility to run, specified as a decimal  
 111001        number with an optional decimal fraction and an optional suffix, which can be:

111002	<b>s</b>	seconds	
111003	<b>m</b>	minutes	
111004	<b>h</b>	hours	
111005	<b>d</b>	days	
111006		If a decimal fraction is present, the application shall ensure that it is separated from	
111007		the units by a <period>. If no suffix is present, the value shall specify seconds.	
111008		If the value is zero, <i>timeout</i> shall not enforce a time limit.	
111009	<i>utility</i>	The name of a utility that is to be executed. If the <i>utility</i> operand names any of the	
111010		special built-in utilities in <a href="#">Section 2.14</a> (on page 2382), the results are undefined.	
111011	<i>argument</i>	Any string to be supplied as an argument when executing the utility named by the	
111012		<i>utility</i> operand.	
111013	<b>STDIN</b>		
111014		Not used.	
111015	<b>INPUT FILES</b>		
111016		None.	
111017	<b>ENVIRONMENT VARIABLES</b>		
111018		The following environment variables shall affect the execution of <i>timeout</i> :	
111019	<i>LANG</i>	Provide a default value for the internationalization variables that are unset or null.	
111020		(See XBD <a href="#">Section 8.2</a> (on page 157) for the precedence of internationalization	
111021		variables used to determine the values of locale categories.)	
111022	<i>LC_ALL</i>	If set to a non-empty string value, override the values of all the other	
111023		internationalization variables.	
111024	<i>LC_CTYPE</i>	Determine the locale for the interpretation of sequences of bytes of text data as	
111025		characters (for example, single-byte as opposed to multi-byte characters in	
111026		arguments and input files).	
111027	<i>LC_MESSAGES</i>		
111028		Determine the locale that should be used to affect the format and contents of	
111029		diagnostic messages written to standard error.	
111030	XSI <i>NLSPATH</i>	Determine the location of messages objects and message catalogs.	
111031	<i>PATH</i>	Determine the search path that is used to locate the utility to be executed. See XBD	
111032		<a href="#">Section 8.3</a> (on page 162).	
111033	<b>ASYNCHRONOUS EVENTS</b>		
111034		The default behavior specified in <a href="#">Section 1.4</a> (on page 2328) shall apply, except that:	
111035		• The <i>timeout</i> utility shall ignore SIGTTIN and SIGTTOU signals.	
111036		• The <i>timeout</i> utility may alter the disposition of SIGALRM if the inherited disposition was	
111037		for it to be ignored.	
111038		• If the signal specified with the <b>-s</b> option, or any signal whose default action is to terminate	
111039		the process, is delivered to the <i>timeout</i> utility, then unless the signal is SIGKILL or	
111040		SIGSTOP, the <i>timeout</i> utility shall immediately send the same signal to the process or	
111041		processes to which it would send a signal when the time limit is reached. If the delivered	
111042		signal is SIGALRM, <i>timeout</i> may behave as if the time limit had been reached instead of	
111043		sending SIGALRM.	

111044 • If the **-f** option is not specified, then if *timeout* sends a signal to its process group, it shall  
 111045 briefly change the disposition of that signal to ignored while it sends the signal, so that it  
 111046 does not receive the signal itself.

111047 With the single exception of the signal specified with the **-s** option, or SIGTERM if **-s** is not  
 111048 used, all signal dispositions inherited by the utility specified by the *utility* operand shall be the  
 111049 same as the disposition that *timeout* inherited.

#### 111050 **STDOUT**

111051 Not used.

#### 111052 **STDERR**

111053 The standard error shall be used only for diagnostic messages.

#### 111054 **OUTPUT FILES**

111055 None.

#### 111056 **EXTENDED DESCRIPTION**

111057 None.

#### 111058 **EXIT STATUS**

111059 If the **-p** option is not specified and the time limit was reached:

111060 • If the **-k** option was not specified or the utility terminated before the time period specified  
 111061 by the *time* option-argument elapsed since the first signal was sent, the exit status shall be  
 111062 124.

111063 • If the **-k** option was specified and the SIGKILL signal was sent, it is unspecified whether  
 111064 the exit status is 124 or the behavior is as if the **-p** option was specified.

111065 Otherwise, if the executed utility terminated by exiting, the exit status of *timeout* shall be that of  
 111066 the utility; if the utility was terminated by a signal, *timeout* shall terminate itself with the same  
 111067 signal while ensuring that a core image is not created.

111068 If an error occurs, the following exit values shall be returned:

111069 125 An error other than the two described below occurred.

111070 126 The utility specified by *utility* was found but could not be executed.

111071 127 The utility specified by *utility* could not be found.

#### 111072 **CONSEQUENCES OF ERRORS**

111073 Default.

#### 111074 **APPLICATION USAGE**

111075 Unlike the *kill* utility, the **-s** option of *timeout* is not required to accept the symbolic name 0 to  
 111076 represent signal value zero.

111077 When the value of *duration* is zero, *timeout* does not time out the utility, but it does still perform  
 111078 signal propagation (including to descendents of the utility if **-f** is not specified).

111079 Regardless of locale, the <period> character (the decimal-point character of the POSIX locale) is  
 111080 the decimal-point character recognized in the *duration* operand and the *time* option-argument.

111081 The *command*, *env*, *nice*, *nohup*, *time*, *timeout*, and *xargs* utilities have been specified to use exit  
 111082 code 127 if a utility to be invoked cannot be found, so that applications can distinguish “failure  
 111083 to find a utility” from “invoked utility exited with an error indication”. The value 127 was  
 111084 chosen because it is not commonly used for other meanings; most utilities use small values for  
 111085 “normal error conditions” and the values above 128 can be confused with termination due to  
 111086 receipt of a signal. The value 126 was chosen in a similar manner to indicate that the utility could

be found, but not invoked. Some scripts produce meaningful error messages differentiating the 126 and 127 cases. The distinction between exit codes 126 and 127 is based on KornShell practice that uses 127 when all attempts to *exec* the utility fail with [ENOENT], and uses 126 when any attempt to *exec* the utility fails for any other reason. The *timeout* utility extends these special exit codes to 125 and 124, with the meanings described in EXIT STATUS. A *timeout* exit status below 124 can only result from passing through the exit status of the executed utility.

#### EXAMPLES

None.

#### RATIONALE

Some *timeout* implementations make themselves a process group leader (when *-f* is not used) in order to be able to send signals to descendents of the child process. However, using this method means that any descendents which change their process group do not receive the signal. To ensure all descendents receive the signal, some implementations instead make use of a feature whereby descendents that are orphaned have their parent process changed to the *timeout* utility—that is, *timeout* becomes their “reaper”—together with the ability of a reaper to send a signal to all of its descendents.

Some historical *timeout* implementations exited with status *128+signal\_number* when the child process was terminated by a signal before the time limit was reached (or when *-p* was used). This is reasonable when *timeout* is invoked from a shell which sets *\$?* to *128+signal\_number*, but not all shells do that. In particular, the KornShell sets *\$?* to *256+signal\_number* and so an exit status of *128+signal\_number* from *timeout* would be misleading. In order to avoid any possible ambiguity, this standard requires that *timeout* mimics the wait status of the child process by terminating itself with the same signal. When it does this it needs to ensure that it does not create a core image, otherwise it could overwrite one created by the invoked utility.

The *timeout* utility ignores SIGTTIN and SIGTTOU so that if the utility it executes reads from or writes to the controlling terminal and this generates a SIGTTIN or SIGTTOU for the process group, *timeout* will not be stopped by the signal and can still time out the utility.

Some historical *timeout* implementations always set the disposition for SIGTTIN and SIGTTOU in the child process to default, even if these signals were inherited as ignored. This could result in processes being stopped unexpectedly. Likewise, they did not ensure that for signals they caught, the disposition inherited by the executed utility was the same as the disposition that was inherited by *timeout*. This meant that, for example, if *timeout* was used in a script that was run with *nohup*, the utility executed by *timeout* would unexpectedly not be protected from SIGHUP. This standard requires that all signal dispositions inherited by the utility specified by the *utility* operand are the same as the disposition that *timeout* inherited, with the single exception of the signal that *timeout* sends when the time limit is reached, which needs to be inherited as default in order for the timeout to take effect (without resorting to SIGKILL if *-k* is specified).

Some historical *timeout* implementations only propagated a subset of the signals whose default action is to terminate the process to the child process if one was delivered to the *timeout* utility. Propagating these signals is beneficial, as otherwise termination of the *timeout* utility by a signal results in the utility it executed being left running indefinitely (unless it also received the signal, for example a terminal-generated SIGINT). There is no reason to select a subset of these signals to be propagated, therefore this standard requires them all to be propagated (except SIGKILL, which cannot). In the event that a user wants to prevent the utility being timed out, sending *timeout* a SIGKILL can be used for this purpose.

111132 **FUTURE DIRECTIONS**

111133 None.

111134 **SEE ALSO**111135 *kill*111136 XBD [Chapter 8](#) (on page 155), [Section 12.2](#) (on page 201), [<signal.h>](#) (on page 329)111137 **CHANGE HISTORY**

111138 First released in Issue 8.

111139

116531 **EXTENDED DESCRIPTION**

116532 None.

116533 **EXIT STATUS**

116534 The following exit values shall be returned:

- 116535 0 All invocations of *utility* returned exit status zero.
- 116536 1-125 A command line meeting the specified requirements could not be assembled, one or  
116537 more of the invocations of *utility* returned a non-zero exit status, or some other error  
116538 occurred.
- 116539 126 The utility specified by *utility* was found but could not be invoked.
- 116540 127 The utility specified by *utility* could not be found.

116541 **CONSEQUENCES OF ERRORS**

116542 If a command line meeting the specified requirements cannot be assembled, the utility cannot be  
116543 invoked, an invocation of the utility is terminated by a signal, or an invocation of the utility exits  
116544 with exit status 255, the *xargs* utility shall write a diagnostic message and exit without  
116545 processing any remaining input.

116546 **APPLICATION USAGE**

116547 The 255 exit status allows a utility being used by *xargs* to tell *xargs* to terminate if it knows no  
116548 further invocations using the current data stream will succeed. Thus, *utility* should explicitly *exit*  
116549 with an appropriate value to avoid accidentally returning with 255.

116550 Note that since input is parsed as lines, <blank> characters separate arguments, and  
116551 <backslash>, <apostrophe>, and double-quote characters are used for quoting, if *xargs* is used to  
116552 bundle the output of commands like *find dir -print* or *ls* into commands to be executed,  
116553 unexpected results are likely if any filenames contain <blank>, <newline>, or quoting characters.  
116554 This can be solved by using *find* to call a script that converts each file found into a quoted string  
116555 that is then piped to *xargs*, but in most cases it is preferable just to have *find* do the argument  
116556 aggregation itself by using *-exec* with a '+' terminator instead of ';' '. Note that the quoting  
116557 rules used by *xargs* are not the same as in the shell. They were not made consistent here because  
116558 existing applications depend on the current rules. An easy (but inefficient) method that can be  
116559 used to transform input consisting of one argument per line into a quoted form that *xargs*  
116560 interprets correctly is to precede each non-<newline> character with a <backslash>. More  
116561 efficient alternatives are shown in Example 2 and Example 5 below.

116562 On implementations with a large value for {ARG\_MAX}, *xargs* may produce command lines  
116563 longer than {LINE\_MAX}. For invocation of utilities, this is not a problem. If *xargs* is being used  
116564 to create a text file, users should explicitly set the maximum command line length with the *-s*  
116565 option.

116566 The *command*, *env*, *nice*, *nohup*, *time*, *timeout*, and *xargs* utilities have been specified to use exit +  
116567 code 127 if an error occurs so that applications can distinguish "failure to find a utility" from  
116568 "invoked utility exited with an error indication". The value 127 was chosen because it is not  
116569 commonly used for other meanings; most utilities use small values for "normal error  
116570 conditions" and the values above 128 can be confused with termination due to receipt of a  
116571 signal. The value 126 was chosen in a similar manner to indicate that the utility could be found,  
116572 but not invoked. Some scripts produce meaningful error messages differentiating the 126 and  
116573 127 cases. The distinction between exit codes 126 and 127 is based on KornShell practice that  
116574 uses 127 when all attempts to *exec* the utility fail with [ENOENT], and uses 126 when any  
116575 attempt to *exec* the utility fails for any other reason.

## 116690 NAME

116691 *xgettext* — extract *gettext* call strings from C-language source files (DEVELOPMENT)

## 116692 SYNOPSIS

```

116693 CD      xgettext [-j] [-n] [-d default-domain] [-K keyword-spec]...
116694          [-p pathname] file...
116695
116696      xgettext -a [-n] [-d default-domain] [-p pathname]
116697          [-x exclude-file] file...

```

## 116698 DESCRIPTION

116699 The *xgettext* utility shall automate the creation of portable messages object source files (dot-po  
 116700 files). A dot-po file shall contain copies of string literals that are found in C-language source  
 116701 code in files specified by *file* operands. The dot-po file can be used as input to the *msgfmt* utility,  
 116702 to produce a messages object file that can be used by applications.

116703 The *xgettext* utility shall write *msgid* argument strings that are passed as string literals in  
 116704 *gettext()*, *gettext\_l()*, *ngettext()*, and *ngettext\_l()* calls in C-language source code to the default  
 116705 output file; this file shall be named **messages.po** unless it is changed by the **-d** option. The  
 116706 *xgettext* utility shall also write *msgid* argument strings that are passed as string literals in  
 116707 *dcgettext()*, *dcgettext\_l()*, *dcngettext()*, *dcngettext\_l()*, *dgettext()*, *dgettext\_l()*, *dngettext()*, and  
 116708 *dngettext\_l()* calls either to the default output file or to the output file *domainname.po* where  
 116709 *domainname* is the first parameter to the call; it is implementation-defined which of those output  
 116710 files is used. A **msgid** directive shall precede each *msgid* argument string. For the functions that  
 116711 have a *msgid\_plural* argument, a **msgid\_plural** directive followed by that argument string shall  
 116712 also be written directly after the corresponding **msgid** directive. A **msgstr** directive or  
 116713 **msgstr[index]** directives with an empty string shall be written after the corresponding **msgid** or  
 116714 **msgid\_plural** directive, respectively. The function names that *xgettext* searches for can be  
 116715 changed using the **-K** option.

116716 The first directive in each created dot-po file shall be a **domain** directive giving the associated  
 116717 domain name, except that this directive is optional in the default output file.

116718 If the **-p pathname** option is specified, *xgettext* shall create the dot-po files in the *pathname*  
 116719 directory. Otherwise, the dot-po files shall be created in the current working directory.

116720 The **msgid** values shall be in the same order that the strings are extracted from each *file* and  
 116721 subsections with duplicate **msgid** values shall be written to the dot-po files as comment lines.

## 116722 OPTIONS

116723 The *xgettext* utility shall conform to XBD [Section 12.2](#) (on page 201).

116724 The following options shall be supported:

116725 **-a** Extract all strings, not just those found in calls to *gettext* family functions. Only one  
 116726 dot-po file shall be created.

116727 **-d default-domain**  
 116728 Name the default output file *default-domain.po* instead of **messages.po**.

116729 **-j** Join messages from C-language source files with existing dot-po files. For each  
 116730 dot-po file that *xgettext* writes messages to, if the file does not exist, it shall be  
 116731 created. New messages shall be appended but any subsections with duplicate  
 116732 **msgid** values except the first (including **msgid** values found in an existing dot-po  
 116733 file) shall either be commented out or omitted in the resulting dot-po file; if  
 116734 omitted, a warning message may be written to standard error. Domain directives  
 116735 in the existing dot-po files shall be ignored; the assumption is that all previous



116736 **msgid** values belong to the same domain. The behavior is unspecified if an existing  
 116737 dot-po file was not created by *xgettext* or has been modified by another application.

116738 **-K keyword-spec**

116739 Specify an additional keyword to be looked for:

- 116740 • If *keyword-spec* is an empty string, this shall disable the use of default
- 116741 keywords for the *gettext* family of functions.
- 116742 • If *keyword-spec* is a C identifier, *xgettext* shall look for strings in the first
- 116743 argument of each call to the function or macro *keyword-spec*.
- 116744 • If *keyword-spec* is of the form *id:argnum* then *xgettext* shall treat the *argnum*-th
- 116745 argument of a call to the function or macro *id* as the *msgid* argument, where
- 116746 *argnum* 1 is the first argument.
- 116747 • If *keyword-spec* is of the form *id:argnum1,argnum2* then *xgettext* shall treat
- 116748 strings in the *argnum1*-th argument and in the *argnum2*-th argument of a call
- 116749 to the function or macro *id* as the *msgid* and *msgid\_plural* arguments,
- 116750 respectively.

116751 For all mentioned forms, the application shall ensure that if *argnum2* is given, it is  
 116752 not equal to *argnum1*. All numeric values shall be converted as specified in item 6  
 116753 in XBD [Section 12.1](#) (on page 199).

116754 **-n**

116755 Add comment lines to the output file indicating pathnames and line numbers in  
 116756 the source files where each extracted string is encountered. These lines shall  
 appear before each **msgid** directive. Such comments should have the format:

116757 `#: pathname1:linenumber1 [pathname2:linenumber2...]`

116758 **-p pathname**

116759 Create output files in the directory specified by *pathname* instead of in the current  
 116760 working directory.

116761 **-x exclude-file**

116762 Specify a file containing strings that shall not be extracted from the input files. The  
 116763 format of *exclude-file* is identical to that of a dot-po file. However, only statements  
 116764 containing **msgid** directives in *exclude-file* shall be used. All other statements shall  
 116765 be ignored.

## 116766 OPERANDS

116767 The following operand shall be supported:

116768 *file* A pathname of an input file containing C-language source code. If '-' is specified  
 116769 for an instance of *file*, the standard input shall be used.

## 116770 STDIN

116771 The standard input shall not be used unless a *file* operand is specified as '- '.

## 116772 INPUT FILES

116773 The input files specified as *file* operands shall be C-language source files. The input file specified  
 116774 as the option-argument for the **-x** option shall be a dot-po file in the format specified as input for  
 116775 the *msgfmt* utility.

## 116776 ENVIRONMENT VARIABLES

116777 The following environment variables shall affect the execution of *xgettext*:

116778	<i>LANG</i>	Provide a default value for the internationalization variables that are unset or null. (See XBD <a href="#">Section 8.2</a> (on page 157) for the precedence of internationalization variables used to determine the values of locale categories.)
116779		
116780		
116781	XSI <i>LANGUAGE</i>	Determine the location of messages objects if <i>NLSPATH</i> is not set or the evaluation of <i>NLSPATH</i> did not lead to a suitable messages object being found.
116782		
116783	<i>LC_ALL</i>	If set to a non-empty string value, override the values of all the other internationalization variables.
116784		
116785	<i>LC_CTYPE</i>	Determine the locale for the interpretation of sequences of bytes of text data as characters (for example, single-byte as opposed to multi-byte characters in arguments and input files).
116786		
116787		
116788	<i>LC_MESSAGES</i>	
116789		Determine the locale name used to locate messages objects, and the locale that should be used to affect the format and contents of diagnostic messages written to standard error.
116790		
116791		
116792	XSI <i>NLSPATH</i>	Determine the location of messages objects and message catalogs.
116793	<b>ASYNCHRONOUS EVENTS</b>	
116794		Default.
116795	<b>STDOUT</b>	
116796		The standard output shall not be used.
116797	<b>STDERR</b>	
116798		The standard error shall be used for diagnostic messages and may be used for warning messages.
116799		
116800	<b>OUTPUT FILES</b>	
116801		The output files shall be dot-po files in the format specified as input for the <i>msgfmt</i> utility. It is unspecified whether each output file includes a header ( <b>msgid ""</b> ) before the content derived from the input C-language source files.
116802		
116803		
116804	<b>EXTENDED DESCRIPTION</b>	
116805		None.
116806	<b>EXIT STATUS</b>	
116807		The following exit values shall be returned:
116808	0	Successful completion.
116809	>0	An error occurred.
116810	<b>CONSEQUENCES OF ERRORS</b>	
116811		Default.
116812	<b>APPLICATION USAGE</b>	
116813		Implementations differ as to whether they write all output to the default output file or split the output into separate per-domain files. Portable applications can either ensure that each C-language source file contains calls to <i>gettext</i> family functions for only a single domain, or force all output to be to the default output file by using the <b>-K</b> option to override the default keywords.
116814		
116815		
116816		
116817		
116818		Some implementations of <i>xgettext</i> are not able to extract cast strings (unless <b>-a</b> is used), for example casts of literal strings to ( <b>const char *</b> ). Use of a cast is unnecessary anyway, since the prototypes in <b>&lt;libintl.h&gt;</b> already specify this type.
116819		
116820		

The *xgettext* utility is not required to handle C preprocessor directives. Therefore if, for example, calls to *gettext* family functions are wrapped by macros, they might not be found unless the **-K** option is used to tell *xgettext* to look for the macro calls.

## EXAMPLES

### Example 1

The following example shows how **-K** can be used to force all output to be to the default output file:

```
xgettext -K "" -K gettext:1 -K dgettext:2 -K dcgettext:2 \
-K ngettext:1,2 -K dngettext:2,3 -K dcngettext:2,3 source.c
```

By overriding the default keywords using the **-K** option as above, the *xgettext* utility is directed to ignore the *domainname* arguments to the *dgettext()*, *dcgettext()*, *dngettext()*, and *dcngettext()* functions. Thus, the utility treats the functions as their respective equivalent without the *d* prefix, ignoring the *domainname* argument and writing generated output to the default output file, **messages.po**. Additional **-K** options would be needed for the variants of the functions with an *\_l* suffix if they are used.

### Example 2

If the source uses a macro definition such as:

```
#define i18n gettext
```

the use of:

```
xgettext -K i18n:1 source.c
```

will pick up **msgid** values from a line such as:

```
fprintf(stdout, i18n("The value is %s"), value1);
```

## RATIONALE

The **-K** option is based on the **-k** option of GNU *xgettext*; the only difference is that GNU's **-k** takes an optional option-argument whereas **-K** in this standard has a mandatory option-argument in order to comply with the syntax guidelines.

The standard developers considered including functionality equivalent to the **-c**, **-m**, and **-M** options in existing implementations. However, those letters could not be used as the syntax differed between implementations. The usual solution of adding an uppercase equivalent of lowercase options with the standard syntax instead was not possible, for obvious reasons for **-m** and **-M**, and as **-C** was already in use for another purpose in one implementation.

The **-s** option is not included as it has been deprecated in at least one implementation because it has been found to deprive translators of valuable context.

## FUTURE DIRECTIONS

A future version of this standard may change the description of the **-n** option to use "shall" instead of "should".

## SEE ALSO

*gettext*, *msgfmt*

XBD [Chapter 8](#) (on page 155), [Section 12.2](#) (on page 201)

XSH *gettext*

116861 **CHANGE HISTORY**  
116862 First released in Issue 8.  
116863

129690	POSIX_DEVICE_SPECIFIC_R: Thread-Safe General Terminal	
129691	<i>ttyname_r()</i>	
129692	POSIX_DYNAMIC_LINKING: Dynamic Linking	
129693	<i>dladdr(), dlclose(), dlderror(), dlopen(), dlsym()</i>	
129694	POSIX_FD_MGMT: File Descriptor Management	
129695	<i>dup(), dup2(), dup3(), fcntl(), fgetpos(), fseek(), fseeko(), fsetpos(), ftell(), ftello(), ftruncate(),</i>	
129696	<i>lseek(), rewind()</i>	
129697	POSIX_FIFO: FIFO	
129698	<i>mkfifo()</i>	
129699	POSIX_FIFO_FD: FIFO File Descriptor Routines	
129700	<i>mkfifoat(), mknodat()</i>	
129701	POSIX_FILE_ATTRIBUTES: File Attributes	
129702	<i>chmod(), chown(), fchmod(), fchown(), umask()</i>	
129703	POSIX_FILE_ATTRIBUTES_FD: File Attributes File Descriptor Routines	
129704	<i>fchmodat(), fchownat()</i>	
129705	POSIX_FILE_LOCKING: Thread-Safe Stdio Locking	
129706	<i>flockfile(), ftrylockfile(), funlockfile(), getc_unlocked(), getchar_unlocked(), putc_unlocked(),</i>	
129707	<i>putchar_unlocked()</i>	
129708	POSIX_FILE_SYSTEM: File System	
129709	<i>access(), chdir(), closedir(), creat(), fchdir(), fpathconf(), fstat(), fstatvfs(), getcwd(), link(),</i>	
129710	<i>mkdir(), mkostemp(), mkstemp(), opendir(), pathconf(), posix_getdents(), readdir(), remove(),</i>	
129711	<i>rename(), rewinddir(), rmdir(), stat(), statvfs(), tmpfile(), tmpnam(), truncate(), unlink()</i>	
129712	POSIX_FILE_SYSTEM_EXT: File System Extensions	
129713	<i>alphasort(), dirfd(), getdelim(), getline(), mkdtemp(), scandir()</i>	
129714	POSIX_FILE_SYSTEM_FD: File System File Descriptor Routines	
129715	<i>faccessat(), fdopendir(), fstatat(), linkat(), mkdirat(), openat(), renameat(), unlinkat(),</i>	
129716	<i>utimensat()</i>	
129717	POSIX_FILE_SYSTEM_GLOB: File System Glob Expansion	
129718	<i>glob(), globfree()</i>	
129719	POSIX_FILE_SYSTEM_R: Thread-Safe File System	
129720	<i>readdir_r()</i>	
129721	POSIX_I18N: Internationalization	
129722	<i>bind_textdomain_codeset(), bindtextdomain(), catclose(), catgets(), catopen(), dcgettext(),</i>	+
129723	<i>dcgettext_l(), dcngettext(), dcngettext_l(), dgettext(), dgettext_l(), dngettext(), dngettext_l(),</i>	+
129724	<i>gettext(), gettext_l(), iconv(), iconv_close(), iconv_open(), ngettext(), ngettext_l(),</i>	
129725	<i>nl_langinfo(), textdomain()</i>	
129726	POSIX_JOB_CONTROL: Job Control	
129727	<i>setpgid(), tcgetpgrp(), tcsetpgrp(), tcgetsid()</i>	
129728	POSIX_MAPPED_FILES: Memory Mapped Files	
129729	<i>mmap(), munmap()</i>	
129730	POSIX_MEMORY_PROTECTION: Memory Protection	
129731	<i>mprotect()</i>	

129820	POSIX_WIDE_CHAR_DEVICE_IO: Device Input and Output	
129821	<i>fgetc()</i> , <i>fgetws()</i> , <i>fputc()</i> , <i>fputws()</i> , <i>fwide()</i> , <i>fwprintf()</i> , <i>fwscanf()</i> , <i>getc()</i> , <i>getwchar()</i> ,	
129822	<i>putc()</i> , <i>putwchar()</i> , <i>ungetc()</i> , <i>vfwprintf()</i> , <i>vfwscanf()</i> , <i>vwprintf()</i> , <i>vwscanf()</i> , <i>wprintf()</i> ,	
129823	<i>wscanf()</i>	
129824	XSI_C_LANG_SUPPORT: XSI General C Library	
129825	<i>a64l()</i> , <i>daylight</i> , <i>drand48()</i> , <i>erand48()</i> , <i>ffs()</i> , <i>ffsl()</i> , <i>ffsll()</i> , <i>getdate()</i> , <i>hcreate()</i> , <i>hdestroy()</i> ,	
129826	<i>hsearch()</i> , <i>initstate()</i> , <i>insque()</i> , <i>jrand48()</i> , <i>l64a()</i> , <i>lcong48()</i> , <i>lfind()</i> , <i>lrand48()</i> , <i>lsearch()</i> ,	
129827	<i>memccpy()</i> , <i>mrnd48()</i> , <i>nrnd48()</i> , <i>random()</i> , <i>remque()</i> , <i>seed48()</i> , <i>setstate()</i> , <i>signgam</i> ,	
129828	<i>srand48()</i> , <i>srandom()</i> , <i>strptime()</i> , <i>swab()</i> , <i>tdelete()</i> , <i>tfind()</i> , <i>timezone</i> , <i>tsearch()</i> , <i>twalk()</i>	
129829	XSI_DBM: XSI Database Management	
129830	<i>dbm_clearerr()</i> , <i>dbm_close()</i> , <i>dbm_delete()</i> , <i>dbm_error()</i> , <i>dbm_fetch()</i> , <i>dbm_firstkey()</i> ,	
129831	<i>dbm_nextkey()</i> , <i>dbm_open()</i> , <i>dbm_store()</i>	
129832	XSI_DEVICE_IO: XSI Device Input and Output	
129833	<i>fntmsg()</i> , <i>readv()</i> , <i>writv()</i>	
129834	XSI_DEVICE_SPECIFIC: XSI General Terminal	
129835	<i>grantpt()</i> , <i>posix_openpt()</i> , <i>ptsname()</i> , <i>unlockpt()</i>	
129836	XSI_FILE_SYSTEM: XSI File System	
129837	<i>basename()</i> , <i>dirname()</i> , <i>lockf()</i> , <i>mknod()</i> , <i>nftw()</i> , <i>realpath()</i> , <i>seekdir()</i> , <i>sync()</i> , <i>telldir()</i>	
129838	XSI_GENERAL_TERMINAL_R: XSI Thread-Safe General Terminal	
129839	<i>ptsname_r()</i>	
129840	XSI_IPC: XSI Interprocess Communication	
129841	<i>flok()</i> , <i>msgctl()</i> , <i>msgget()</i> , <i>msgrcv()</i> , <i>msgsnd()</i> , <i>semctl()</i> , <i>semget()</i> , <i>semop()</i> , <i>shmat()</i> , <i>shmctl()</i> ,	
129842	<i>shmdt()</i> , <i>shmget()</i>	
129843	XSI_MATH: XSI Maths Library	
129844	<i>j0()</i> , <i>j1()</i> , <i>jn()</i> , <i>y0()</i> , <i>y1()</i> , <i>yn()</i>	
129845	XSI_MULTI_PROCESS: XSI Multiple Process	
129846	<i>getpriority()</i> , <i>getrlimit()</i> , <i>getrusage()</i> , <i>nice()</i> , <i>setpriority()</i> , <i>setrlimit()</i>	
129847	XSI_SIGNALS: XSI Signal	
129848	<i>killpg()</i> , <i>sigaltstack()</i>	
129849	XSI_SINGLE_PROCESS: XSI Single Process	
129850	<i>gethostid()</i> , <i>putenv()</i>	
129851	XSI_SYSTEM_DATABASE: XSI System Database	
129852	<i>endpwent()</i> , <i>getpwent()</i> , <i>setpwent()</i>	
129853	XSI_SYSTEM_LOGGING: XSI System Logging	
129854	<i>closelog()</i> , <i>openlog()</i> , <i>setlogmask()</i> , <i>syslog()</i>	
129855	XSI_THREADS_EXT: XSI Threads Extensions	
129856	<i>pthread_attr_getstack()</i> , <i>pthread_attr_setstack()</i>	
129857	XSI_USER_GROUPS: XSI User and Group	
129858	<i>endgrent()</i> , <i>endutxent()</i> , <i>getgrent()</i> , <i>getresgid()</i> , <i>getresuid()</i> , <i>getutxent()</i> , <i>getutxid()</i> ,	+
129859	<i>getutxline()</i> , <i>pututxline()</i> , <i>setgrent()</i> , <i>setregid()</i> , <i>setresgid()</i> , <i>setresuid()</i> , <i>setreuid()</i> , <i>setutxent()</i>	+
129860	XSI_WIDE_CHAR: XSI Wide-Character Library	
129861	<i>wcswidth()</i> , <i>wcwidth()</i>	