



Optimizing MemorySanitizer

Ideas, Victories, and Dead Ends

Gui Andrade



What is MemorySanitizer?

MSAN is a compiler pass, it protects you from using uninitialized memory

Uninitialized memory causes nasty bugs!

MSAN instruments your code with shadow memory that tags init status

How does Google use MSAN?

Chromium testing

User-facing memory vulns are a disaster!



Google3 testing

Business-facing memory vulns are a disaster!

MemorySanitizer

OSS-Fuzz

Adaptive checking for bugs in open-source software

oss-fuzz

Why optimize it? Space



Large binaries get inflated even further with MSAN (typically 3-5x).

Standard x86_64 memory models address code within 2GB region

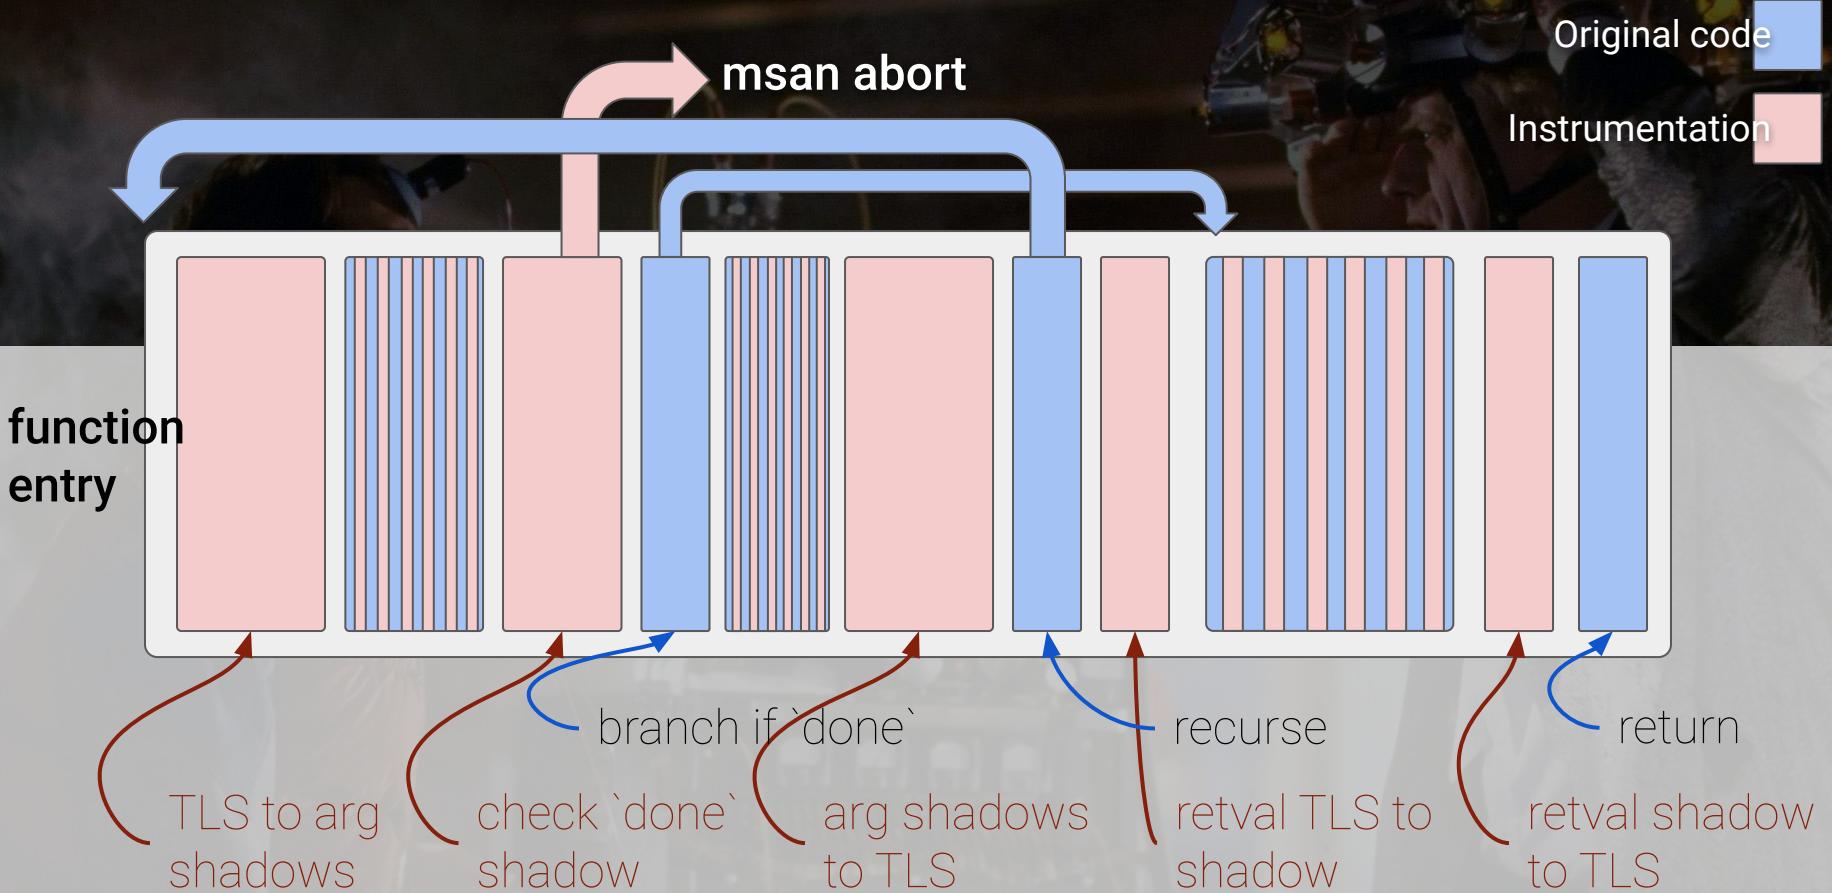
Large Google binaries + MSAN may approach/exceed 2GB

Why optimize it? Runtime

Instrumented Google3 testing costs
\$\$ in compute resources

Instrumented binary runtime directly
affects fuzzing ops/second. In the long run,
faster binary => more bugs caught!

MSAN Instrumentation - a Primer



Optimizing MSAN = Cutting Overhead

Instrumentation

Call site
shadows

Origin
tracking

Function
body
shadows

Runtime support

Stack
unwinding

LibCall
interception

Origin
chaining

Note: discussing optimizations for MSAN+origin tracking

Case Study - Binary Search Sqrt

```
sqrt(int):
    mov eax, edi
    shr eax, 31
    add eax, edi
    sar eax
    movsxd rdx, eax
    imul rdx, rdx
    movsxd r8, edi
    sub rdx, r8
    je .LBB0_7
    xor esi, esi
.LBB0_2:
    test rdx, rdx
    jg .LBB0_3
    mov esi, eax
    jmp .LBB0_5
.LBB0_3:
    mov edi, eax
    .LBB0_5:
        .LBB0_7:
            ret
```

```
int sqrt(int x) {
    int min_r = 0, max_r = x;
    long long old_loss, loss = 0;
    int r;

    #pragma nounroll
    while (true) {
        r = (max_r - min_r) / 2 + min_r;
        old_loss = loss;
        loss = (long long)r * r - x;
        if (!loss || old_loss == loss)
            return r;
        if (loss > 0) {
            max_r = r;
        } else {
            min_r = r;
        }
    }
    return r;
}
```

```

sqrt(int):
pushq %rbp
pushq %r15
pushq %r14
pushq %r13
pushq %r12
pushq %rbx
subq $56, %rsp
movq

```

```

    msan param tls@GOTTPPO
    FF(%rip), %rax
movslq %fs:(%rax), %rdx

```

```

movl %edi, %eax
shrl $31, %eax
addl %edi, %eax
sarl %eax
movslq %eax, %r10
imulq %r10, %r10
movslq %edi, %rcx
movq %rcx, 48(%rsp)
subq %rcx, %r10
movq

```

```

    msan param origin tls
    @GOTTPOFF(%rip), %rcx

```

```

movl %fs:(%rcx), %ecx
movl %ecx, 12(%rsp)

```

```

movq %rdx, 24(%rsp)

```

```

testq %rdx, %rdx
je .LBB0_2

```

```

movq 24(%rsp), %rcx
notq %rcx

```

```

andq %r10, %rcx
je .LBB0_13

```

```

.LBB0_2:
testq %r10, %r10
je .LBB0_3
xorl %ebp, %ebp

```

```

movq 24(%rsp), %rdx
movq %rdx, %r12

```

```

movl 12(%rsp), %esi
movl %esi, %ecx

```

```

movl %edx, %r15d

```

```

movl %esi, %r13d
movl $0, 16(%rsp)

```

```

xorl %ebx, %ebx
movl %edx, %r9d
movl %esi, 20(%rsp)

```

```

movl %ebp, %r8d

```

```
.LBB0_6:
```

```

movl %edi, %edx

```

```

xorl %eax, %edx

```

```

movl %r9d, %esi

```

```

orl %r15d, %esi

```

```

orl %edx, %esi

```

```

movl %ebx, %edx

```

```

xorl %eax, %edx

```

```

orl %r15d, %r8d

```

```

orl %edx, %r8d

```

```

testq %r10, %r10

```

```

jg .LBB0_7

```

```

movl %eax, %ebx

```

```

movl 20(%rsp), %edx

```

```

jmp .LBB0_9

```

```
.LBB0_7:
```

```

movl %r15d, %r9d

```

```

movl %r13d, %edx

```

```
# ... (col omitted)
```

```

cmovnel 12(%rsp), %ecx
subq 48(%rsp), %r14

```

```

movq %rbp, %rdx

```

```

notq %rdx

```

```

testq %rdx, %r14

```

```

sete %dl

```

```

testl %ebp, %ebp

```

```

setne %dl

```

```

andb %dl, %dl

```

```

testq %r14, %r14

```

```

setne %r11b

```

```

movq %r14, %rdx

```

```

xorq %r10, %rdx

```

```

orq %rbp, %r12

```

```

movq %r10, %rbx

```

```

movq %r12, %r10

```

```

notq %r10

```

```

testq %r10, %rdx

```

```

sete %r10b

```

```

testq %r12, %r12

```

```

setne %dl

```

```

andb %r10b, %dl

```

```

movq %rbx, %r10

```

```

cmpq %r14, %rbx

```

```

setne %bl

```

```

testb %bl, %dl

```

```

jne .LBB0_14

```

```

orb %r11b, %dl

```

```

andb %dl, %dl

```

```

jne .LBB0_14

```

```

testq %r14, %r14

```

```

je .LBB0_4

```

```

movq %rbp, %r12

```

```

movl %r8d, %ebp

```

```

cmpq %r14, %r10

```

```

movq %r14, %r10

```

```

movl 44(%rsp), %edi

```

```

movl 40(%rsp), %ebx

```

```

jne .LBB0_6

```

```

jmp .LBB0_4

```

```
.LBB0_3:
```

```

movq 24(%rsp), %rcx

```

```

movl %ecx, %r15d

```

```

movl 12(%rsp), %r13d

```

```
.LBB0_4:
```

```

movq

```

```

    msan retval tls@GO
    TTPOFF(%rip), %rcx

```

```

movl %r15d,
    %fs:(%rcx)

```

```

movq

```

```

    msan retval origin
    tls@GOTTPOFF(%rip),
    %rcx

```

```

movl %r13d,
    %fs:(%rcx)

```

```

addq $56, %rsp

```

```

popq %rbx

```

```

popq %r12

```

```

popq %r13

```

```

popq %r14

```

```

popq %r15

```

```

popq %rbp

```

```

retq

```

```
.LBB0_14:
```

```

testl %esi, %esi
movl 36(%rsp), %edi
cmovnel %ecx, %edi
testb %dl, %dl
cmovel %ecx, %edi
callq

```

msan warning with
origin_noreturn

```
.LBB0_13:
movl 12(%rsp), %edi
callq

```

msan warning with
_origin_noreturn

Call site shadows

Shadow checking

Shad. propagation

Orig. propagation

-fsanitize=memory

-fsanitize-memory

-track-origins

Cutting instrumentation overhead?

For non-trivial functions, most runtime+space overhead comes from shadow prop.

Every operation needs its own shadow operation!

Can we *factor out* shadow prop?

Eager Checks: Cutting Shadows

Don't pass shadows over TLS; check
them instead!

Force clean param/retval shadows
Let constant propagation do the work
for us!

Eager Checks: Catching Errors Early

Current MSAN defers checking as long as possible
With more checks, bugs emerge nearer to their origin

Some bugs may emerge that MSAN wouldn't have
caught (e.g. atomic_store erasing shadow)

Catches UB from violating language semantics

Easier said than done...

Can't willy nilly check everything -- sometimes passing
uninit values is legal! (retval in C, padded types, etc.)

Language-specific rules. Need some info from the
frontend!

New LLVM parameter attribute: `noundef`
(actually more broadly useful to LLVM pass authors!)

Case Study - Binary Search Sqrt

With eager checking:

of instructions: 153 => 132 (-13.7%)

uninstrumented version: 30 instructions

Object file size: 641B => 545B (-15.0%)

uninstrumented version: 134B

But! Theoretically all shadow prop should go away...

Running opt a second time?

Wow! But it doesn't generalize :(

~~DEAD END~~

uninstrumented	eager + opt
<pre>gtr(%rdx,%rcx) mov %eax,%esi shr %eax,31 add %eax,%esi sarl %eax,%esi movsxq %eax,%rcx imul %rcx,%rax movxdq r8,%rcx add %rax,%rcx sub rdx,0 je .LBB0_7 xor esi,esi .LBB0_2: test rdx,rdx jg .LBB0_3 mov esi,eax jmp .LBB0_5 .LBB0_3: mov edi,eax .LBB0_7: ret</pre>	<pre>_Z4sqrti: movl %edi,%eax shrl \$31,%eax addl %edi,%eax sarl %eax,%eax movslq %eax,%rdx imulq %rax,%rdx movslq %edi,%eax subq %rax,%rdx xorl %esi,%eax .LBB0_5: movl %edi,%ecx subl %esi,%ecx movl %ecx,%eax shl \$31,%eax addl %ecx,%eax sarl %eax,%eax addl %esi,%eax movslq %eax,%rcx imulq %rcx,%rcx subq %r8,%rcx je .LBB0_3 movl %eax,%esi jmp .LBB0_5 .LBB0_3: movl %eax,%edi jmp .LBB0_2 .LBB0_7: retq</pre>

Extra wins on the margins #1

MSAN checked shadows with a call to `_msan_warning` and passed origins over TLS to the runtime

Checks were now extremely common, TLS store insts are bulky => have `_msan_warning` take origin as a formal parameter instead

Measured 5% space savings on Clang

Extra wins on the margins?

~~DEAD END~~

MSAN poisons any local variables on the stack, usually inline (e.g. with memset)

Could force it to poison var + set origin with one rt call (outlining alloca poisoning)

1.5% space savings... 20% increase in runtime 😕

Extra wins on the margins #3

Large stack variables call `_msan_set_alloca_origin` to register their stack trace

If variable immediately stored to, this call is redundant

Make MSAN optimize it away?

Variable space savings; 8% for grep, <1% for Clang

Bookkeeping and performance

MSAN has to track stack traces where uninitialized memory comes from, and chain together their “origins”

Some programs spend large % of runtime on this bookkeeping (for Clang, ~30%)

% depends on size of program, amount of mallocs done, amount of large allocas done

Bookkeeping process

Malloc

Alloc

Poisoned Store

Unwind stack trace

expensive!

Unwind stack trace

Stack Depot, acquire ID

expensive!

Stack Depot, acquire ID

expensive!

Chain Depot, put edge

Cached stack unwinding?

What's our *index*?

Idea 1: some combination of PC and SP (and/or BP)

Too many collisions! Need some extra info...

Idea 2: every callee XORs its UUID into a TLS variable at

entry, and again at exit (construct a call chain hash)

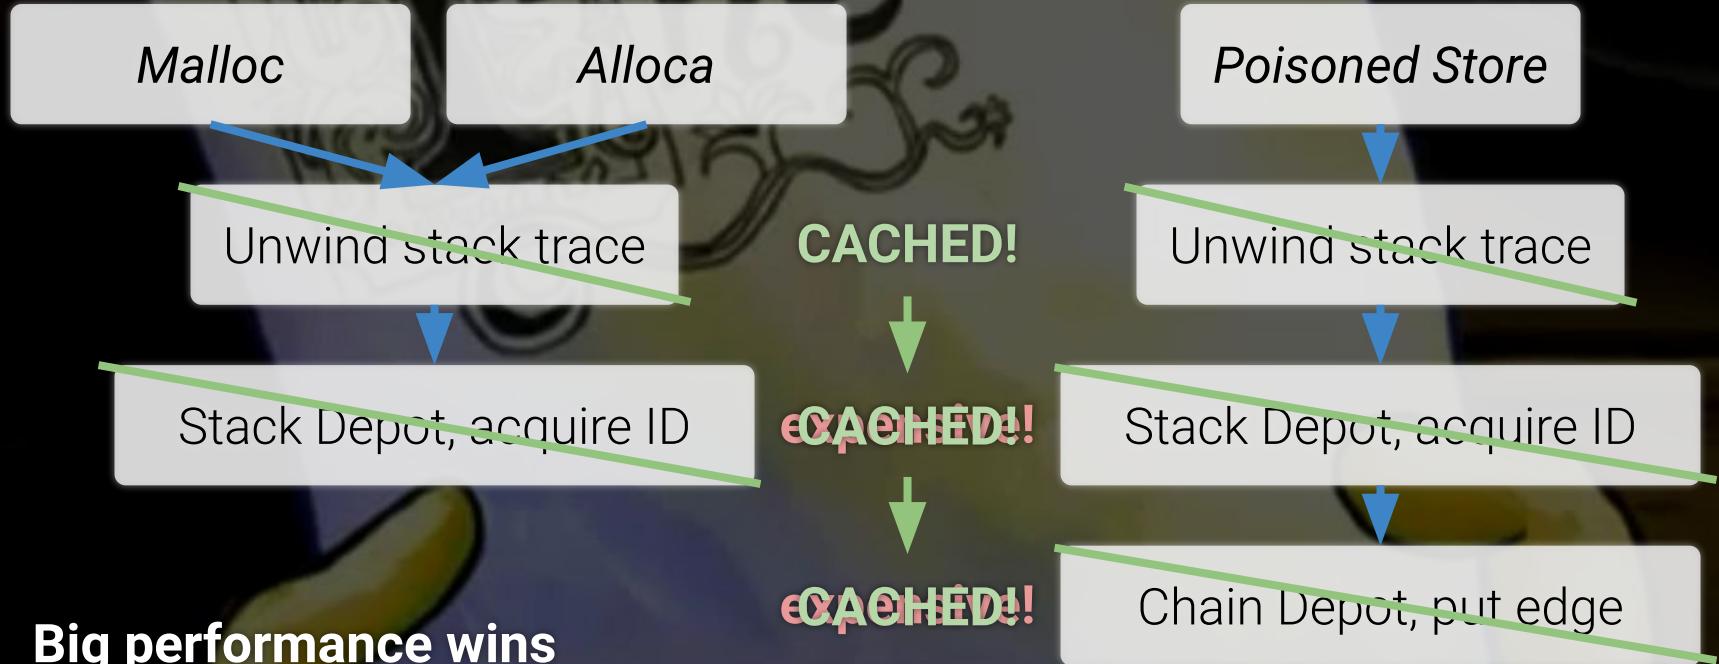
Collision if same parent calls same child twice...

Cached stack unwinding?

Idea 3: same idea, but construct a stack trace hash by XORing return-addr and then rotating by 5. Much more precise!

Then we can use this *trace hash* to index into a direct-mapped *trace cache*

Cached bookkeeping process



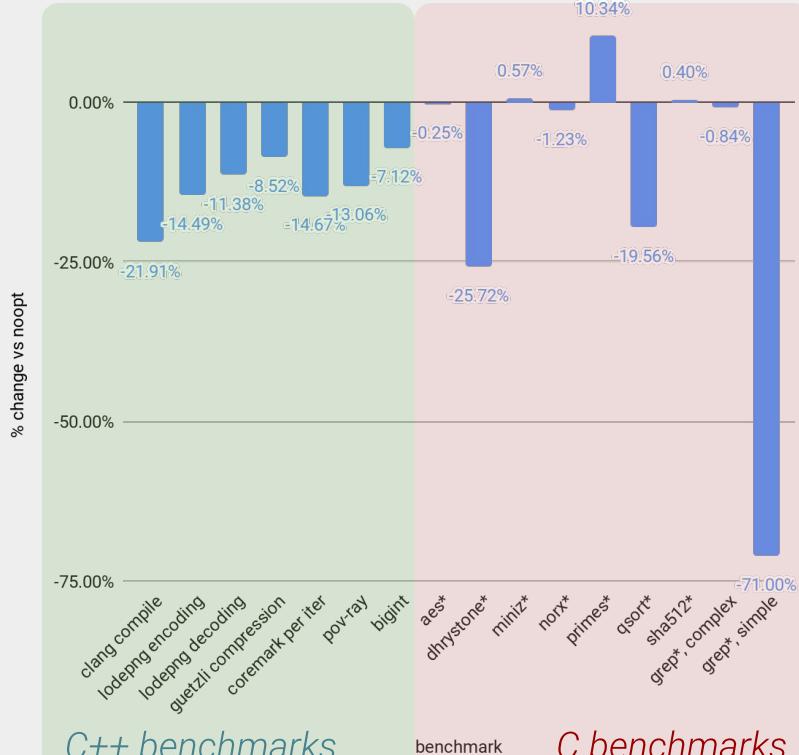
Big performance wins
for Clang, >17%

Combined Results

% change in binary size (excluding fixed cost)

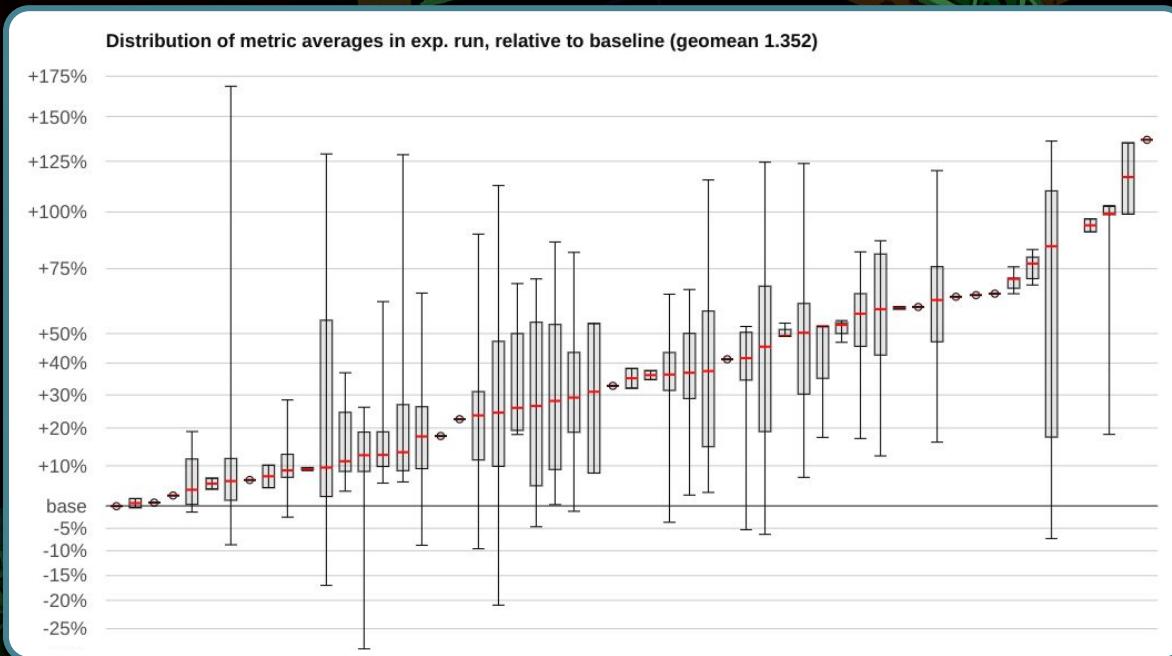


% change in runtime



Eager checks most effective in C++
Cached traces best with large programs

Internal benchmarks? Promising!



Execution frequency benchmarks to evaluate compiler releases. Higher=better!

More Google-Specific Results

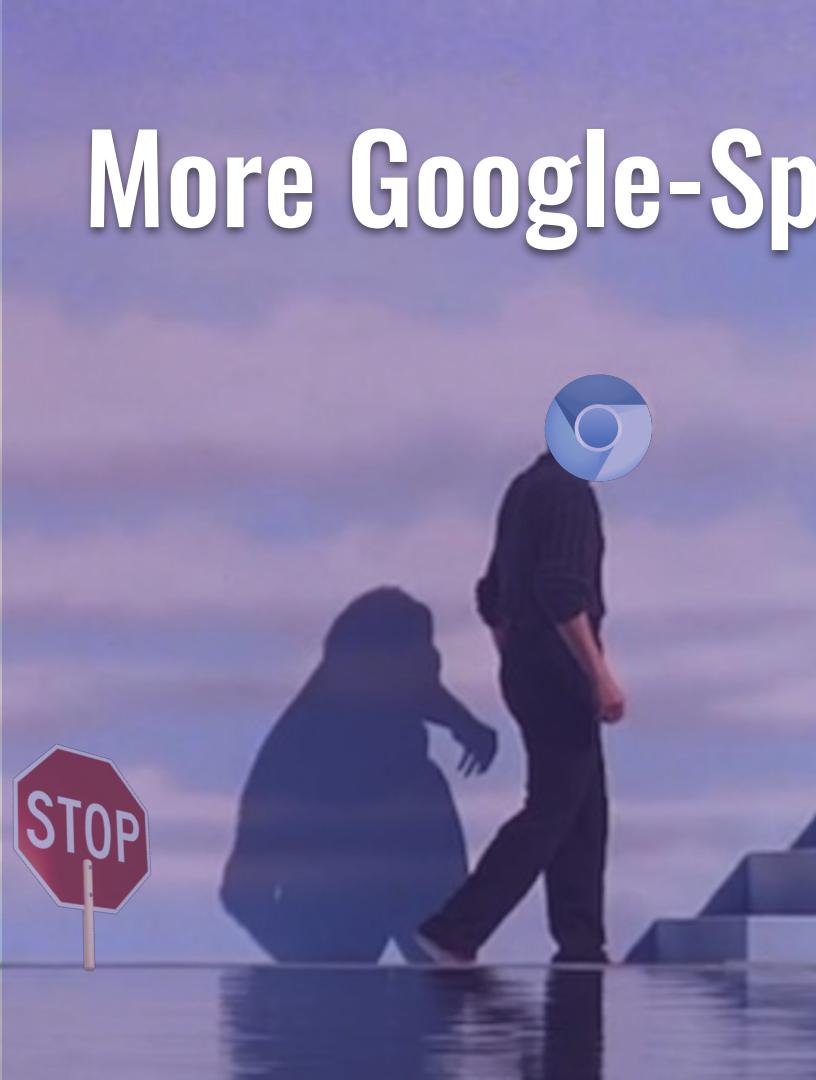
Chromium binary size

0.749 GB => 0.592 GB. **-21.1%**

Internal RPC benchmark

61k => 75.5k req/sec. **+23.7%**

Internal server binary size...



Reining in Binary Size



Google internal server binary

1.900 GB => 1.687 GB

Moderate savings: **-11.2%**

Enough to keep it comfortably
in the 2GB code size limit!

(no origin tracking, as the binary would be too big)

Challenges

Sadly, optimizations don't necessarily stack

Precise standard compliance for `noundef` rules

Community agreement around `noundef` changes

Integrating Clang/MSAN to ad-hoc build systems :(

Next steps, next ideas

Shadow frame pointer

Computing shadow pointer for stack item
is a common and space-inefficient op

Reserve a register as shadow fp instead?

Next steps, next ideas

Noundef improvements

Certain classes of types are never `noundef` in the initial implementation (for simplicity)

Structs, coerced types: if these were marked noundef they could be checked in MSAN

Acknowledgements



Hosts **Evgenii Stepanov and Vitaly Buka**

Huge support and guidance!

Richard Smith, Johannes Doerfert, other reviewers

Patience and dedication to correctness

Rest of the dynamic tools team!